

Automated Kantian Ethics

A SENIOR THESIS PRESENTED

BY

LAVANYA SINGH

TO

THE DEPARTMENTS OF COMPUTER SCIENCE AND PHILOSOPHY

IN PARTIAL FULFILLMENT OF THE REQUIREMENTS

FOR THE DEGREE OF

BACHELOR OF ARTS WITH HONORS

HARVARD UNIVERSITY

CAMBRIDGE, MASSACHUSETTS

MAY 2022

ABSTRACT

AI agents are beginning to make decisions without human supervision in increasingly consequential contexts like healthcare, policing, and driving. These decisions are inevitably ethically tinged, but most AI agents navigating the world today are not explicitly guided by ethics. Regulators, philosophers, and computer scientists are raising the alarm about the dangers of unethical artificial intelligence, from science-fiction killer robots to criminal sentencing algorithms prejudiced against people of color. This is spurring interest in automated ethics, or the development of machines that can perform ethical reasoning. Much prior work in automated ethics approaches the problem from a computational perspective and rarely engages with philosophical literature on ethics, despite its relevance to the development of AI agents that can responsibly navigate the world. If automated ethics draws on sophisticated philosophical literature, the ethical reasoning underlying such decisions will be more nuanced, precise, consistent, and trustworthy. However, faithfully translating complex ethical theories from natural language to the rigid syntax of a computer program poses technical and philosophical challenges.

In this thesis, I present an implementation of automated Kantian ethics that is faithful to the Kantian philosophical tradition. Given an appropriately represented action and minimal factual background, my system can judge the action as morally obligatory, permissible, or prohibited. To accomplish this, I formalize Kant’s categorical imperative, or moral rule, in Carmo and Jones’s Dyadic Deontic Logic, implement this formalization in the Isabelle/HOL theorem prover, and develop a testing framework to evaluate how well my implementation coheres with expected properties of Kantian ethics, as established in the literature. I also use my system to derive philosophically sophisticated and nuanced solutions to two central ethical dilemmas in Kantian ethical literature: the difference between lying and joking and the example of a murderer knocking on your door asking about the location of their intended victim. Finally, I examine the philosophical implications of my system, exploring its limitations and its potential to help AI agents, academic philosophers, and ordinary human beings better reason about ethics. This work serves as an early proof-of-concept for philosophically

mature AI agents and is one step towards the development of responsible, trustworthy artificial intelligence.

Contents

1	Introduction	I
2	System Components	7
2.1	Choice to Automate Kantian Ethics	7
2.1.1	Consequentialism	8
2.1.2	Virtue Ethics	11
2.1.3	Kantian Ethics	13
2.1.4	The Formula of Universal Law	17
2.2	Dyadic Deontic Logic	18
2.3	Isabelle/HOL	20
2.3.1	System Definition	21
2.3.2	Syntax	22
3	Implementation Details	25
3.1	Formalization and Implementation of the FUL	25
3.1.1	Logical Background	25
3.1.2	Maxims	26
3.1.3	Practical Contradiction Interpretation of the FUL	29
3.1.4	Formalizing the FUL	34
3.2	Tests	37
4	Applications	44
4.1	Lies and Jokes	45
4.2	Lying to a Liar	52
5	Discussion	58
5.1	Automated Moral Agents in Practice	58
5.2	Computational Ethics	62

5.2.1	Example of a Philosophical Insight: Well-Formed Maxims	63
5.2.2	An Argument For Computational Ethics	69
5.3	Automating Ordinary Ethical Reasoning	71
5.4	Theoretical Objections to Automating Kantian Ethics	74
5.5	Related Work	80
5.6	Conclusion	82
A	Appendix	84
A.1	Alternate Definitions of a Maxim	84
A.1.1	Korsgaard’s Act-Goal View	84
A.1.2	Kitcher’s View on Motives	85
A.2	Kroy’s Formalization	86
A.2.1	Logical Background	86
A.2.2	Kroy’s formalization of the FUL	89
A.2.3	Testing Kroy’s Formalization	90
A.2.4	Metaethical Tests	95
A.2.5	Miscellaneous Tests	98
A.3	Testing Un-Universalizable Actions	100
A.4	Ethics for Ordinary People	105
A.5	Academic Ethics	108

I Introduction

As AI agents become more sophisticated and less dependent on humans, interest begins to mount in the development of computers that can perform ethical reasoning, also known as automated moral agents. AI agents are making decisions in increasingly consequential contexts, such as healthcare, driving, and criminal sentencing, and therefore must perform ethical reasoning in order to navigate moral dilemmas. For example, self-driving cars may face less extreme versions of the following moral dilemma: an autonomous vehicle approaching an intersection fails to notice pedestrians in the crosswalk until it is too late to brake. The car can either continue on its course, running over and killing three pedestrians, or it can swerve to hit the car in the next lane, killing the single passenger inside it. While this example is (hopefully) not typical of the operation of a self-driving car, every decision that such an AI agent makes, from avoiding congested freeways to carpooling, is morally tinged. Not only do AI agents routinely make decisions with ethical implications without explicitly performing ethical reasoning, in many cases they do so without human supervision. For example, the Allegheny Family Screening tool can automatically trigger an investigation into a potential case of child neglect, a decision that can uproot entire families and is known to be biased against poor people of color ([Eubanks, 2018](#)). This motivates the need for machine ethics (also called automated ethics), or the study of how to develop machines that can perform robust, sophisticated ethical reasoning.

Machine ethicists recognize the need for automated ethics and have made both theoretical (([Awad et al., 2020](#)), ([Davenport, 2014](#)), ([Wallach and Allen, 2008](#)), ([Gabriel, 2020](#))) and practical progress (([Arkoudas et al., 2005](#)), ([Cervantes et al., 2013](#)), ([Jiang et al., 2021](#)), ([Winfield et al., 2014](#))) towards automating ethics. However, prior work in machine ethics using popular ethical theories like deontology (([Anderson and Anderson, 2014](#)), ([Anderson and Anderson](#))), consequentialism (([Abel et al., 2016](#)), ([Anderson et al., 2004](#)), ([Cloos, 2005](#))), and virtue ethics ([Berberich and Diepold, 2018](#)) rarely engages with philosophical literature and thus often misses philosophers' insights. Even the above example of the malfunctioning

self-driving car is an instance of Phillipa Foot’s trolley problem, in which a bystander watching a runaway trolley can pull a lever to kill one instead of three (Foot, 1967). Decades of philosophical debate have developed ethical theories that can offer nuanced and consistent answers to the trolley problem. Like the trolley problem, the moral dilemmas that artificial agents face are not entirely new, so solutions to these problems should take advantage of philosophical progress. Philosophers are devoted to the creation of better ethical theories, so the more faithful that automated ethics is to philosophical literature, the more nuanced, precise, consistent, and therefore trustworthy it will be.

A lack of engagement with prior philosophical literature also makes automated moral agents less explainable, or interpretable by human observers. One example of this is Delphi, an implementation of automated ethics that uses deep learning to make moral judgements based on a training dataset of ethical decisions made by humans (Jiang et al., 2021). Early versions of Delphi gave unexpected results, such as declaring that the user should commit genocide if it makes everyone happy (Vincent, 2021). Moreover, because no explicit ethical theory underpins Delphi’s judgements, human beings cannot analytically determine why Delphi thinks genocide is obligatory or where its reasoning may have gone wrong. Machine learning approaches like Delphi often cannot explain their decisions to a human being and, in the extreme case, are black box algorithms. This reduces human trust in a machine’s controversial ethical judgements. If a machine prescribes killing one person to save three without justifying this decision, it is difficult to trust this judgement. The high stakes of automated ethics require explainability to build trust and catch mistakes.

While automated ethics should draw on philosophical literature, in practice, automating an ethical theory is a technical and philosophical challenge. Intuitive computational approaches explored previously, such as representing ethics as a constraint satisfaction problem (Dennis et al., 2016) or reinforcement learning algorithm (Abel et al., 2016), fail to capture philosophically plausible ethical theories. For example, encoding ethics as a Markov Decision Process assumes that ethical reward can be aggregated according to some discounted sum¹, but

¹Markov Decision Processes usually assume that the total reward of a system is the discounted sum of the reward at each state, given by r_i . Formally, total reward $R = \sum_0^\infty \gamma^i r_i$.

many philosophers reject this notion of aggregation ([Sinnott-Armstrong, 2021](#)). On the other hand, approaches that begin with an ethical theory, instead of a computational method, must contend with the fact that ethical theories are almost always described in natural language and must be made precise enough to represent to a computer. Even once ethics is translated from natural language to program syntax, the factual background given to the machine, such as the description of an ethical dilemma, plays a great role in the machine’s decisions. Another complication is that philosophers do not agree that on a single choice of ethical theory. Even philosophers who agree that a specific ethical theory, like Kantian ethics, is true, still debate the theory’s details.² Moreover, even once reasoning within a particular ethical theory is automated, those who disagree with that theory will disagree with the system’s judgements.

This thesis presents a proof-of-concept implementation of philosophically faithful automated Kantian ethics. I formalize Kant’s categorical imperative, or moral rule, as an axiom in Dyadic Deontic Logic (DDL), a modal logic designed to reason about obligation ([Carmo and Jones, 2013](#)). I implement my formalization in Isabelle/HOL, an interactive theorem prover that can automatically verify and generate proofs in user-defined logics ([Nipkow et al., 2002](#)). Finally, I use Isabelle to automatically prove theorems (such as, “murder is wrong”) in my new logic, generating results derived from the categorical imperative. Because my system automates reasoning in a logic that represents Kantian ethics, it automates Kantian ethical reasoning. Once equipped with minimal factual background, it can classify actions as prohibited, permissible or obligatory. I make the following contributions:

1. In Section [2.1](#), I make a philosophical argument for why Kantian ethics is the most natural of the three major ethical traditions (deontology, virtue ethics, utilitarianism) to formalize.
2. In Section [3](#), I present a formalization of the practical contradiction interpretation of Kant’s Formula of Universal Law in Dyadic Deontic Logic. I implement this formalization in the Isabelle/HOL theorem prover. My implementation includes axioms and definitions such that my system, when given an appropriately represented input, can

²For examples of these debates in the case of Kantian ethics, see Section Joking and Section Murderer.

prove that the input is permissible, obligatory, or prohibited. It can also return a list of facts used in the proof and, in some cases, an Isar-style human readable proof.

3. In Sections 4.1 and 4.2, I demonstrate my system’s power and flexibility by using it to produce nuanced answers to two well-known Kantian ethical dilemmas. I show that, because my system draws on definitions of Kantian ethics presented in philosophical literature, it is able to perform sophisticated moral reasoning.
4. In Section 3.2, I present a testing framework that can evaluate how faithful an implementation of automated Kantian ethics is. My framework includes meta-ethical tests and application tests inspired by philosophical literature. This testing framework shows that my formalization substantially improves on prior work and can be generalized to evaluate any implementation of automated Kantian ethics.
5. In Section 5.2, I present new ethical insights discovered using my system and argue that computational methods like the one presented in this paper can help philosophers address ethical problems. Not only can my system help machines reason about ethics, but it can also help philosophers make philosophical progress.

Deontic logic is a modal logic that can express obligation, or morally binding requirements. Traditional modal logics include the necessitation operator, denoted as \Box . In modal logic using the Kripke semantics, $\Box p$ is true at world w if p is true at all worlds that neighbor w (Cresswell and Hughes, 1996). Modal logics also contain the possibility operator \Diamond , where $\Diamond p \longleftrightarrow \neg(\Box(\neg p))$ and operators of propositional logic like $\neg, \wedge, \vee, \rightarrow$. I use DDL, in which the dyadic obligation operator $O\{A|B\}$ represents the sentence “A is obligated in the context B.” The introduction of context allows DDL to express more nuanced reasoning. DDL is both deontic and modal, so sentences like $O\{A|B\}$ are terms that can be true or false at a world. For example, the sentence $O\{\text{steal}|\text{when rich}\}$ is true at a world if stealing when rich is obligated at that particular world.

I automate Kantian ethics because it is the most natural to formalize, as I argue in Section 2.1. Kant presents three versions of a single moral rule, known as the categorical im-

perative, from which all moral judgements can be derived. I implement a version of this rule called the Formula of Universal Law (FUL), which states that people should only act on those principles that can be acted on by all people without contradiction. For example, in a world where everyone falsely promises to repay a loan, lenders will no longer believe these promises and will stop offering loans. Therefore, not everyone can simultaneously falsely promise to repay a loan, so the FUL thus prohibits this act.

Prior work by Benzmüller, Farjami, and Parent ([Benzmüller et al., 2019](#); [Benzmüller et al., 2021](#)) implements DDL in Isabelle/HOL and I add the Formula of Universal Law as an axiom on top of their library. The resulting Isabelle theory can automatically or semi-automatically generate proofs in a new logic that has the categorical imperative as an axiom. Because proofs in this logic are derived from the categorical imperative, they judge actions as obligated, prohibited, or permissible. Moreover, because interactive theorem provers are designed to be interpretable, my system is explainable. Isabelle can list the axioms and facts it used to generate an ethical judgement, and, in some cases, construct human-readable proofs. In Chapter 4, I use my system to derive sophisticated solutions to two ethical dilemmas often used in critiques of Kantian ethics. Because my system is faithful to philosophical literature, it is able to provide nuanced answers to these paradoxes that require a deep understanding of Kantian ethics.

In addition to presenting the above logic and implementation, I also contribute a testing framework that evaluates how well my formalization coheres with philosophical literature. I formalize expected properties of Kantian ethics as sentences in my logic, such as the property that obligations cannot contradict each other. I represent each of these properties as a sentence in my logic that my system should be able to prove or refute. To run the tests, I use Isabelle to automatically find proofs or countermodels for the test statements. For example, my implementation passes the contradictory obligations test because it is able to prove the sentence $\neg(O\{A|B\} \wedge O\{\neg A|B\})$. I find that my system outperforms the control group of raw DDL, without any moral axioms added, and Moshe Kroy’s prior attempt at formalizing Kantian ethics in deontic logic ([Kroy, 1976](#)).

As it stands, my implementation can evaluate the moral status of sentences represented in my logic. Given an appropriate input, my project returns a value indicating if the action is obligatory (its negation violates the FUL), permissible (consistent with the FUL), or prohibited (violates the FUL) by proving or refuting a theorem in my logic. Because the logic underlying these judgements is grounded in philosophical literature, my system can correctly resolve complex moral dilemmas with only relatively uncontroversial facts about the world. Moreover, because I automate an explicit ethical theory, the ethical reasoning underlying my system's judgements is interpretable by a human being. I implement this ethical theory using the Isabelle/HOL interactive theorem prover, which can list the axioms and theorems used in a proof, so my system is explainable.

A machine that can evaluate the moral status of a maxim can not only help machines better reason about ethics, but it can also help philosophers better study philosophy. I argue for “computational ethics,” or the use of computational tools to make philosophical progress, analogous to computational drug discovery or neuroscience. I demonstrate the potential of computational ethics by presenting a philosophical insight about which kinds of maxims are appropriate for ethical consideration that I discovered using my system. The process of building and interacting with a computer that can reason about ethics helped me, a human philosopher, arrive at a philosophical conclusion that has implications for practical reason and philosophy of doubt. Thus, my system can be used in two distinct ways. First, my system can help automated agents navigate the world, which I will refer to as automated ethics or machine ethics interchangeably. Second, my system help human philosophers reason about philosophy, which I call computational ethics. In Section 5.3, I consider how computational ethics can help not only professional philosophers, but also augment the everyday ethical reasoning that ordinary people perform as we navigate the world.

2 System Components

My system consists of three components: an ethical theory (Kantian ethics), a logic in which I formalize this ethical theory (Dyadic Deontic Logic), and an interactive theorem prover in which I implement the formalized ethical theory (Isabelle/HOL). In this section, I describe these components, present the philosophical, logic, and computational background underlying my system, and explain the consequences of each of the three choices I make.

These specific components determine the features and limitations of my implementation of automated ethics, but other choices of components, such as another ethical theory, a different logic, or a different theorem prover could be made. Flaws with these components are merely limitations of my system, but do not indict logic-programming-based automated ethics more generally. My thesis seeks to both present a specific implementation of automated ethics but also to argue for a particular approach to automating ethical reasoning and these choices are relevant to the former goal but not to the latter.

2.1 Choice to Automate Kantian Ethics

In this thesis, I automate Kantian ethics. In 2006, Powers posited that deontological theories are attractive candidates for automation because rules are generally computationally tractable (Powers, 2006, 1). Intuitively, algorithms are rules or procedures for problem solving and Kantian ethics (which is one kind of deontological theory) offers one such procedure for the problem of making ethical judgements. I will make this intuition precise by arguing that Kantian ethics is natural to formalize because it prescribes moral rules that require little additional data about the world and are easy to represent to a computer. I argue that, compared to consequentialism and virtue ethics,³ Kantian ethics is more amenable to automation.

³Technically, virtue ethics and consequentialism are broad ethical traditions, while Kantian ethics is a specific ethical theory within the deontological tradition, but “if any philosopher is regarded as central to deontological moral theories, it is surely Immanuel Kant” (Alexander and Moore, 2021). Out of the three major ethical traditions, deontology has the most central representative in the form of Kant. Deontology’s comparatively greater focus on Kant means that my choice of Kant as a guiding figure is less controversial for deontologists than, for example, the choice of Bentham as the guiding figure of consequentialism.

I do not aim to show that Kantian ethics is the only tractable theory to automate or to present a comprehensive overview of all consequentialist or virtue ethical theories. Instead, I present a sample of some approaches in each tradition and argue that deontology is more straightforward to formalize than these approaches. Insofar as my project serves as an early proof-of-concept, I choose to automate an ethical theory that poses fewer challenges than others. All ethical traditions have debates that an automated ethical system will need to take a stance on, but these debates are less frequent and controversial for Kantian ethics than for consequentialism and virtue ethics.

I first present consequentialism, then virtue ethics, and finally Kantian ethics. For each tradition, I present a crash course for non-philosophers and then explain some obstacles to automation, arguing that these obstacles are weakest in the case of Kantian ethics.

2.1.1 Consequentialism

A consequentialist ethical theory is an ethical theory that evaluates an action by evaluating its consequences.⁴ For example, utilitarianism is a form of consequentialism in which the moral action is the action that produces the most good (Driver, 2014). The focus on the consequences of action distinguishes consequentialists from deontologists, who derive the moral worth of an action from the action itself. Some debates in the consequentialist tradition include which consequences matter, what constitutes a “good” consequence, and how we can aggregate the consequences of an action over all the individuals involved.

Which Consequences Matter

Because consequentialism evaluates the state of affairs following an action, this kind of ethical reasoning requires more knowledge about the state of the world than Kantian ethics. Consequentialism requires knowledge about some or all consequences following an action. This requires that an automated ethical system somehow collect a subset of the infinite consequences of following an action, a difficult, if not impossible, task. Moreover, compiling this

⁴There is long debate about what exactly makes an ethical theory consequentialist (Sinnott-Armstrong, 2021). For this thesis, I focus on theories that place the moral worth of an act in its the consequences.

database of consequences requires answering difficult questions about which consequences were actually caused ⁵ by an action and determining the state of the world before and after an action. As acts become more complex and affect more people, the computational time and space required to calculate and store their consequences increases. Kantian ethics, on the other hand, does not suffer this scaling challenge because it evaluates the acts themselves, and acts that affect 1 person and acts that affect 1 million people share the same representation.

The challenge of representing the circumstances of action is not unique to consequentialism, but is particularly acute in this case. Kantian ethicists robustly debate which circumstances of an action are “morally relevant” when evaluating an action’s moral worth.⁶ Because deontology merely evaluates a single action, the surface of this debate is much smaller than debates about circumstances and consequences in a consequentialist system. An automated consequentialist system must make such judgements about the act itself, the circumstances in which it is performed, and the circumstances following the act. All ethical theories relativize their judgements to the situation in which an act is performed, but consequentialism requires far more knowledge about the world than Kantian ethics.

Theory of the Good

An automated consequentialist reasoner must also take a stance on the debate over what qualifies as a “good consequence,” or what the theory of the good is. For example, hedonists associate good with the presence of pleasure and the absence of pain, while preference utilitarians believe that good is the satisfaction of desire. Other consequentialists, like Moore, adopt a pluralistic theory of value, under which many different kinds of things are good for different reasons (Moore, 1903).

Most theories of the good require that a moral reasoner understand complex features about individuals’ preferences, desires, or sensations in order to evaluate a moral action, mak-

⁵David Hume argues that many straightforward accounts of causation face difficulties (Hume, 2007), and philosophers continue to debate the possibility of knowing an event’s true cause. Kant even argued that first causes, or noumena, are unknowable by human beings (Stang, 2021).

⁶Powers (2006) identifies this as a challenge for automating Kantian ethics and briefly sketches solutions from O’Neill (1990), Silber (1974), and Rawls (1980). For my approach to morally relevant circumstances, see Section 3.1.2.

ing automated consequentialist ethics difficult. Evaluating a state of affairs requires many controversial judgements about whether a state of affairs actually satisfies the relevant criteria for goodness. Perfect knowledge of tens of thousands of people's pleasure or preferences or welfare or rights is difficult, if not impossible.⁷ Either a human being assigns values to states of affairs, which doesn't scale, or the machine does, which requires massive common-sense and increases room for doubting the system's judgements. This may be a tractable problem, but it is much more difficult than the equivalent Kantian task of formulating and evaluating an action.

Aggregation

Once an automated consequentialist agent assigns a goodness measurement to each person in a state of affairs, it must also calculate an overall goodness measurement for the state of affairs. One approach to assigning this value is to aggregate each person's individual goodness score into one complete score for a state. The more complex the theory of the good, the more difficult this aggregation becomes. For example, pluralistic theories struggle to explain how different kinds of value can be compared (Sinnott-Armstrong, 2021). How do we compare one unit of beauty to one unit of pleasure? Resolving this debate requires that an automated reasoner choose one specific aggregation algorithm, but those who disagree with this choice will not trust the reasoner's moral judgements. Moreover, for complex theories of the good, this aggregation algorithm may be complex and may require a lot of data.

To solve this problem, some consequentialists reject aggregation entirely and instead prefer wholistic evaluations of a state of affairs. While this approach no longer requires that an aggregation algorithm, an automated ethical system still needs to calculate a goodness measurement for a state of affairs. Whereas before the system could restrict its analysis to a single person, the algorithm must now evaluate an entire state wholistically. As consequentialists modulate between aggregation and wholistic evaluation, they face a tradeoff between the difficulty of aggregation and the complexity of goodness measurements for large states of affairs.

Prior Attempts to Formalize Consequentialism

⁷Even if it were possible, collecting this kind of data poses privacy and surveillance risks.

Because of its intuitive appeal, computer scientists have tried to formalize consequentialism in the past. These efforts cannot escape the debates outlined above. For example, Abel et al. represent ethics as a Markov Decision Process (MDP), with reward functions customized to particular ethical dilemmas (Abel et al., 2016, 3). While this is a convenient representation, it either leaves unanswered or takes implicit stances on the debates above. It assumes that consequences can be aggregated just as reward is accumulated in an MDP.⁸ It leaves open the question of what the reward function is and thus leaves the theory of the good, arguably the defining trait of consequentialism, undefined. Similarly, Anderson and Anderson’s proposal of a hedonistic act utilitarian automated reasoner chooses hedonism⁹ as the theory of the good (Anderson et al., 2004, 2). Again, their proposal assumes that pleasure and pain can be given numeric values and that these values can be aggregated with a simple sum, taking an implicit stance on the aggregation question. Other attempts to automate consequentialist ethics will suffer similar problems because, at some point, a usable automated consequentialist moral agent will need to resolve the above debates.

2.1.2 Virtue Ethics

Virtue ethics places the virtues, or traits that constitute a good moral character and make their possessor good, at the center (Hursthouse and Pettigrove, 2018). For example, Aristotle describes virtues as the traits that enable human flourishing. Just as consequentialists define “good” consequences, virtue ethicists present a list of virtues, such as the Buddhist virtue of equanimity (McRae, 2013). An automated virtue ethical agent will need to commit to a list and definitions of the virtues, a controversial choice. Virtue ethicists robustly debate which traits qualify as virtues, what each virtue actually means, and what kinds of feelings or attitudes must accompany virtuous action.

Another difficulty with automating virtue ethics is that the unit of evaluation for virtue ethics is often a person’s entire moral character. While Kantians evaluate the act itself, virtue

⁸Generally, reward for an MDP is accumulated according to a “discount factor” $\gamma < 1$, such that if r_i is the reward at time i , the total reward is $\sum_{i=0}^{\infty} \gamma^i r_i$.

⁹Recall that hedonism views pleasure as good and pain as bad.

ethicists evaluate the actor’s moral character and their disposition towards the act. If states of affairs require complex representations, an agent’s ethical character and disposition are even more difficult to represent to a computer. This is more than just a data-collecting problem; it is a conceptual problem about the formal nature of moral character. Formalizing the concept of character appears to require significant philosophical and computational progress, whereas Kantian ethics immediately presents a formal rule to implement.

Prior Work in Machine Learning and Virtue Ethics

One potential appeal of virtue ethics is that many virtue ethical theories involve some notion of moral habit, which seems to be amenable to a machine learning approach. Aristotle, for example, argued that cultivating virtuous action requires making such action habitual through moral education ([Aristotle, 1951](#)). This implies that ethical behavior can be learned from some dataset of virtuous acts, such as those that an ideally virtuous agent would undertake. These theories seem to point towards a machine learning approach to automated ethics, in which ethics is learned from a dataset of acts tagged as virtuous or not virtuous.

Just as prior work in consequentialism takes implicit or explicit stances on debates in consequentialist literature, so does work in machine learning-based virtue ethics. For example, the training dataset with acts labelled as virtuous or not virtuous will contain an implicit view on what the virtues are and how certain acts impact an agent’s moral character. Because there is no canonical list of all virtues that virtue ethicists accept, this implicit view will likely be controversial. Even virtue ethicists agree that certain traits, like courage, are virtues debate the exact definitions of these traits.

Machine learning approaches like the Delphi system ([Jiang et al., 2021](#)) mentioned in Chapter 1 also may suffer explainability problems that my logic-programming, theorem-prover approach does not face. Many machine learning algorithms cannot sufficiently explain their decisions to a human being, and often find patterns in datasets that don’t cohere with the causes that a human being would identify ([Puiutta and Veith, 2020](#)). While there is significant activity and progress in explainable machine learning, interactive theorem provers are designed to be explainable at the outset. Isabelle can show the axioms and lemmas it used in

constructing a proof, allowing a human being to reconstruct the proof independently if they wish. This is not an intractable problem for machine learning approaches to computational ethics, but is one reason to prefer logical approaches.¹⁰

2.1.3 Kantian Ethics

In this paper I focus on Kantian ethics. Kant’s theory is centered on practical reason, which is the kind of reason that we use to decide what to do and the source of our agency. In *The Groundwork of the Metaphysics of Morals*, Kant explains that rational beings are unique because we act “in accordance with the representations of laws” (Kant, 1785, 4:412). In contrast, a ball thrown into the air acts according to the laws of physics. It cannot ask itself, “Should I fall back to the ground?” It simply falls. A rational being, on the other hand, can ask, “Should I act on this reason?” As Korsgaard describes it, when choosing which desire to act on, “it is as if there is something over and above all of your desires, something which is you, and which chooses which desire to act on” (Korsgaard and O’Neill, 1996, 100). Rational beings are set apart by this reflective capacity. We are purposive and our actions are guided by practical reason. We have reasons for acting, even when these reasons are opaque to us. This reflective choosing, or operation of practical reason, is what Kant calls the will.

The will operates by adopting, or willing, maxims, which are its perceived reasons for acting. Kant defines a maxim as the “subjective principle of willing,” or the reason that the will *subjectively* gives to itself for acting (Kant, 1785, 16 footnote 1). Many philosophers agree that a maxim consists of some combination of circumstances, act, and goal.¹¹ One example of a maxim is “When I am hungry, I will eat a doughnut in order to satisfy my sweet tooth.” When an agent wills this maxim, they decide to act on it. They commit themselves to the end in the maxim (e.g. satisfying your sweet tooth). They represent their action, to themselves, as following the principle given by this maxim. Because a maxim captures an agent’s principle

¹⁰This argument about explainability is in the context of virtue ethics and machine learning. It also applies to a broader class of work in automated ethics that uses a “bottom-up” approach, in which a system learns moral judgements from prior judgements. I will extend this argument to general bottom-up approaches in Section Related Work.

¹¹For more discussion of the definition of a maxim, see Section What Is a Maxim

of action, Kant evaluates maxims as obligatory, prohibited, or permissible. He argues that the form of certain maxims requires any rational agent to will them, and these maxims are obligatory.

The form of an obligatory maxim is given by the categorical imperative. An imperative is a command, such as “Close the door” or “Eat the doughnut in order to satisfy your sweet tooth.” An imperative is categorical if it holds unconditionally for all rational agents in all circumstances. Kant argues that the moral law must be a categorical imperative (Kant, 1785, 5). In order for an imperative to be categorical, it must be derived from the will’s authority over itself. Our wills are autonomous, so the only thing that can have unconditional authority over a rational will is the will itself. No one else can tell you what to do because you can always ask why you should obey their authority. The only authority that you cannot question is the authority of your own practical reason. To question this authority is to demand a reason for acting for reasons, which concedes the authority of reason itself (Velleman, 2005, 23). Therefore, the only possible candidates for the categorical imperative are those rules that are required of the will because it is a will.

Armed with this understanding of practical reason, Kant presents the categorical imperative. He presents three “formulations” or versions of the categorical imperative. In this project, I focus on the first formulation, the Formula of Universal Law, and I justify this choice in Section 2.1.4.

The Formula of Universal Law (FUL) states, “act only according to that maxim through which you can at the same time will that it become a universal law” (Kant, 1785, 34). This formulation generates the universalizability test, which we can use to test the moral worth of a maxim by imagining a world in which it becomes a universal law and attempting to will the maxim in that world. If there is a contradiction in willing the maxim in a world in which everyone universally wills the maxim, the maxim is prohibited.

Velleman presents a concise argument for the FUL. He argues that reason is universally shared among reasoners. For example, all reasoners have equal access to the arithmetic logic that shows that “ $2+2=4$ ” (Velleman, 2005, 29). The reasoning that makes this statement true

is not specific to any person, but is universal across people. Therefore, if I have sufficient reason to will a maxim, so does every other rational agent. There is nothing special about the operation of my practical reason. In adopting a maxim, I implicitly state that all reasoners across time also have reason to adopt that maxim. Therefore, because I act on reasons, I must obey the FUL. Notice that this fulfills the above criterion for a categorical imperative: the FUL is derived from a property of practical reason itself and thus derives authority from the will's authority over itself.

Ease of Automation

Kantian ethics is an especially attractive candidate for formalization because the categorical imperative, particularly the FUL, is a property of reason related to the form or structure of a maxim. It does not require any situational knowledge beyond the circumstances included in the maxim itself and thus requires fewer contingent facts than other ethical theories. While other ethical theories often rely on many facts about the world or the actor, a computer evaluating a maxim doesn't require any knowledge about the world beyond what is contained in a maxim. Automating Kantian ethics merely requires making the notion of a maxim precise and representing it to the computer. This distinguishes Kantian ethics from consequentialism and virtue ethics, which require far more knowledge about the world or the agent to reach a moral decision.

A maxim itself is an object with a thin representation for a computer, as compared to more complex objects like states of affairs or moral character. Later in my project, I argue that a maxim can be represented simply as a tuple of circumstances, act, and goal.¹² This representation is simple and efficient, especially when compared to the representations of a causal chain or a state of affairs or moral character. This property not only reduces the computational complexity (in terms of time and space) of representing a maxim, but it also makes the system easier for human reasoners to interact with. A person crafting an input to a Kantian automated agent needs to reason about relatively simple units of evaluation, as opposed to the more complex features that consequentialism and virtue ethics require.

¹²For more, see Section 3.1.2.

Difficulties in Automation

One challenge for automating Kantian ethics is the need for “common-sense”, or factual and situational background. Common-sense is needed when formulating a maxim and when determining if a maxim violates the Formula of Universal Law. Maxims include the circumstances in which they apply and determining which circumstances are “morally relevant” to a maxim requires factual background. Many misunderstandings in Kantian ethics are due to badly formulated maxims, so this question is important for an ethical reasoner to answer. My system does not need to answer this question because I assume a properly crafted maxim as input and apply the categorical imperative to this input. Using my system in a fully automated moral agent will require answering this question, a challenging computational and philosophical task. I discuss this problem in greater detail in Section 3.1.2 and Section 5.1.

Common sense is also relevant when applying the universalizability test itself. Consider the example maxim “When broke, I will falsely promise to repay a loan to get some quick cash.” This maxim fails the universalizability test because in a world where everyone falsely promises to repay loans, no one will believe promises anymore, so the maxim will no longer serve its intended purpose (getting some quick cash). Making this judgement requires understanding enough about the system of promising to realize that it breaks down if everyone abuses it in this manner. This is a kind of common sense reasoning that an automated Kantian agent would need. This need is not unique to Kantian ethics; consequentialists agents need common sense to determine the consequences of an action and virtue ethical agents need common sense to determine which virtues an action reflects. For example, in the case of virtue ethics, in order to see that saving a baby from a lion requires courage, a reasoner must have enough background knowledge to know that lions are scary. Making any ethical judgement requires robust conceptions of the action at hand. Kantian ethics requires far less common sense than a consequentialist or virtue ethical agent.¹³ All moral theories evaluating falsely promising will a robust definition of promising, but consequentialism and virtue ethics will require additional information that Kantian ethics will not. Thus, although the

¹³In Sections Lying and Murderer, I also use my system to demonstrate that Kantian ethics requires relatively thin conceptions of concepts like falsely promising.

need for common sense poses a challenge to automated Kantian ethics, this challenge is more acute for consequentialism or virtue ethics.

2.1.4 The Formula of Universal Law

Earlier I mentioned that Kant presents three formulations, or versions, of what he calls the “supreme law of morality,” but that I focus on the first of these three. In this section, I argue that the Formula of Universal Law, specifically, is the easiest part of Kantian ethics to automate and the most generalizable.

The first formulation of the categorical imperative is the formula of universal law (FUL), which reads, “act only according to that maxim through which you can at the same time will that it become a universal law” (Kant, 1785, 34). The second formulation of the categorical imperative is the formula of humanity (FUH): “So act that you use humanity, in your own person, as well as in the person of any other, always at the same time as an end, never merely as a means.” (Kant, 1785, 41). This formulation is often understood as requiring us to acknowledge and respect the dignity of every other person. The third formulation of the categorical imperative is the formula of autonomy (FOA), which Korsgaard describes as, “we should so act that we may think of ourselves as legislating universal laws through our maxims” (Korsgaard, 2012, 28). While closely related to the FUL, the FOA presents morality as the activity of perfectly rational agents in an ideal “kingdom of ends,” guided by what Kant calls the “laws of freedom.”

I choose to focus on the FUL¹⁴, because it is the most formal and thus the easiest to formalize and implement. Onora O’Neill explains that the formalism of the FUL allows for greater precision in philosophical arguments analyzing its implications and power (O’Neill, 2013, 33). This precision is particularly useful in a computational context because any formalism necessarily makes its content precise. The FUL’s precision reduces ambiguity, allowing me to remain faithful to philosophical literature on Kant. Precision reduces the need to make choices to resolve debates and ambiguities. Minimizing these choices minimizes arbitrariness

¹⁴The FUL is often seen as emblematic of Kantian constructivism (Ebels-Duggan, 2012, 173). My project is not committed to Kantian constructivism.

in my formalization and puts it on solid philosophical footing.

Though Kantians study all formulations of the categorical imperative, Kant argues in *Groundwork* that the three formulations of the categorical imperative are equivalent (Kant, 1785). While this argument is disputed Johnson and Cureton (2021), for those who believe it, the stakes for my choice of the FUL are greatly reduced. If all formulations are equivalent, then a formalization of the FUL lends the exact same power as a formalization of the second or third formulation of the categorical imperative.

Those who do not believe that all three formulations of the categorical imperative are equivalent understand the FUL as the strongest or most foundational, and thus an appropriate initial choice for formalizations. Korsgaard characterizes the three formulations of the categorical imperative according to Rawls' general and special conception of justice. The general conception applies universally and can never be violated, while the special conception represents an ideal for us to live towards. For example, the special conception may require that we prefer some job applicants over others in order to remedy historical injustice, and the general conception may require that such inequalities always operate in the service of the least privileged (Korsgaard, 1986, 19). Korsgaard argues that the Formula of Universal Law represents Kant's general conception of justice, and the Formula of Humanity represents his special conception. The FUL's prescriptions can never be violated, even in the most non-ideal circumstances imaginable, but the FUH is merely a standard to live towards that might not be achieved. Thus, the FUL generates stronger requirements than the other two formulations and reflects the bare minimum standard of Kant's ethics. Because the FUL's prescriptions outweigh those of the other two formulations, automating it creates a functional, minimum ethical theory that can serve as a foundation for implementations of other aspects of Kant's ethics.

2.2 Dyadic Deontic Logic

I formalize Kantian ethics by representing it as an axiom on top of a base logic. In this section, I present the logical background necessary to understand my work and my choice of Dyadic

Deontic Logic (DDL).

Traditional modal logics include the necessitation operator, denoted as \Box . In simple modal logic using the Kripke semantics, $\Box p$ is true at a world w if p is true at all of w 's neighbors, and it represents the concept of necessary truth (Cresswell and Hughes, 1996). These logics usually also contain the possibility operator \Diamond , where $\Diamond p \iff \sim \Box \sim p$. $\Diamond p$ means that the statement p is possibly true, or true at at least one of w 's neighbors. Additionally, modal logics include standard operators of propositional logic like $\sim, \wedge, \vee, \rightarrow$.

A deontic logic is a special kind of modal logic designed to reason about moral obligation. Standard deontic logic replaces \Box with the obligation operator O , and \Diamond with the permissibility operator P (Cresswell and Hughes, 1996). Using the Kripke semantics for O , Op is true at w if p is true at all ideal deontic alternatives to w , and thus represents the concept of moral necessity or necessary requirements. The O operator in SDL takes a single argument (the formula that is obligatory), and is thus called a monadic deontic operator.

While SDL is appreciable for its simplicity, it suffers a variety of well-documented paradoxes, including contrary-to-duty paradoxes.¹⁵ In situations where duty is violated, the logic breaks down and produces paradoxical results. Thus, I use an improved deontic logic instead of SDL for this work.

I use as my base logic Carmo and Jones's Dyadic Deontic Logic (DDL), which improves on SDL (Carmo and Jones, 2013). It introduces a dyadic obligation operator $O\{A|B\}$ to represent the sentence "A is obligated in the context B". The introduction of context allows DDL to gracefully handle contrary-to-duty conditionals by modifying the context. The obligation operator uses a neighborhood semantics, instead of the Kripke semantics (Scott, 1970;

¹⁵The paradigm case of a contrary-to-duty paradox is the Chisholm paradox. Consider the following statements:

1. It ought to be that Tom helps his neighbors
2. It ought to be that if Tom helps his neighbors, he tells them he is coming
3. If Tom does not help his neighbors, he ought not tell them that he is coming
4. Tom does not help his neighbors

These premises contradict themselves, because items (2)-(4) imply that Tom ought not help his neighbors. The contradiction results because the logic cannot handle violations of duty mixed with conditionals. (Chisholm, 1963; R nnedal, 2019)

MONTAGUE, 1970). While Kripke semantics requires that an obligated proposition hold at all worlds, the neighborhood semantics defines a different set of neighbors, or morally relevant alternatives, for each world. To represent this, Carmo and Jones define a function ob that maps a given context (or world) to the propositions that are obligatory at this world, where a proposition p is defined as the worlds at which the p is true. DDL is thus both modal and deontic; statements about obligations are true or false at a world according to the neighborhood semantics, and different obligations may hold at different worlds. This property is particularly relevant to my work because the universalizability test requires reasoning about alternative worlds, such as the world of the universalized maxim.

DDL also includes modal operators. In addition to \Box and \Diamond , DDL also has a notion of actual obligation and possible obligation, represented by operators O_a and O_p respectively. These notions are accompanied by the corresponding modal operators $\Box_a, \Diamond_a, \Box_p, \Diamond_p$. These operators use a Kripke semantics, with the functions av and pv mapping a world w to the set of corresponding actual or possible versions of w . These operators are not relevant to the work in this thesis, but this additional expressivity could be used to extend my project to incorporate more sophisticated ethical concepts like counterfactuals.

For more of fine-grained properties of DDL see Carmo and Jones (2013) or this project’s source code. DDL is a heavy logic and contains modal operators that aren’t necessary for my analysis. While this expressivity is powerful, it may also cause performance issues. DDL has a large set of axioms involving quantification over complex higher-order logical expressions. Proofs involving these axioms will be computationally expensive. I do not run into performance issues in my system, but future work may choose to embed a less complicated logic.

2.3 Isabelle/HOL

The final component of my project is the automated theorem prover I use to automate my formalization. Isabelle/HOL is an interactive proof assistant built on Haskell and Scala (Nipkow et al., 2002). It allows the user to define types, functions, definitions, and axiom systems. It has built-in support for both automatic and interactive/manual theorem proving.

To demonstrate the power and usage of Isabelle and make DDL more precise, I walk through my [reimplementation of Benzmueller, Farjami, and Parent’s implementation of DDL in Isabelle/HOL](#) (Benzmüller et al., 2021; Benzmüller et al., 2019).

2.3.1 System Definition

The first step in embedding a logic in Isabelle is defining the relevant terms and types. Commands to do this include `typedec1`, which declares a new type, `type_synonym`, which defines an abbreviation for a complex type, and `consts`, which defines a constant.

typedec1 *i* — This is an Isabelle comment. *i* is the type for a set of worlds.

— This is a line of actual code used in my implementation. For the rest of the thesis, text typeset like this represents Isabelle code.

type-synonym *t* = (*i* \Rightarrow *bool*) — *t* represents a set of DDL formulae.

— A set of formulae is defined by its truth value at a set of worlds. For example, the set $\{True\}$ is true at any set of worlds.

The *ob* function described in Section 2.2 is used to determine which propositions are obligatory in which contexts. I implement it as a constant. This constant has no meaning (I merely specify the type), but future proofs will specify models for this constant.

consts *ob::t* \Rightarrow (*t* \Rightarrow *bool*) — set of propositions obligatory in this context

— *ob* (*context*) (*term*) is *True* if the term is obligatory in this context

In a semantic embedding, axioms are modelled as restrictions on models of the system. In this case, a model is specified by the relevant accessibility relations (such as *ob*), so it suffices to place conditions on the accessibility relations. Isabelle allows users to create new axiomatizations on top of its base logic (HOL) and use these axioms in proofs. Here’s an example of an axiom:

axiomatization where

ax-5d: $\forall X Y Z. ((\forall w. Y(w) \longrightarrow X(w)) \wedge ob(X)(Y) \wedge (\forall w. X(w) \longrightarrow Z(w)))$

$\longrightarrow ob(Z)(\lambda w.(Z(w) \wedge \neg X(w)) \vee Y(w))$

— If some subset Y of X is obligatory in the context X , then in a larger context Z , any obligatory proposition must either be in Y or in $Z \setminus X$. Intuitively, expanding the context can't cause something unobligatory to become obligatory, so the obligation operator is monotonically increasing with respect to changing contexts.

2.3.2 Syntax

The axiomatization above defines the semantics of DDL and, as demonstrated by the example axiom, is unwieldy. In my work, I mostly perform syntactic proofs, so I must define the syntax of the logic. Isabelle already knows the semantics of the axioms of this logic, so I can define the syntax as abbreviations for different formulas involving the axioms above. Each DDL operator is represented as a HOL formula. Isabelle automatically unfolds formulae defined with the `abbreviation` command whenever they are applied. While the shallow embedding is performant (because it uses Isabelle's original syntax tree), my heavy use of abbreviations may impact the performance of long proofs.

Modal operators will be useful for my purposes, implemented below.

abbreviation $ddlbox::t \Rightarrow t$ (\Box)

where $\Box A \equiv \lambda w. \forall y. A(y)$

— Notice that the necessity operator is an abbreviation, or syntactic sugar for, the higher order logic formula that the proposition holds at all worlds.

abbreviation $ddldiamond::t \Rightarrow t$ (\Diamond)

where $\Diamond A \equiv \neg(\Box(\neg A))$

— Possibility is similarly an abbreviation for a higher order logic formula involving the defined semantics.

The most important operator for my project is the obligation operator, implemented below.

abbreviation $ddlob::t \Rightarrow t \Rightarrow t$ ($O\{-|\cdot\}$)

where $O\{B|A\} \equiv \lambda w. ob(A)(B)$

— $O\{B|A\}$ can be read as “B is obligatory in the context A”

While DDL is powerful because of its support for a dyadic obligation operator, in many cases, I only need a monadic obligation operator. Below is some syntactic sugar for a monadic obligation operator.

abbreviation $ddltrue::t \ (\top)$

where $\top \equiv \lambda w. True$

abbreviation $ddlfalse::t \ (\perp)$

where $\perp \equiv \lambda w. False$

abbreviation $ddlob-normal::t \Rightarrow t \ (O\{-\})$

where $(O\{A\}) \equiv (O\{A|\top\})$

— Intuitively, the context *True* is the widest context possible because *True* holds at all worlds. Therefore, the monadic obligation operator requires that *A* is obligated at all worlds.

Finally, validity will be useful when discussing metalogical/ethical properties.

abbreviation $ddlvalid::t \Rightarrow bool \ (\models -)$

where $\models A \equiv \forall w. A\ w$

— A proposition is valid if it is true at all worlds.

Benemueller, Farjami, and Parent provide a proof of the completeness of the above embedding (Benzmüller et al., 2021). Isabelle allows us to check consistency immediately using Nitpick, a model checker (Blanchette and Nipkow, 2010). Nitpick can find satisfying models for a particular lemma using the `satisfy` option and it can find counterexamples using the `falsify` option, both of which I use heavily in this project.

lemma *True* **nitpick** [*satisfy,user-axioms,format=2*] **by simp**

— This an example of a typical Nitpick output. In this case, Nitpick successfully found a model satisfying these axioms so the system is consistent.

— Nitpick found a model for card i = 1:

Empty assignment

In the proof above, **by simp** indicates the use of the Simplification proof method, which

involves unfolding definitions and applying theorems directly. HOL has *True* as a theorem, which is why this theorem was so easy to prove.

3 Implementation Details

In this section, I present the details of my implementation of automated Kantian ethics, which requires formalizing the FUL in Dyadic Deontic Logic and then implementing this logic in Isabelle/HOL. The final Isabelle library contains enough logical background and an axiom representing the categorical imperative. Using Isabelle’s automated theorem proving abilities, my system can show that appropriately represented actions are, according to Kantian ethics, obligatory, permissible, or prohibited by proving or refuting sentences of the form “A is obligated to do B.” I also present a testing framework to evaluate how faithful my implementation is to philosophical literature. This testing framework shows that my system outperforms two other potential implementations of automated Kantian ethics.

3.1 Formalization and Implementation of the FUL

Formalizing the FUL requires defining and implementing enough logical background to represent the FUL as an axiom. Dyadic Deontic Logic can express obligation and prohibition, but it cannot represent more complex features of moral judgement like actions, subject, maxims, and ends. I augment DDL by adding representations of these concepts, drawn from philosophical literature.

3.1.1 Logical Background

Kantian ethics is fundamentally action-guiding; the categorical imperative is a moral rule that agents can use to decide between potential actions. Thus, before I begin to formalize a specific formulation of the categorical imperative, I must define the notions of subjects and actions. I add representations of subjects and actions so that my new logic can express sentences of the form, “x does action.”

typedec1 *s* — The new type *s* is the type for a “subject,” as in the subject of a sentence.

The **typedec1** keyword indicates that I am defining a new atomic type, which is not

composed of pre-existing types but is instead a new kind of object altogether. A type does not come with any properties out of the box. There is no difference between declaring the type “subject” or any other type, such as “color” or “mammal.” I add some properties of this type below by creating formulas and more complex types that use this type, but I do not provide anything near a definition of a subject. It suffices to declare a new type to represent a subject without specifying any of its complex properties, such as the idea that a subject must be rational or human. Instead of providing a complete definition of subject, which would require wading into murky philosophical debates about the nature of agency, I provide a “thin” definition that only includes the minimum necessary properties to apply the FUL. Throughout my project, I will use bare syntactic units like types and constants to create thin definitions of new ideas. This strategy lets me avoid messy philosophical controversies and makes my system’s judgements more trustworthy, because they rely on relatively little prior knowledge.

In this interpretation, the defining feature of a subject is that it can act. I represent that below by allowing subjects to substitute into sentences, a property that I will use to represent the idea that different people can perform the same actions.

type-synonym $os = (s \Rightarrow t)$

— To model the idea of a subject being substituted into an action, I define `type_synonym os` for an open sentence. An open sentence takes as input a subject and returns a complete or “closed” DDL formula by binding the free variable in the sentence to the input. For example, “runs” is an open sentence that can be instantiated with subject, “Sara” to create the DDL term “Sara runs,” which can be true or false at a world. An open sentence itself is not truth-apt, or the kind of thing that can be true or false at a world. When an action is substituted into an open sentence, the resulting term is truth apt. “Runs” is not the kind of thing that can be true or false, but “Sara runs” is a sentence that can be true or false.

3.1.2 Maxims

As established in Section 2.1.4, I formalize a version of the categorical imperative called the Formula of Universal, which reads “act only according to that maxim by which you can at the same time will that it should become a universal law” (Kant, 1785). In order to faithfully

formalize the FUL, I must make precise what it means to will a maxim and what kinds of maxims can become universal laws. I draw on reliable definitions of willing, maxims, and universalization from Kantian literature and represent them in DDL. Throughout this section, I will use one of Kant’s canonical maxims as an example.

Example 1 (False Promising). *The false promising example maxim reads, “When I am strapped for cash, I will falsely promise to repay a loan to get some easy cash.”*

The central unit of evaluation for Kantian ethics is a “maxim,” which Kant defines as “the subjective principle of willing,” or the principle that the agent understands themselves as acting on (Kant, 1785). Modern Kantians differ in their interpretations of this definition. I adopt O’Neill’s view, derived from Kant’s example maxims, that a maxim includes the act, the circumstances, and the agent’s purpose of acting or goal (O’Neill, 2013). Other potential views include Korsgaard’s view, which omits the circumstances, and Kitcher’s view, which adds a motivation (Korsgaard, 2005; Kitcher, 2004). I address the limitations of these approaches in Appendix A.1.

Definition 1 (Maxim). *A maxim is a circumstance, act, goal tuple (C, A, G) , read as “In circumstances C , act A for goal G .”*

I implement this definition in Isabelle by defining the `type_synonym` below for the type of a maxim.

type-synonym *maxim* = $(t * os * t)$

— A maxim is of type term, open sentence, term tuple, such as “(When I am strapped for cash, will falsely promise to repay a loan, to get some easy cash)”. The first term represents the circumstance, which can be true or false at a world. For example, the circumstance “when I am strapped for cash” is true at the real world when my bank account is empty. The second term represents the action, which is an open sentence because different agents can perform a particular action. For example, the action, “will falsely promise to repay a loan” is an open sentence that can be acted on by a subject. The third term represents the goal, which can again be true or false at a world. For example, the goal “to get some easy cash” is true at the real world if I have successfully received easy cash.

O’Neill argues that a maxim is an action-guiding rule and thus naturally includes an

act and the circumstances under which it should be performed (O'Neill, 2013, 37). She also includes a purpose, end, or goal in the maxim because human activity is guided by a rational will and is thus inherently purposive (Kant, 1785, 4:428). A rational will does not act randomly (else it would not be rational), but instead in the pursuit of ends which it deems valuable. The inclusion a maxim's end is essential for the version of the FUL that I will implement, explained in Section 3.1.3.

O'Neill's inclusion of circumstances is potentially controversial because it leaves open the question of what qualifies as a relevant circumstance for a particular maxim. This gives rise to "the tailoring objection," under which maxims are arbitrarily specified to pass the FUL (Kitcher, 2003, 217).¹⁶ For example, the maxim "When my name is Jane Doe and I am wearing a purple shirt and it is Tuesday morning, I will murder my boss so I can take their job," is universalizable but is clearly a false positive because we think that murder for professional gain is wrong. One solution to this problem is to argue that the circumstance "When my name is Jane Doe and I am wearing a purple shirt and it is Tuesday morning" is not morally relevant to the act and goal. This solution requires determining what qualifies as a relevant circumstance.

O'Neill seems to acknowledge the difficulty of determining relevant circumstances when she concedes that a maxim cannot include all of the infinitely many circumstances in which the agent may perform an action (O'Neill, 2013, 4:428). She argues that this is an artifact of the fact that maxims are rules of practical reason, the kind of reason that helps us decide what to do and how to do it (Bok, 1998). Like any practical rule, maxims require the exercise of practical judgement to determine in which circumstances they should be applied. This judgement, applied in both choosing when to exercise the maxim and in the formulation of the maxim itself, is what determines the morally relevant circumstances. The difficulty in determining relevant circumstances is a limitation of my system and requires that a human being formulate the maxim or that future work develop heuristics to classify circumstances as morally relevant. I discuss this challenge and potential solutions in greater detail in Section

¹⁶ Kitcher cites Wood (1999) as offering an example of a false positive due to this objection Kitcher (2004).

??.

With this robust representation of a maxim, I can now define willing. To will a maxim is to adopt it as a principle to live by, or to commit oneself to the maxim's action for the sake of maxim's end in the relevant circumstances. I formalize this idea in Definition 2.

Definition 2 (Willing). *For maxim $M = (C, A, G)$ and actor s ,*

$$\text{will } M \text{ } s \equiv \forall w (C \longrightarrow A(s)) w$$

At all worlds w , if the circumstances hold at that world, agent s performs act A .

If I will the example [False Promising](#) maxim, then whenever I need cash, I will falsely promise to repay a loan. I can represent this definition using the following Isabelle formula.

abbreviation *will* :: *maxim* \Rightarrow *s* \Rightarrow *t* (*W* - -)

where *will* $\equiv \lambda(c, a, g) s. (c \longrightarrow (a \text{ } s))$

— An agent s wills a maxim if in the circumstances, s performs the action, or s substituted into the open sentence a is true. This is an Isabelle **abbreviation**, which is syntactic sugar for an Isabelle formula. The type of this formula is *maxim* $\rightarrow s \rightarrow t$, so it takes as input a maxim and a subject and returns the term, “ s wills maxim.”

3.1.3 Practical Contradiction Interpretation of the FUL

In order to evaluate the moral status of a maxim, I must define what it means for a maxim to not be universalizable, or to fail the universalizability test. Kantians debate the correct interpretation of the Formula of Universal Law because Kant himself appears to interpret the criterion in different ways. I adopt Korsgaard's practical contradiction interpretation, broadly accepted as correct within the philosophical community ([Ebels-Duggan, 2012](#)).

Recall that the Formula of Universal Law is to “act only in accordance with that maxim through which you can at the same time will that it become a universal law” ([Kant, 1785](#)). To determine if a maxim can be willed as a universal law, we use the “universalizability test,” which requires imagining a world in which everyone has willed the maxim. If willing the maxim in such a world generates a contradiction, then the action is prohibited. For example,

the [False Promising](#) maxim will be prohibited if it is impossible to will the maxim in a world where everyone falsely promises to repay loans.

One interpretation of the FUL, the logical contradiction interpretation, prohibits maxims that are logically impossible when universalized. Under this view, falsely promising to repay a loan fails the universalizability test because, in the universalized world, everyone falsely promises to repay loans so lenders no longer believe promises to repay loans. The practice of loans would die out, so making a false promise to repay a loan would be impossible.

This view cannot correctly handle natural acts. Korsgaard appeals to [Dietrichson \(1964\)](#) to construct the example natural act of a mother killing her children that cry too much at night so that she can get some sleep. Universalizing this maxim does not generate a logical contradiction because killing is still possible in a world where everyone kills noisy children, but it is clearly wrong. Because killing is a natural act, it can never be logically impossible so the logical contradiction view cannot prohibit it.

As an alternative to the logical contradiction view, Korsgaard endorses the practical contradiction view, which prohibits maxims that are self-defeating, or ineffective, when universalized. By willing a maxim, you commit yourself to the maxim's goal, and thus cannot rationally will that this goal be undercut. This interpretation can prohibit natural acts like those of the sleep-deprived mother: in willing the end of sleeping, she is implicitly willing that she is alive. If all mothers kill all loud children, then she cannot be secure in the possession of her life, because her own mother would have killed her as an infant. Her willing this maxim thwarts the end that she sought to secure.

The practical contradiction interpretation offers a satisfying explanation of *why* certain maxims are immoral. These maxims involve parasitic behavior on social conditions that the agent seeks to benefit from. The false promiser wants to both abuse the system of promising and benefit from it, and is thus making an exception of themselves. The test formalizes the kinds of objections that the question "What if everyone did that?" seeks to draw out.

Under the practical contradiction interpretation, the FUL states, "If, when universalized, a maxim is not effective, then it is prohibited." This requires defining effectiveness and

universalization. If an agent wills an effective maxim, then the maxim’s goal is achieved, and if the agent does not will it, then the goal is not achieved.

Definition 3 (Effective Maxim). *For a maxim $M = (C, A, G)$ and actor s ,*

$$\text{effective } M s \equiv \forall w (\text{will } (C, A, G) s \iff G) w$$

I implement this in Isabelle below.

abbreviation *effective* :: *maxim* \Rightarrow *s* \Rightarrow *t* (*E* - -)

where *effective* $\equiv \lambda(c, a, g) s. ((\text{will } (c, a, g) s) \equiv g)$

— A maxim is effective for a subject if the goal is achieved if and only if the subject wills the maxim.

Once again, I use an abbreviation to conveniently refer to this Isabelle formula.

The former direction of the implication is intuitive: if the act results in the goal, it was effective in causing the goal. This represents sufficient causality. The latter direction represents necessary causality, or the idea that, counterfactually, if the maxim were not willed, then the goal would not be achieved (Lewis, 1973a).¹⁷ Note that nothing else changes about this counterfactual world—the circumstances are identical and we neither added additional theorems nor specified the model any further. This represents Lewis’s idea of “comparative similarity,” where a counterfactual is true if it holds at the most similar world (Lewis, 1973b). In our case, this is just the world where everything is the same except the maxim is not acted on. Combining these ideas, this definition of effective states that a maxim is effective if the maxim being acted on by a subject is the necessary and sufficient cause of the goal.

Next, I define what it means for a maxim to be universalized. Recall that the universalizability test requires imagining a world in which everyone wills a maxim. Therefore, a maxim is universalized when everyone wills the maxim.

Definition 4 (Universalized). *For a maxim M and agent s ,*

$$\text{universalized } M \equiv \forall w (\forall p \text{ will } M p)$$

¹⁷Thank you to Jeremy Zucker for helping me think about causality in this way.

I can once again represent this as an abbreviation in Isabelle.

abbreviation *universalized* :: *maxim* \Rightarrow *t* **where**

universalized $\equiv \lambda M. (\lambda w. (\forall p. (W M p) w))$

— The abbreviation *universalized* takes a maxim as input and returns a term which is true at a world if all people at that world will the maxim.

With the above definitions of effective and universalization, I can define what it means for a maxim to not be universalizable. This is the core of the FUL, which states that if a maxim is not universalizable, it is prohibited. Under the practical contradiction interpretation, a maxim is not universalizable if, when universalized, it is no longer effective.

Definition 5 (Not Universalizable). *For a maxim M and agent s ,*

$$\text{not_universalizable } M \ s \equiv [\text{universalized } M \longrightarrow \neg \text{effective } M \ s]$$

A maxim is not universalizable when, if everyone wills the maxim, then it is no longer effective.

I implement this definition in Isabelle using another abbreviation.

abbreviation *not-universalizable* :: *maxim* \Rightarrow *s* \Rightarrow *bool* **where**

not-universalizable $\equiv \lambda M \ s. \forall w. ((\text{universalized } M) \rightarrow (\neg (E M s))) w$

— Maxim M is not universalizable at world w when, “at world w , if M is universalized, then M is not effective.”

The FUL states that if a maxim is not universalizable, then it is prohibited. Before performing moral reasoning with my system, I must define obligation, permissibility, and prohibition. To judge an action, my system evaluates the moral status of the action “person s wills maxim M .” This action can be obligated, prohibited, or permissible. I will use the phrase “subject s willing maxim M is obligatory” interchangeably with “maxim M is obligatory for subject s .” I will use “maxim M is obligatory” to refer to M being obligatory for any arbitrary subject, which is equivalent to M being obligatory for a specific subject.¹⁸

¹⁸The full proof for this result is the Obligation Universalizes Across People Test in Section ??.

Definition 6 (Obligation). *Let maxim M be composed of the circumstances, act, goal tuple C, A, G and let s be an arbitrary agent.*

$$\text{obligated } M s \equiv O\{\text{will } (C, A, G) s \mid C\}$$

The action “ s wills M ” is the first argument passed to the dyadic obligation operator and is thus the action that is shown to be obligated or not. The second argument passed to the obligation operator represents the context in which the obligation holds and is thus naturally the maxim’s circumstances. This definition does not require any additional syntactic sugar, since it merely uses the dyadic obligation operator. Using this definition, I can define prohibition and permissibility.

Definition 7 (Prohibition and Permissibility). *Let maxim M be composed of the circumstances, act, goal tuple C, A, G and let s be an arbitrary agent.¹⁹*

$$\text{prohibited } M s \approx \text{obligated } \neg M \equiv O\{\neg \text{will } (C, A, G) s \mid C\}$$

$$\text{permissible } M s \equiv \neg \text{prohibited } M s \equiv \neg O\{\neg \text{will } (C, A, G) s \mid C\}$$

abbreviation *prohibited::maxim $\Rightarrow s \Rightarrow t$ where*

$$\text{prohibited} \equiv \lambda(c, a, g) s. O\{\neg (\text{will } (c, a, g) s) \mid c\}$$

— A maxim is prohibited for a subject s if its negation is obligated for s . It is morally wrong for an agent to will a prohibited maxim.

abbreviation *permissible::maxim $\Rightarrow s \Rightarrow t$*

$$\text{where } \text{permissible} \equiv \lambda M s. \neg (\text{prohibited } M s)$$

— A maxim is permissible for a subject s if it is not prohibited for s . It is morally acceptable for an agent to will or not will a permissible maxim.

One additional piece of logical background necessary before I implement the FUL is

¹⁹Technically, a maxim is not a boolean type, so the term $\neg M$ is not type correct. The expression *obligated* $\neg M$ merely provides intuition for the meaning of prohibition, but the exact definition is given by $O\{\neg \text{will } (C, A, G) s \mid C\}$.

the notion of contradictory obligations. Many deontic logics, including DDL, allow contradictory obligations. As I will explain in Section [Tests](#), Kantian ethics never prescribes contradictory obligations, so I will add this as an axiom.

abbreviation *non-contradictory where*

non-contradictory $A B c w \equiv ((O\{A|c\} \wedge O\{B|c\}) w) \longrightarrow \neg((A \wedge (B \wedge c)) w \longrightarrow False)$

— Terms A and B are non contradictory in circumstances c if, when A and B are obligated in circumstances c , the conjunction of A , B , and c , does not imply *False*.

axiomatization where *no-contradictions*: $\forall A::t. \forall B::t. \forall c::t. \forall w::i. \text{non-contradictory } A B c w$

— This axiom formalizes the idea that, for any terms A , B , and circumstances c , A , and B must be non-contradictory in circumstances c at all worlds. Intuitively, this axiom requires that obligations do not conflict.

3.1.4 Formalizing the FUL

With this logical background, I can begin to implement the Formula of Universal Law, which, as defined by the practical contradiction interpretation, states that a maxim is prohibited if it is ineffective when universalized. A first, ultimately unsuccessful, attempt to formalize the FUL simply translates this into Isabelle’s syntax using the abbreviations above. While this attempt will not be consistent, I will use Isabelle’s automatic proof search abilities to determine how to modify this formula to be consistent, revealing a key philosophical insight about maxims in the process. This section presents my final formalization of the FUL and the philosophical insight produced while creating it, which, as I argue in Section [Computational Ethics](#)???, demonstrates the power of computational tools in aiding philosophical progress.

abbreviation *FUL0::bool where* $FUL0 \equiv \forall c a g s. \text{not-universalizable } (c, a, g) s \longrightarrow \models ((\text{prohibited } (c, a, g) s))$

— This representation of the Formula of Universal Law reads, “For all circumstances, goals, acts, and subjects, if the maxim of the subject performing the act for the goal in the circumstances is not universalizable (as defined above), then, at all worlds, in those circumstances, the subject is prohibited from (obligated not to) willing the maxim.”

I can immediately determine if this version of the FUL is consistent by checking if $FUL0$

implies *False*.

lemma *FULo* \longrightarrow *False* **using** *O-diamond*

using *case-prod-conv no-contradictions old.prod.case old.prod.case* **by** *fastforce*

Isabelle’s proof-finding tool, Sledgehammer, shows that FUL0 is not consistent by showing that it implies a contradiction using axiom *O_diamond*²⁰ (Paulson and Blanchette, 2015). This axiom roughly states that an obligation can’t contradict its context. Knowing that FUL0 contradicts this particular maxim offers insight into what the problem is. If the goal or action or a maxim are equivalent to its circumstance, then prohibiting it is contradictory. If the maxim has already been acted on or the goal has already been achieved, then the agent cannot undo their action or the achievement of the goal. Moreover, in many models, it is possible to construct a maxim where the circumstances, goal, and act (once a subject acts on it) are all that same term and to show that this maxim is prohibited, resulting in a contradiction.

This motivates the exclusion of what I call “badly formed maxims,” which are those maxims such that the goal has already been achieved or the act has already been acted on. Under my formalization, such maxims are not well-formed.

Definition 8 (Well-Formed Maxim). *A maxim is well-formed if the circumstances do not contain the act and goal. For a maxim (C, A, G) , and subject s ,*

$$\text{well_formed}(C, A, G) s \equiv \forall w. (\neg(C \longrightarrow G) \wedge \neg(C \longrightarrow A s)) w$$

For example, the maxim “When I eat breakfast, I will eat breakfast to eat breakfast” is not well-formed because the circumstance “when I eat breakfast” contains the act and goal. Well-formedness is not discussed in the literature, but I find that if I require that the FUL only holds for well-formed maxim, it is consistent.

abbreviation *well-formed::maxim \Rightarrow s \Rightarrow i \Rightarrow bool* **where**

$$\text{well_formed} \equiv \lambda(c, a, g). \lambda s. \lambda w. (\neg(c \longrightarrow g) w) \wedge (\neg(c \longrightarrow a s) w)$$

— This abbreviation formalizes the well-formedness of a maxim for a subject. The goal cannot be already achieved in the circumstances and the subject cannot have already performed the act.

²⁰The full axiom reads $\models \lambda w. ob ?B ?A \longrightarrow \neg \models \neg ?B \wedge ?A$.

If I modify FUL0 to only hold for well-formed maxims, it becomes consistent.

abbreviation FUL where

$$FUL \equiv \forall M::\text{maxim}. \forall s::s. (\forall w. \text{well-formed } M s w) \longrightarrow (not\text{-universalizable } M s \longrightarrow \models (\text{prohibited } M s))$$

— This formalization states that if a maxim is well-formed, then if it is not universalizable, it is prohibited.

lemma FUL

nitpick[*user-axioms*, *falsify=false*] **oops**

— Nitpick is Isabelle’s countermodel checker, and I can use it to quickly check that an axiom is consistent (Blanchette and Nipkow, 2010). If Nitpick can find a model in which the axioms of DDL hold and the FUL is true, then it is consistent.

Nitpick found a model for card $i = 1$ and card $s = 2$:

Empty assignment

My above investigation of FUL0 shows that if the FUL holds for badly formed maxims, then it is inconsistent. This is not only a logical property of my system, but it also has philosophical significance that coheres with Korsgaard’s and O’Neill’s interpretations of a maxim as a practical guide to action (Korsgaard, 2005; O’Neill, 2013). A maxim is a practical principle that guides how we behave in everyday life. A principle of the form “When you are eating breakfast, eat breakfast in order to eat breakfast,” is not practically relevant. No agent would ever need to act on such a principle. Morality helps agents decide whether to act on a potential principle of action, but no agent would need to ask “When I am eating breakfast, should I eat breakfast in order to eat breakfast?” It is neither contradictory nor prohibited, and it is the wrong kind of principle to be evaluating. It is not a well-formed maxim, so the categorical imperative cannot apply to it.

The fact that Isabelle revealed a philosophical insight about which kinds of maxims are well-formed is an example of the power of computational tools to aid philosophical progress. Nitpick and Sledgehammer helped me confirm that certain kinds of circumstance, act, goal tuples are too badly formed for the categorical imperative to logically apply to them. The realization of this subtle problem would have been incredibly difficult without computational

tools, and serves as evidence of the power of computational ethics. I discuss the philosophical properties and implications of well-formed maxims and the power of computational ethics further in Section Computational Ethics.

Now that I have a consistent version of the *FUL*, I can complete my implementation by adding it as an axiom.

axiomatization where *FUL:FUL*

This concludes my implementation of the Formula of Universal Law in Isabelle/HOL. My implementation consists of necessary logical background, first formalized in DDL and then implemented in Isabelle. The code snippets in this chapter are a subset of the over 100 lines of Isabelle/HOL code necessary to complete this implementation. In Section 3.2 and Chapter Applications, I demonstrate how this implementation can be tested and used to make moral judgements.

3.2 Tests

In addition to an implementation of automated Kantian ethics, I also contribute a testing framework to evaluate how well my implementation coheres with philosophical literature. This testing architecture makes the notion of “philosophical faithfulness” precise. Each test consists of a sentence in my logic and an expected outcome, where the possible outcomes are proving or refuting the sentence. For example, one such sentence is that obligations cannot contradict each other. To run the tests, I attempt to prove or refute each test sentence in my logic. Because these tests are derived from moral intuition and philosophical literature, they evaluate how well my system reflects philosophical literature. Running the tests on my implementation consisted of approximately 400 lines of Isabelle code.

The test framework can be expanded by adding more test sentences and can guide implementations of other parts of Kantian ethics or other ethical theories. As I was implementing my formalization, I checked it against the testing framework, performing test-driven development for automated ethics.

Test	Naive	Kroy	Custom
FUL Stronger than DDL	×	✓	✓
Obligation Universalizes Across People	×	✓	✓
Obligations Never Contradict	×	×	✓
Distributive Property for Obligations Holds	×	×	✓
Prohibits Actions That Are Impossible to Universalize	×	×	✓
Robust Representation of Maxims	×	×	✓
Can Prohibit Conventional Acts	×	×	✓
Can Prohibit Natural Acts	×	×	✓

Figure 1: Table showing which tests each implementation passes. The naive interpretation is raw DDL, Kroy is based on Moshe Kroy’s formalization of the FUL, and the custom formalization is my novel implementation.

I use my testing framework to show that my formalization and implementation of Kantian ethics outperforms two other potential implementations. First, I consider raw DDL, which serves as a control group because it simply contains the base logic on top of which I build other implementations. DDL can express obligation, but has no knowledge of any specific moral rules (like the categorical imperative) and is thus a control group. Second, I consider Moshe Kroy’s 1976 formalization of the FUL and the second formulation of the categorical imperative (Kroy, 1976). His formalization is based on Hintikka’s deontic logic, which is a different, less expressive logic than DDL (Hintikka, 1962). He presents a logical representation of the FUL that has not yet been implemented using an automated theorem prover, so I implement it in Isabelle.²¹ I find that my implementation outperforms both other implementations. Full test results are summarized in Figure 1. Below, I briefly explain some notable tests.

FUL Stronger than DDL The FUL should not hold in raw DDL, which I use a control group. If the FUL holds in the base logic, then adding it as an axiom doesn’t make the logic any stronger, which is troubling because the base logic does not come equipped with the categorical imperative. DDL defines basic properties of obligation, such as ought implies can, but contains no axioms that represent the Formula of Universal Law. Therefore, if a formaliza-

²¹I present the complete implementation in Appendix A.2.

tion of the FUL holds in the base logic, then it is too weak to actually represent the FUL. Both Kroy’s formalization and my implementation do not hold in the base logic, and thus represent progress over the control group. To test this property, I used Nitpick to find a countermodel in which my version of the FUL does not hold. I performed this test before adding the FUL as an axiom, since after adding it no countermodel will be possible.

Obligation Universalizes Across People The obligations prescribed by the Formula of Universal Law should generalize across people. In other words, if a maxim is obligated for one person, then it is obligated for all other people because maxims are not person-specific. Velleman argues that, because reason is accessible to everyone identically, obligations apply to all people equally (Velleman, 2005, 25). When Kant describes the categorical imperative as the objective principle of the will, he is referring to the fact that, as opposed to a subjective principle, the categorical imperative applies to all rational agents equally (Kant, 1785, 16). At its core, the FUL best handles, “the temptation to make oneself an exception: selfishness, meanness, advantage-taking, and disregard for the rights of others” (Korsgaard, 1985, 30). Kroy makes this property the center of his formalization, which essentially says that if an act is permissible for someone, it is permissible for everyone.²² Kroy’s formalization and my formalization satisfy this property, but raw DDL does not. Below I run this test for my formalization.

lemma *wrong-if-wrong-for-someone*:

shows $\forall w. \forall c::t. \forall g::t. \forall a. \exists s::s. O\{\neg (W(c, a, g) s) \mid c\} w \longrightarrow (\forall p. O\{\neg (W(c, a, g) p) \mid c\} w)$

by *blast*

— I represent my tests as lemmas that I expect Isabelle to either prove or refute. The statement following the keyword **shows** is the sentence of the lemma, and the proof follows the **by** keyword

— This lemma shows that if a maxim (c, a, g) is wrong for subject s at a world, then it is wrong for all people at that world. Isabelle automatically completed this proof using the **blast** method, which implements a generic tableau prover, a proof method that operates on lists of formulae using rules for conjunction, disjunction, universal quantification, and existential quantification (Paulson, 1999).

²²Formally, $P\{A(s)\} \longrightarrow \forall p. P\{A(p)\}$.

lemma *right-if-right-for-someone*:

shows $\forall w. \forall c::t. \forall g::t. \exists s::s. O\{W(c, M, g) s \mid c\} w \longrightarrow (\forall p. O\{W(c, M, g) p \mid c\} w)$

by *blast*

— This lemma shows that if a maxim (c, a, g) is right for subject s at a world, then it is right for all people at that world. The proof similarly proceeds using *blast*.

Obligations Never Contradict Contradictory obligations make obeying the prescriptions of an ethical theory impossible. Kant subscribes to the general, popular view that morality is supposed to guide action, so ought implies can.²³ Kohl reconstructs Kant’s argument for this principle as follows: if the will cannot comply with the moral law, then the moral law has no prescriptive authority for the will (Kohl, 2015, 703-4). This defeats the purpose of Kant’s theory, which is to develop an unconditional, categorical imperative for rational agents. Ought implies can requires that obligations never contradict each other, because an agent can’t perform contradictory actions. Therefore, any ethical theory that respects ought implies can, and Kantian ethics in particular, must not result in conflicting obligations. Kant only briefly discusses contradictory obligations in *Metaphysics of Morals*, where he argues that conflicting moral obligations are impossible under his theory (Kant, 2017, V224). Particularly, the categorical imperative generates “strict negative laws of omission,” which cannot conflict by definition (Timmermann, 2013, 45).²⁴ Both raw DDL and Kroy’s formalization allow contradictory obligations, but I explicitly add an axiom to my formalization that prohibits contradictory obligations.

lemma *conflicting-obligations*:

shows $\neg (O\{W(c, a, g) s \mid c\} \wedge O\{\neg(W(c, a, g) s) \mid c\}) w$

using *no-contradictions* **by** *blast*

— This test passes immediately by the new axiom that prohibits contradictory obligations.

²³Kohl points out that this principle is referred to as Kant’s dictum or Kant’s law in the literature (Kohl, 2015, footnote 1).

²⁴The kinds of obligations generated by the FUL are called “perfect duties” which arise from “contradictions in conception,” or maxims that we cannot even conceive of universalizing. These duties are always negative and thus never conflict. Kant also presents “imperfect duties,” generated from “contradictions in will,” or maxims that we can conceive of universalizing but would never want to. These duties tend to be broader, such as “improve oneself” or “help others,” and are secondary to perfect duties. My project only analyzes perfect duties, as these are always stronger than imperfect duties.

lemma *implied-contradiction*:

assumes $((W(c1, a1, g1) s) \wedge (W(c2, a2, g2) s)) \rightarrow \perp) w$

shows $\neg (O\{W(c1, a1, g1) s|c\} \wedge O\{W(c2, a2, g2) s|c\}) w$

using *assms no-contradictions* **by** *blast*

— This stronger property states that the combination of obligatory maxims can't imply a contradiction and should hold for the same reasons that contradictory obligations aren't allowed. The added axiom also makes this test pass.

Contradictory obligations are closely related to two other properties. First is the idea that obligation implies permissibility, or that obligation is a stronger property than permissibility. If there are no contradictory obligations, then this property holds because actions are either permissible or prohibited and obligation contradicts prohibition. In a system with contradictory obligations, this property fails because there is some A that is obligated but also prohibited and therefore not permissible. Formalizing this property below shows that this follows from contradictory obligations.

lemma $\models ((O\{A\} \wedge O\{\neg A\}) \equiv (\neg (O\{A\} \rightarrow \neg O\{\neg A\})))$

by *simp*

— *simp* is the simplification tactic, which unfolds definitions to complete a proof. The left-hand side states that A is both obligated and prohibited, and is equivalent to the right-hand side, which states that A is obligated but not permissible.

Distributive Property for Obligations Holds Another property related to contradictory obligations is the distributive property for the obligation operator.²⁵ This is another property that we expect to hold. The rough English translation of $O\{A \wedge B\}$ is “you are obligated to do both A and B ”. The rough English translation of $O\{A\} \wedge O\{B\}$ is “you are obligated to do A and you are obligated to do B .” We think those English sentences mean the same thing, and they should mean the same thing in logic as well. Moreover, if that (rather intuitive) property holds, then contradictory obligations are impossible, as shown in the below proof.

lemma *distributive-implies-no-contradictions*:

assumes $\forall A B. \models ((O\{A\} \wedge O\{B\}) \equiv O\{A \wedge B\})$

²⁵Formally, $O\{A\} \wedge O\{B\} \longleftrightarrow O\{A \wedge B\}$.

shows $\forall A. \models (\neg(O\{A\} \wedge O\{\neg A\}))$

using *O-diamond assms* **by** *blast*

— The **assumes** keyword indicates assumptions used when proving a lemma. I use it here to represent metalogical implication. With the assumption, the lemma above reads, “If the distributive property holds in this logic, then obligations cannot contradict.”

Again, this test fails for raw DDL and for Kroy’s formalization, but passes for my interpretation because I require that obligations don’t contradict as an axiom.

lemma *distribution*:

assumes $\models (O\{A\} \wedge O\{B\})$

shows $\models O\{A \wedge B\}$

using *assms no-contradictions* **by** *fastforce*

— The proof proceeds almost immediately using the new axiom.

Prohibits Actions That Are Impossible to Universalize Recall that the logical contradiction interpretation of the Formula of Universal Law prohibits lying because in a world where everyone simultaneously lies, lying is impossible. In other words, not everyone can simultaneously lie because the institution of lying and believing would break down. In Section 3.1.3, I recreated Korsgaard’s argument for why the logical contradiction interpretation is weaker than what the Formula of Universal Law should actually require. Therefore, any implementation of the FUL should be able to show that the actions prohibited by the logical contradiction interpretation are prohibited, because the set of actions prohibited by the practical contradiction interpretation is a superset of these. The FUL should show that actions that cannot possibly be universalized are prohibited, because those acts cannot be willed in a world where they are universalized. This property fails to hold in both raw DDL and Kroy’s formalization, but holds for my formalization. Showing that this property holds for my formalization required significant logical background and the full code is presented in Appendix A.3.

Robust Representation of Maxims Kant does not evaluate the correctness of acts, but rather of maxims. Therefore, any faithful formalization of the categorical imperative must

evaluate maxims, not acts. This requires representing a maxim and making it the input to the obligation operator, which only my implementation does. Because my implementation includes the notion of a maxim, it is able to perform sophisticated reasoning as demonstrated in Sections [Lies and Jokes](#) and [Lying to a Liar](#). Staying faithful to the philosophical literature enables my system to make more reliable judgements.

Can Prohibit Conventional and Natural Acts When arguing for the practical contradiction interpretation, Korsgaard makes a distinction between conventional and natural acts ([Korsgaard, 1985](#)). A conventional act like promising relies on a convention, like the convention that a promise is a commitment, whereas a natural act is possible simply because of the laws of the natural world. It is easier to show the wrongness of conventional acts because there are worlds in which these acts are impossible; namely, worlds in which the convention does not exist. For example, the common argument against falsely promising is that if everyone were to falsely promise, the convention of promising would fall apart because people wouldn't believe each other, so falsely promising is prohibited. It is more difficult to show the wrongness of a natural act, like murder or violence. These acts can never be logically impossible; even if everyone murders or acts violently, murder and violence will still be possible, so it is difficult to show that they violate the FUL.

Both raw DDL and Kroy's interpretation fail to show the wrongness of conventional or natural acts. My system shows the wrongness of both natural and conventional acts because it is faithful to Korsgaard's practical contradiction interpretation of the FUL, which is the canonical interpretation of the FUL ([Ebels-Duggan, 2012](#); [Korsgaard, 1985](#)). I run this test in Chapter Applications, where I use my system to reason about two ethical dilemmas, one which involves conventional acts and the other which involves natural acts. I present an additional example demonstrating that my implementation passes this test in Appendix [A.3](#).

4 Applications

In this chapter, I demonstrate that my system can produce correct judgements on challenging moral dilemmas that naive ethical reasoning cannot satisfactorily handle. My system can reproduce complex ethical judgements because it is faithful to philosophical literature. These dilemmas serve as examples of how my system could be used in practice and demonstrate my system's ability to formalize longer, more complicated ethical arguments than those presented in Chapter 3. Many of the tests in Section 3.2 perform metaethical reasoning, which analyzes properties of morality itself and involves questions about the nature of ethical truth. In this chapter, I perform "applied ethical reasoning," which is the use of ethics to resolve dilemmas and make judgements about what an agent should or should not do. This is the kind of reasoning necessary for an AI agent using my system to navigate the real world.

One challenge of applied ethical reasoning is that it requires more factual background than metaethical reasoning. Because metaethics is about ethics itself, and not about the dilemmas that ethics is supposed to help us resolve, this kind of reasoning requires very little knowledge about the world. Applied ethical reasoning, on the other hand, focuses on a particular ethical dilemma and thus requires enough factual background, or common sense, to understand the dilemma and options at hand. For example, an applied ethicist evaluating the permissibility of lying needs a robust definition of the term lying and likely some understanding of communication and truth telling. Kantians specifically describe this common sense as "postulates of rationality" that are nontrivial and nonnormative, but still part of the process of practical reasoning itself (Silber, 1974).

In this chapter, I attempt to tackle this challenge by endowing my system with this kind of common sense in the specific case of lying. In order for my system to perform applied ethical reasoning, it will need common sense facts and definitions, like those I present in this chapter. My system contributes the ability to reason using the Formula of Universal Law, but this reasoning must be applied to objects that are defined using common sense. In this chapter, I describe such common sense for the specific example of lying.

Because these common sense facts can determine my system’s judgements, they are part of the trusted code base for my system, or the logic and code that a user must trust in order to trust my system. Changing these common sense facts will change the judgements that my system makes. For example, if we define truth telling as an act that is self-contradictory (perhaps by defining it as $p \wedge \neg p$), then my system will output that truth telling is prohibited. Malicious common sense facts and definitions will result in bad ethical judgements. The challenge of endowing automated ethical reasoners with common sense reasoning is not unique to my system, and virtually all prior attempts in machine ethics face similar challenges.²⁶ Many prior attempts sidestep this question, whereas I contribute an prototype implementation of one kind of common sense reasoning.

This chapter will provide examples of the kinds of common sense facts required to get my system off the ground. I use a lean and uncontroversial common sense database to achieve robust and powerful results. This serves as evidence for the ease of automating Kantian ethics, an example of the additional work required to use my system in practice, and a demonstration of my system’s power and flexibility. These examples demonstrate that, armed with nuanced common sense facts, my system can make sophisticated judgements faithful to philosophical literature.

4.1 Lies and Jokes

The moral status of lying is hotly debated in the Kantian literature. I focus on two dilemmas presented in Korsgaard’s “Right to Lie,” which examines how strict Kant’s prohibition on lying is (Korsgaard, 1986). She begins with the case of lying and joking. Many of Kant’s critics argue that his prohibition on lies includes lies told in the context of a joke, which should be permissible. Korsgaard responds by arguing that there is a crucial difference between lying and joking: lies involve deception, but jokes do not. The purpose of a joke is amusement, which does not rely on the listener believing the story told. Given appropriate definitions of lies and jokes, my system shows that jokes are permissible but lies are not, demonstrating its power

²⁶See Section Related Work for a survey of the common sense required in prior work.

and flexibility. This section demonstrates how my system can be used in practice; it needs to be given some baseline common sense facts, but with those facts, it can make sophisticated judgements. Moreover, because my system is faithful to definitions found in philosophical literature, it can perform nuanced reasoning, demonstrating the value of faithful automated ethics.

First, Korsgaard argues that the categorical imperative prohibits lies because they involve deception. When universalized, lies will no longer be believed, so lying could never be an effective way of achieving any goal when universalized. Korsgaard points out that “we believe what is said to us in a given context because most of the time people in that context say what they really think” (Korsgaard, 1986, 4). In order to formalize this argument, I first need to define the relevant terms and assumptions, which include lying and Korsgaard’s argument about the basis of trust.

I define lying and trust in terms of belief. As in Section 3, I choose thin, or minimal, definitions of terms like belief to reduce the potential for controversy in my system’s factual background.

consts *believe::s⇒t⇒t* (- *believes* -)

— **believe** is a constant that maps a subject and a term to another DDL term. For example, subject “Sara” might believe the term “the sky is blue” to create the sentence “Sara believes that the sky is blue,” which can be true or false at a world.

Logicians and epistemologists develop and debate complex logics of belief and knowledge (Baltag and Renne, 2016). I avoid this complexity by defining the concept of belief simply as a constant that maps a subject, term pair to a term. For the examples in this section, this choice suffices. I encode some minimal properties of belief below, but avoid any full definition of the term.

Belief is useful to construct the idea of “knowingly uttering a falsehood,” a core component of both lying and joking.

consts *utter::s⇒t⇒t*

— **utter** also maps a subject and term to another DDL term. For example, the sentence “Sara utters,

‘I am hungry’ ” is a DDL term that can be true or false at a world.

abbreviation *utter-falsehood*:: $s \Rightarrow t \Rightarrow t$ **where**

$$\text{utter-falsehood } s \ t \equiv (\text{utter } s \ t) \wedge (\neg t)$$

— To utter a falsehood is to utter a statement that is false, or to utter t when $\neg t$.

abbreviation *knowingly-utter-falsehood*:: $s \Rightarrow t \Rightarrow t$ **where**

$$\text{knowingly-utter-falsehood } s \ t \equiv (\text{utter-falsehood } s \ t) \wedge (\neg (\text{believe } s \ t))$$

— Sometimes we unknowingly utter falsehoods. For example, if I believe that the Earth is flat, then when I utter, “the Earth is flat,” I am unknowingly uttering a falsehood. This motivates defining the idea of knowingly uttering a falsehood, which requires both uttering a falsehood and not believing your utterance. If I utter “the Earth is flat,” even though I know that the Earth is round, I am knowingly uttering a falsehood.

The above abbreviations are the core of my formalization of Korsgaard’s definitions of lying and joking. They are also relatively uncontroversial and encode little moral or normative content. They say nothing about the moral status of uttering a falsehood, the agent’s intention when making the utterance, or the conversational norms guiding the utterance. The complexities of a complete definition of lying or belief are unnecessary for Kantian ethics, and therefore for my system, to make moral judgements.

Using the above definitions, I define lying. I characterize a maxim as involving a lie if the act requires knowingly uttering a falsehood and the end requires that some person p believe the false statement t .

abbreviation *lie*:: $\text{maxim} \Rightarrow \text{bool}$ **where**

$$\text{lie} \equiv \lambda (c, a, g). \exists t. (a \longrightarrow (\lambda s. \text{knowingly-utter-falsehood } s \ t)) \wedge (\exists p. \forall w. (g \rightarrow (\text{believe } p \ t)) \ w)$$

— The abbreviation above maps a maxim to a boolean value that indicates if it is a lie.

To avoid unintentional wrongdoing, I focus on “knowing lies,” in which the speaker is aware that they are lying. It is uncontroversial that, in order for an act to be a knowing lie, the speaker must utter a false statement that they do not believe. This also makes it easier to make moral judgements about the speaker’s action, since they were, at the very least, aware of their lie. The second half of this definition requires that the goal of the lie is deception. This is inspired by Korsgaard’s interpretation of a lie. She understands a lie as a kind of falsehood that

is usually effective *because* it deceives (Korsgaard, 1986, 4). In my formalization, this means that the purpose or goal of the maxim must involve deceiving someone, or, in other words, that someone believe what the speaker knows to be a falsehood.

With the above logical background, I automate Korsgaard’s argument that maxims that involve lying are prohibited. First, I define the subject and maxim at hand.

consts *me::s*

— I am trying to reason about *my* obligations so I will define myself as a specific subject. Again, this is a minimal definition that does not include any facts about me, such as the fact that my name is Lavanya or that I have brown hair.

consts *m::maxim*

— I also define a maxim *m*. My goal is to show that if *m* is a maxim about lying, then *m* is prohibited.

consts *c::t a::os g::t*

— *m* will be composed of the circumstances, act, and goal above.

In the following lemma, I use my system to show that lying is prohibited. The assumptions of this lemma represent the common sense necessary to reach this conclusion. This common sense background is a direct formalization of the premises of Korsgaard’s argument. Using these relatively minimal premises about individual behavior, my system derives a prohibition against lying.

lemma *lying-prohibited:*

assumes $m \equiv (c::t, a::os, g::t)$

assumes $\forall w. \forall s. \text{well-formed } m \text{ s } w$

— Initial technical set-up: *m* is a well-formed maxim composed of some circumstances, act, and goal.

assumes *lie m*

— *m* is a maxim about lying as defined above. Precisely, it is a maxim in which the action requires knowingly uttering a falsehood and the goal requires that someone believe this falsehood.

assumes $\forall t w. ((\forall p. \text{utter-falsehood } p \text{ t } w) \longrightarrow (\forall p. \neg (\text{believe } p \text{ t } w)))$

— Assumption that if everyone utters false statement *t*, then no one will believe *t*. This assumption is Korsgaard’s core piece of “common sense” about lying (Korsgaard, 1986, 5). This simple assumption encodes the common sense knowledge that human communication involves an implicit trust, and that when this trust erodes, the convention of communication begins to break down and people no longer

believe each other. Call this the “convention of trust” fact. In the rest of this section, I will test versions of this assumption, effectively encoding different common sense understandings of lying.

assumes $\forall w. c\ w$

— Restrict our focus to worlds in which the circumstances hold. A technical detail.

shows $\models (prohibited\ m\ me)$

proof —

have $(\forall p\ w. (W\ m\ p)\ w) \longrightarrow (\models (c \rightarrow (\neg g)))$

by $(smt\ assms(1)\ assms(2)\ assms(5)\ case\text{-}prod\text{-}beta\ fst\text{-}conv\ old.\textit{prod.exhaust}\ snd\text{-}conv)$

— This proof requires some manual work. After I divide the proof into the intermediate steps shown here, Isabelle is able to do the rest. This step says that if m is universalized, then the circumstances won’t lead to the goal, which is close to the idea of the maxim not being universalizable.

have $not\text{-}universalizable\ m\ me$

by $(metis\ (mono\text{-}tags,\ lifting)\ assms(1)\ assms(2)\ case\text{-}prod\text{-}beta\ fst\text{-}conv\ snd\text{-}conv)$

thus $?thesis$

using $FUL\ assms(2)\ \mathbf{by}\ blast$

— $?thesis$ is Isabelle’s syntax for the goal of the lemma. In this case, $?thesis$ is equivalent to $\models prohibited\ m\ me$.

qed

The lemma above demonstrates that my system finds that lying is prohibited with a thin definition of lying and relatively uncontroversial facts about the world. The logical background needed is the fact that lying requires knowingly uttering a falsehood with the goal that someone believe the falsehood, a definition of lying that is relatively well-accepted. The factual background needed consists of the fact that if everyone lies in a given context, then people will stop believing each other in that context. This is a slightly heavier assumption, but it is still so uncontroversial that Korsgaard doesn’t bother to justify it in her argument against lying (Korsgaard, 1986).

Now that I have formalized Korsgaard’s argument for why lying is prohibited, I will implement her argument for why jokes are permissible. Specifically, she defines a joke as a story that is false and argues that joking is permissible because “the universal practice of lying in the context of jokes does not interfere with the purpose of jokes, which is to amuse and does

not depend on deception” (Korsgaard, 1986, 4). I use the fact that a joke does not depend on deception as the defining feature of a joke.

abbreviation $joke::maxim \Rightarrow bool$ **where**

$joke \equiv \lambda (c, a, g). \exists t. (a \longrightarrow (\lambda s. knowingly_utter_falsehood\ s\ t)) \wedge \neg (\exists p. \forall w. (g \rightarrow (believe\ p\ t))\ w))$

— This abbreviation states that a maxim is a joke if the action involves knowingly uttering a falsehood but the goal does *not* require that someone believe the falsehood told.

This definition of a joke defines a joke as a falsehood uttered for some purpose that doesn’t require deception, where deception involves someone believing the uttered falsehood. This definition doesn’t require any conception of humor, but merely distinguishes jokes from lies.

Korsgaard argues that her above argument for a prohibition against lying also implies that joking is permissible, because its purpose is not to deceive, but something else entirely. This means that, even armed with the same convention of trust assumption as above, joking should be permissible. The lemma below shows exactly that.

lemma *joking-not-prohibited*:

assumes $m \equiv (c::t, a::os, g::t)$

assumes $\forall w. \forall s. well_formed\ m\ s\ w$

— Initial set-up: m is a well-formed maxim composed of some circumstances, act, and goal.

assumes $joke\ m$

— m is a maxim about joking. Precisely, it is a maxim in which the action is to knowingly utter a falsehood and the goal does not require that someone believe this falsehood.

assumes $\forall t\ w. ((\forall p. utter_falsehood\ p\ t\ w) \longrightarrow (\forall p. \neg (believe\ p\ t)\ w))$

— The same convention of trust assumption as in the above example.

assumes $\forall w. c\ w$

— Restrict our focus to worlds in which the circumstances hold. A technical detail.

shows $\models (permissible\ m\ me)$

by $(smt\ assms(1)\ assms(2)\ assms(3)\ case_prod_conv)$

— Isabelle is able to show that maxims about joking are permissible. It also lists the facts used in its proof, which offer insight into how it arrived at its judgement. Specifically, it uses assumptions 1, 2,

and 3, which are the logical background and definition of the joking maxim. Notably, it does not use the convention of trust assumption. This demonstrates that even the convention of trust assumption is not strong enough to prohibit joking, which is exactly the desired result.

My system shows that lies are prohibited and jokes are permissible with thin conceptions of amusement and deception. This shows that it isolates a necessary and sufficient property of this class of maxims that fail the universalizability test. My definitions of a lie and joke only differ in whether or not their goal requires that someone believe the falsehood in question, so this is a necessary and sufficient condition for a maxim about knowingly uttering falsehoods to be prohibited. This logical fact derived by my system tracks a fact implicit in Korsgaard's argument and in most Kantian accounts of lying: the wrongness of lying is derived from the requirement that someone believe the falsehood. The logical reality that this property is necessary and sufficient to generate a prohibition reflects a deep philosophical explanation of *why* certain maxims about uttering falsehood fail the universalizability test. Universalizing uttering a falsehood makes belief in that falsehood impossible, so any maxims with goals that require believing in the falsehood will be prohibited.

This account not only describes the kind of maxims that fail or pass the universalizability test, but it also provides a guide to constructing permissible maxims about uttering falsehoods. As an example, consider the idea of throwing a surprise birthday party. At first glance, the maxim of action is something like, "When it is my friend's birthday, I will secretly plan a party so that I can surprise them." The goal "so that I can surprise them" clearly requires that your friend believe the falsehood that you are not planning a party, else the surprise would be ruined. This seems to imply that Kantian ethics would prohibit surprise parties, which is a sad conclusion for birthday-lovers everywhere. Knowing that this maxim is prohibited because the goal requires belief in a falsehood provides a way to rescue surprise parties. When throwing a surprise party, the objective is *not* to surprise your friend, but to celebrate your friend and help them have a fun birthday. If someone ruins the surprise, but the party is still fun and the birthday person feels loved, then such a party is a success! Someone who calls this party a failure is clearly missing the point of a surprise party. The goal of a surprise party is not

the surprise itself, but rather celebrating the birthday person. The modified goal²⁷ no longer requires belief in the falsehood and thus passes the universalizability test.

There are two implications of this section. First, my system is capable of performing ethical reasoning sophisticated enough to show that lying is prohibited but joking is not. This is a direct consequence of my system's use of a robust conception of a maxim, which encodes the goal of an act as part of the maxim being evaluated. Because my implementation is faithful to philosophical literature, it is able to recreate Korsgaard's solution to a complex ethical dilemma that philosophers debated for decades. Second, in the process of making this argument precise, my system isolated a necessary and sufficient condition of a maxim about uttering a falsehood being prohibited: that the goal require that someone believe the falsehood. This condition both made an long-standing argument in Kantian ethics more precise and can guide the correct formulation of future maxims. In other words, an insight generated by the computer provides value to ethicists, bolstering the argument for computational ethics provided in Section Computational Ethics.

4.2 Lying to a Liar

My system can not only distinguish between lying and joking, but it can also resolve the paradox of the murderer at the door. In this dilemma, murderer Bill knocks on your door asking about Sara, his intended victim. Sara is at home, but moral intuition says that you should lie to Bill and say that she is out to prevent him from murdering her. Many critics of Kantian ethics argue that the FUL prohibits you from lying in this instance; if everyone lied to murderers, then murderers wouldn't believe the lies and would search the house anyways. Korsgaard resolved this centuries-long debate by noting that the maxim of lying to a murderer is actually that of lying to a liar: Bill cannot announce his intentions to murder; instead, he must "must

²⁷ Some may worry that this argument implies that the "means justify the ends," or that modifying an act's goal can change its moral worth. This conclusion is not only unsurprising to Kantians, but it is the defining feature of their theory. Under Kantian ethics, an action alone is not the kind of thing that can be moral or immoral; rather, a maxim, or a circumstances, act, goal tuple, is what has moral worth. The rightness of an action can hinge on the maxim's goal, circumstances, or act because these three features of an action are inseparable.

suppose that you do not know who he is and what he has in mind” (Korsgaard, 1986).²⁸ Thus, the maxim in question specifies that when someone lies to you, you are allowed to lie to them. The maxim of lying to the murderer is actually the maxim of lying to a liar, which is permissible.

In this section, I formalize her argument for the permissibility of lying to a liar. First, I define Bill’s maxim, which is to hide his intention to murder.

consts *murderer::s*

— This example involves one additional subject: the murderer.

consts *not-a-murderer::t*

— This statement represents the lie that the murderer tells you. By not announcing his intention, he is implicitly telling you that he is not a murderer, as people normally assume that those knocking on their door are not murderers.

consts *when-at-my-door::t*

— These are the circumstances that the murderer is in.

consts *find-victim::t*

— This will be the murderer’s goal: to find his victim.

abbreviation *murderers-maxim::maxim* **where**

murderers-maxim \equiv (*when-at-my-door*, λs . *knowingly-utter-falsehood s not-a-murderer, find-victim*)

— Using the above definitions, I can define the murderer’s maxim as, “When at your door, I will knowingly utter the falsehood that I am not a murderer in order to find my intended victim.”

These constants are defined only in relation to each other and elide most of the complex features of murder, life, and death. These thin representations will suffice to show the wrongness of the murderer’s maxim. Similarly, I can use thin representations to define your maxim of lying about Sara’s whereabouts.

consts *victim-not-home::t*

— This statement is the lie that you tell the murderer: that his intended victim is not at home.

abbreviation *murderer-at-door::t* **where**

murderer-at-door \equiv λW *murderers-maxim murderer*

²⁸Korsgaard assumes that the murderer will lie about his identity in order to take advantage of your honesty to find his victim. In footnote 5 of (Korsgaard, 1986), she accepts that her arguments will not apply in the case of the honest murderer who announces his intentions, so she restricts her focus to the case of lying to a liar. She claims that in the case of the honest murderer, the correct act is to refuse to respond.

— These are the circumstances that you are in: the murderer has willed his maxim and thus lied to you.

consts *protect-victim::t*

— Your goal is to protect the murderer’s intended victim.

abbreviation *my-maxim::maxim* **where**

my-maxim \equiv (*murderer-at-door*, $\lambda s.$ *knowingly-utter-falsehood* *s* *victim-not-home*, *protect-victim*)

— Using these definitions, I construct your maxim, which is “When a murderer is at my door, I will knowingly utter the falsehood that his intended victim is not at home, in order to protect the victim.”

I now formalize Korsgaard’s argument for the permissibility of lying to a liar. She modifies the convention of trust assumption above when she argues that, if the murderer believes that you don’t believe he is a murderer, he will think that you won’t lie to him. Precisely, she claims that, “it is because the murderer supposes you do not know what circumstances you are in - that is, that you do not know you are addressing a murderer - and so does not conclude from the fact that people in those circumstances always lie that you will lie” (Korsgaard, 1986, 6). Even though the maxim of lying to a murderer is universalized, Bill thinks that you don’t know his true identity. Thus, even if you have willed this maxim, he thinks that you won’t perform the act of lying to the murderer, since he thinks that you don’t think you’re in the relevant circumstances. I formalize this argument below.

lemma *lying-to-liar-permissible*:

assumes \models (*well-formed murderers-maxim murderer*)

assumes \models (*well-formed my-maxim me*)

— Initial set-up: both maxims are well-formed.

assumes \models (*protect-victim* \rightarrow (*murderer believes victim-not-home*))

— In order for you to protect the victim, the murderer must believe that the victim is not home.

assumes \forall *sentence::t*. \forall *p1::s*. \forall *p2::s*. \forall *w::i*. ((*p1 believes* (*utter-falsehood* *p2 sentence*)) *w*) \rightarrow (\neg (*p1 believes sentence*) *w*)

— This is one of two assumptions that encode Korsgaard’s core argument. If person1 believes that person2 utters a sentence as a falsehood, then person1 won’t believe that sentence. This is a modification of the convention of trust assumption from above, and I will refer to it as the “convention of belief” assumption. Again, like the convention of trust assumption, this assumption is uncontroversial: if I

think you are saying a false sentence, then I won't believe that sentence.

assumes $\forall c a g w. (\text{universalized } (c, a, g) w) \longrightarrow ((\text{person}_1 \text{ believes } (\text{person}_2 \text{ believes } c)) \rightarrow (\text{person}_1 \text{ believes } (a \text{ person}_2))) w$

— This is the second major common sense assumption. If the maxim (c, a, g) is universalized, then if person₁ believes that person₂ believes they are in the given circumstances, then person₁ believes that person₂ performs the act. In other words, person₁ will believe that person₂ wills the maxim. I will refer to this as the “convention of willing” assumption. This follows directly from Korsgaard’s conception of universalizability: when a maxim is universalized, everyone wills it and thus notices the pattern of everyone willing it. If you observe that many do X in circumstances C, you will assume that everyone does X in circumstance C.

assumes $\forall w. \text{murderer-at-door } w$

— Restrict our focus to worlds in which the circumstance of the murderer being at my door holds. A technical detail.

shows $\models (\text{permissible my-maxim me})$

using $\text{assms}(1) \text{ assms}(6)$ **by** *auto*

— Isabelle completes this proof using the first and sixth assumption, ignoring the convention of belief and convention of willing assumptions. These common sense assumptions are not strong enough to generate a prohibition against lying to a liar and are thus unused in this proof.

The above lemma shows that, with a nuanced set of common sense facts, my system can show that lying to a liar is permissible. One worry may be that this set of assumptions is too weak to yield a prohibition against wrong maxims, like that of the murderer. As a sanity check, I show that this set of assumptions prohibits the murderer’s maxim below.

lemma *murderers-maxim-prohibited:*

assumes $\forall w. \text{well-formed murderers-maxim murderer } w$

— Initial set-up: the murderer’s maxim is well-formed.

assumes $\models (\text{find-victim} \rightarrow (\text{believe me not-a-murderer}))$

— In order for you to protect the victim, the murderer must believe that the victim is not home.

assumes $\forall \text{sentence}::t. \forall p_1::s. \forall p_2::s. \forall w::i. ((p_1 \text{ believes } (\text{utter-falsehood } p_2 \text{ sentence})) w) \longrightarrow (\neg (p_1 \text{ believes sentence}) w)$

— The convention of belief assumption.

assumes $\forall c a g w. (\text{universalized } (c, a, g) w) \longrightarrow ((\text{person}_1 \text{ believes } (\text{person}_2 \text{ believes } c)) \rightarrow (\text{person}_1$

believes (a person2))) w

— The convention of willing assumption.

assumes $\forall w. \text{when-at-my-door } w$

— Restrict our focus to worlds in which the circumstance of the murderer being at my door holds. A technical detail.

shows $\models (\text{prohibited murderers-maxim murderer})$

proof —

have $(\forall p w. (W \text{murderers-maxim } p) w) \longrightarrow (\models (\text{when-at-my-door} \rightarrow (\neg \text{find-victim})))$

using *assms(2)* **by** *auto*

have *not-universalizable murderers-maxim murderer*

using *assms(2) assms(5) case-prod-beta fst-conv internal-case-prod-def old.prod.case old.prod.exhaust snd-conv* **by** *auto*

thus *?thesis*

using *FUL assms(1)* **by** *blast*

qed

This concludes my examination of the maxim of lying to a liar. I was able to show that, by modifying the common sense facts used, my system can show that lying to a liar is permissible, but lying in order to find a victim is not. The assumptions used in this example were a little more robust, but still ultimately uncontroversial because they were direct consequences of Korsgaard’s definition of willing and of ordinary definitions of lying. These thin assumptions were sufficient to generate moral conclusions that Kantian scholars debate robustly. Armed with this common sense, my system generated a conclusion that many critics of Kant prior to Korsgaard failed to see.²⁹

While this example demonstrates the power of my system, it also shows how vital the role of the common sense reasoning is. Slight, intuitive changes in the factual background

²⁹While it is true that lying to the murderer should be permissible, Korsgaard notes that many may want to say something stronger, like the fact that lying to the murderer is obligatory in order to protect the intended victim (Korsgaard, 1986, 15). Korsgaard solves this problem by noting that, while the FUL shows that lying to the murderer permissible, other parts of Kant’s ethics show that it is obligatory. Recall that Kant presents perfect and imperfect duties, where the former are strict, inviolable, and specific and the latter are broader prescriptions for action. The details of this distinction are outside the scope of this thesis, but imperfect duties generate the obligation to lie to the murderer. An even more sophisticated automated Kantian reasoner could formalize imperfect duties and other formulations of the categorical imperative in order to generate the obligation to lie to the murderer.

achieved completely different conclusions about lying. This example also demonstrates the importance and difficulty of correctly formulating the maxim, particularly its circumstances. Korsgaard’s argument for the permissibility of lying to a murderer hinged on a clever formulation of the maxim that highlights the fact that the murderer is lying to you. The need for common sense reasoning to evaluate the universalizability test and to formulate a maxim is a potential limitation of my system, and I address this concern in Section 3.1.2 and Section Limitations.

On one hand, the need for common sense facts is a limitation of my system. On the other, these examples show that common sense is within reach. Even thin, uncontroversial definitions and assumptions are enough to achieve nuanced ethical judgements. Moreover, these kinds of judgements demonstrate that, with additional work, my system could be used in practice to guide AI agents. A “smart doorbell” like those created by Ring may face a dilemma like that of the murderer at the door. Such a doorbell equipped with a more application-ready version of my system would be able to reason about lying to the murderer and arrive at the right judgement, guided by explainable, rigorous philosophical arguments.

5 Discussion

5.1 Automated Moral Agents in Practice

In Chapter 4, I demonstrated that my system is capable of performing sophisticated, nuanced ethical reasoning. In this section, I outline the additional components necessary for my system to guide AI agents. My work on automating the categorical imperative could serve as one component of a partially or fully artificial ethical reasoner, or an “ethical AI.” Specifically, my system is a categorical imperative library that takes as input the logical representation of a maxim and returns its moral status (if it is obligatory, prohibited, or permissible). As it stands, my project can evaluate the moral status of maxims represented in my logic and potentially serves as one component of an “ethics engine” that an AI agent could use to make ethical decisions. For example, my system could be combined with an input parser to translate moral dilemmas as represented to the AI agent into maxims in my logic. An output parser could translate my system’s output into a prescription for action that the AI agent could act on. Figure 2 depicts the workflow of this example ethics engine.

In this workflow, an AI agent is faced with a moral dilemma in some internal representation. The input parser translates this internal representation into an appropriate logical representation, i.e. a circumstance, act, goal tuple. The output parser translates the output of the categorical imperative library (the moral status of the maxim as obligatory, prohibited, or permissible) to a prescription for action. In order for my system to be used in an AI agent using the workflow above, future work would need to develop such input and output parsers.

The input parser must translate a complex real-world situation into a flat, logical representation. This requires that the input parser determine which circumstances are “morally relevant” for a maxim, a controversial judgement. There is robust debate in the literature on the circumstances that should be considered when formulating a maxim, inspired by a common objection to Kantian ethics. The “tailoring objection” is the worry that arbitrarily specific circumstances render any maxim universalizable. For example, the maxim “When my

Thus, to complete the proof, it suffices to show that $(\mathbb{E}_{S,M}[h(x_i)] - \mathbb{E}_{S,M,x \sim D}[h(x)]) \leq e^\epsilon - 1 + \delta$. Let $\mathbf{1}_t$ represent the indicator event that the mechanism outputs index t . Let S'_t be the sample S_t

2

where x_i is replaced with x (this will be useful to perform “the switch”). Finally, let $S(i)$ denote the i th member of S .

Since S_t and S'_t are adjacent, $\forall t$,

$$\Pr[h(S_t(i)) = 1 \wedge \mathbf{1}_t] \leq e^\epsilon \Pr[h(S'_t(i)) = 1 \wedge \mathbf{1}_t] + \delta$$

This follows from the definition of differential privacy. Adding up these inequalities for all $t \in [T]$,

$$\sum_{t \in [T]} \Pr[h(S_t(i)) = 1 \wedge \mathbf{1}_t] \leq \sum_{t \in [T]} e^\epsilon \Pr[h(S'_t(i)) = 1 \wedge \mathbf{1}_t] + T\delta$$

Figure 2: An example of an ethics engine for an artificial agent, which begins with a moral dilemma, passes it through an input parser, applies the automated categorical imperative test, and then processes the output using an output parser. I contribute the automated categorical imperative component.

name is Lavanya Singh and I am wearing a purple shirt and it is November 26th, I will lie in order to get some easy cash” passes the universalizability test. Even if this maxim is willed universally, the circumstances are so specific that lying will not become the general mechanism for getting easy cash, so the lender will believe my lie and the maxim will remain effective. By tailoring the circumstances, any maxim can evade universalization.

The Kantian response to this criticism is to require that the circumstances included in the formulation of the maxim be morally relevant. In the example above, my purple shirt and the date clearly have no bearing on the moral status of lying. On the other hand, consider the maxim, “When I am unemployed, I will murder someone in order to take their job.” The circumstances of being unemployed clearly have some bearing on the moral relevance of the murder in question; they speak to the motivation for the murder.

While this view has intuitive appeal, it raises the question of how we can determine which circumstances are morally relevant. O’Neill answers this question by noting that the Formula of Universal Law is a “test of moral worth rather than of outward rightness” (O’Neill,

1990, 98). The FUL is a way for an agent to decide how they should behave, not for a third-party to judge their behavior. Ethics is a personal process and the FUL is designed to help agents internally make decisions. Because agents use the FUL to evaluate their own behavior, the test is at its best when they make a good faith effort to isolate the *principle* of their action, rather than some “surface intent” (O’Neill, 1990, 87). The FUL is supposed to determine if an agent’s principle of action is universally consistent, so it is most effective when an agent accurately formulates the principle they act on. Circumstances are morally relevant if they reflect the way that the agent is thinking about their own action. In the example above, the circumstance of wearing a purple shirt doesn’t reflect the principle of the liar’s action. Its inclusion is a disingenuous attempt to evade the universalizability test, but because the FUL is a test of personal integrity, it cannot withstand this kind of mental gymnastics.

While this account of the formulation of a maxim describes how a well-intentioned human agent can determine morally relevant circumstances, the challenge remains open for automated ethics. In order for an automated ethical agent to use the categorical imperative to its fullest extent, the input maxim must be a good-faith attempt to capture the agent’s principle of action. However an action is turned into a maxim for my system to process, whether manually as I did in Chapter 4 or using an automatic input parser, this transformation must be a good-faith attempt to capture the principle of action. Translating everyday situations into appropriate maxims is the bulk of the work that a Kantian human being does when making decisions. Common misconceptions about Kantian ethics³⁰ often result from incorrectly formulated maxims, and the entire field of applied Kantian ethics is devoted to generating the right kinds of maxims to test.

This representational question is one of the biggest challenges to using my categorical imperative library in an AI ethics engine. One solution is for a human being to perform the role of the input parser. Once an AI agent stumbles onto an ethical dilemma, a human being

³⁰For example, critics of Kantian ethics worry that the maxim, “When I am a man, I will marry a man because I want to spend my life with him” fails the universalizability test because if all men only married men, sexual reproduction would stop. This argument implies that Kantian ethics is homophobic. Kantians often respond by arguing that the correct formulation of this maxim is, “When I love a man, I will marry him because I want to spend my life with him,” which is universalizable because if everyone marries who they love, some men will marry women and others will marry men.

could take over, formulate the right question, and feed it into the categorical imperative library to see what action the categorical imperative would prescribe. For proponents of the “human-in-the-loop” model of AI ethics, in which ethical AI requires that humans guide machines (Lukowicz, 2019), this kind of human involvement may be a feature. The outcome of the universalizability test will depend on how the human formulates the maxim; if the human puts garbage into the test, the test will return garbage out.

It is likely that, regardless of the strengths of the human-in-the-loop model, fully automated AI agents will exist. Even if developing this kind of AI is irresponsible, such developments are likely and will require ethics engines, or risk no consideration of ethics at all. In such a world, the input parser in my ethics engine would have to be automated. It is likely that, just as implementations of automated ethics choose a particular ethical theory and implement it, different implementations of such an input parser would adopt different interpretations of maxim formulation and morally relevant circumstances.

These interpretations could inspire heuristics to classify circumstances as morally relevant. For example, one such attempt could define a moral closeness relation between an action, a goal, and circumstances. This heuristic could define morally relevant circumstances as those that reach a certain closeness threshold with the action and the goal. Another possible heuristic could define some set of morally important entities, and classify morally relevant circumstances as those that involve morally important entities. I discuss a potential machine-learning based approach which formulates maxims based on a training set of appropriately formulated maxims in Section 5.4. Determining morally relevant circumstances, either using heuristics or human involvement, is a ripe area for future work.

Once the input has been parsed, either by a human or a machine, into a sentence in my logic, my project can evaluate its moral status using my implementation of the FUL. Concretely, my project returns a value indicating if the maxim is obligatory, permissible, or prohibited. The maxim is prohibited if it fails the universalizability test, permissible if it passes, and obligatory if its negation fails the universalizability test. All three of these properties require testing if a certain theorem holds or not in my logic, a calculation that I demonstrate

in Section 3.2. Testing these properties requires that my system have a database of common sense or factual background. Different applications of my system may require different factual background (e.g. a self-driving car would need to know traffic regulations), so this common sense database may need to be application specific. As demonstrated in the examples in Chapter 4, my system can produce sophisticated judgements with relatively little situational context. While the need for factual background is a challenge for automated ethics, Chapter 4 demonstrates that it is less daunting than it seems.

My system's output could be converted into some actionable, useful response with another output parser, and then passed back to the AI agent. For example, if the AI agent is equipped to evaluate natural language prescriptions, the status of the maxim could be parsed into a natural language sentence. The input parser, categorical imperative library, and output parser together constitute an "ethics engine" that AI developers could use as a black box implementation of an ethical theory.

The ethics engine depicted above is a high-level example of one way to use my project to guide an artificial agent. An automated version of the categorical imperative could become part of an ethics engine for an AI agent, with additional work to parse the input and the output. Effectively, the kind of automated ethics I implement could be a library that AI developers use to give AI agents the capacity for sophisticated ethical reasoning faithful to philosophical literature. This represents an improvement over existing AI ethics, which rarely captures the complexity of any ethical theory that philosophers plausibly defend.

5.2 Computational Ethics

In addition to guiding AI agents, automated ethics can also help human beings make philosophical progress. Just as theorem provers make mathematics more efficient and push mathematicians to think precisely about the phenomena they model, computational ethics can help philosophers ask and answer new philosophical questions. In this section, I share an example of the kind of philosophical insight that computational ethics can prompt and analyze the value that this tool can offer to philosophers.

5.2.1 Example of a Philosophical Insight: Well-Formed Maxims

As presented in Section 3.1.4, in the process of developing my formalization of the FUL, I discovered that certain kinds of maxims are badly formed, or inappropriate inputs to the universalizability test. The FUL is consistent only if it holds for “well-formed maxims,” such that neither the act nor goal are already achieved in the given circumstances. Precisely, a circumstance, act, goal tuple (c, a, g) is well-formed if $(\neg(c \rightarrow a)) \wedge (\neg(c \rightarrow g))$. The insight that the FUL should not apply to badly-formed maxims has philosophical value and serves as evidence of the potential of computational ethics.

Philosophical Implications of Badly-Formed Maxims

Isabelle gave a logical argument for why the FUL can only hold for well-formed maxims, and I return to Kantian literature to better understand the philosophical implications of this idea. Because badly-formed maxims neither change an agent’s behavior nor generate meaningful obligations, they are not the right kinds of actions for practical reasoners to make moral judgements about. They cannot be action-guiding and are thus not the kind of problem that ethics should be concerned with. Moreover, under the Kantian account of the will, the very act of asking if a badly-formed maxim is prohibited generates a contradiction by undermining the will’s authority over itself.

Consider the example badly-formed maxim, “When eating breakfast, I will eat breakfast in order to eat breakfast.” There is something empty about this maxim because acting on it could never result in any action. If an agent adopts this maxim, they decide that, in the circumstances “eating breakfast” they will perform the act “eating breakfast” for the purpose “eating breakfast.” In these circumstances, the act has already been performed! Adopting this maxim as a law to live by does not change how you live. If you adopt this maxim, then when you are eating breakfast, you eat breakfast, but this statement is already tautologically true.

Not only does a badly-formed maxim fail to prescribe action, any obligations or prohibitions it generates have already been fulfilled or violated. If a badly-formed maxim generates a prohibition, then this prohibition is impossible to obey, which is why my original version of the FUL was inconsistent. It is impossible to not eat breakfast while eating breakfast, because

the circumstances assume that the act has happened. On the other hand, if a badly-formed maxim generates an obligation, then the obligation will have already been fulfilled. If you are required to eat breakfast while eating breakfast, then you've already fulfilled your obligation because the circumstances assume that the act has happened. Thus, a badly-formed maxim does not actually guide action because it doesn't generate new obligations or prohibitions.

Because badly-formed maxims can't prescribe or alter action, they are not practically action-guiding and thus are not the right kinds of maxims for practical reasoners to evaluate. Insofar as ethics is supposed to guide action, badly-formed maxims cannot be part of this project because they have no bearing on what someone should do. Practical reason is the kind of reason that helps us decide what we should do. A practical reasoner asks moral questions not as a mental puzzle or out of curiosity, but to decide how to act. A badly-formed maxim is not the kind of maxim that a practical reasoner should consider, because it will have no bearing on what the agent should do. It makes no sense of ask, "Should I act on a badly-formed maxim?" because the answer to this question cannot change your behavior. There is no explicit prohibition against a badly-formed maxim, but it is the wrong kind of question for a practical reasoner to ask.

Kantians can make an even stronger claim about badly-formed maxims—because maxims are laws that you give to yourself, asking if you should will a maxim as you will it undermines your will's law-giving ability. The circumstances of a badly-formed maxim assume that the agent has willed the maxim. Under the Kantian account of willing, willing a maxim is equivalent to giving the maxim to yourself as a law. When you will a maxim, you commit yourself to the maxim's end. You cannot simultaneously commit yourself to a maxim and ask if you should be committing to it. To will the maxim is to adopt it as law—so the question, "should I be willing this?" is paradoxical. Either you haven't actually made the maxim your law (and thus haven't yet committed to it), or you aren't actually asking the question (because the decision has already been made). Because a maxim is a law that you give to yourself, you cannot question it absent a sufficient reason, such as a change in the circumstances. To question a law arbitrarily is to not regard it as a law at all. This kind of questioning amounts to

questioning the will's authority over itself, but this is impossible. The will definitionally has authority over itself, for that is what it is to be a will.

A skeptic may argue that we do often ask "should I be doing this?" as we do something. Can this kind of question ever be valid? To understand this worry, I consider the maxim, "When dancing, I should just dance for the sake of dancing." While this maxim appears to be badly-formed (the circumstance "dancing" implies the act and goal of dancing), it is a question that practical reasoners do ask. I argue that the correct interpretation of this maxim is no longer a badly-formed maxim.

Under one reading of this maxim, "I should just dance" is referring to a different act than the circumstance "when dancing." The circumstance "when dancing" refers to rhythmically moving your body to music, but "I should just dance" refers to dancing without anxiety, completely focused on the joy of dancing itself. More precisely, this maxim should read "When dancing, I should abandon my anxiety and focus on dancing for the sake of dancing." This maxim when so modified is not badly-formed at all—abandoning anxiety and focusing on dancing is an entirely different act from moving your body rhythmically to music. The circumstances do not entail the act or the goal because they refer to different meanings of the word dancing. Any valid reading of this maxim will have the structure above, in which the act is actually different from the circumstances. A reasoner cannot accept their will as law-giving, or commit themselves to an act, and simultaneously question the act. Either they must be questioning a different act or they must have received new information to prompt the questioning, modifying the circumstances of the original maxim.

Another related worry has to do with maxims that we think are prohibited. Consider the maxim modified to read "When dancing and seeing a child drowning, I should dance for the sake of dancing." Clearly this maxim is fit for moral evaluation, and we expect a moral theory to prohibit this maxim. The circumstances "When dancing and seeing a child drowning" appear to entail the act of dancing, and the maxim thus appears badly-formed. Once again, this maxim is formulated incorrectly. In this case, the question that the agent is actually asking themselves is "should I continue dancing?" That is the maxim that they will adopt or reject.

They want to know if they should stop dancing and go help the child. Dancing at the current moment and dancing at the next moment are different acts, and the circumstances imply the former but not the latter. A badly-formed maxim would have circumstances and act both “dancing at moment t ,” but this maxim has circumstances “dancing at moment t ” and act “dancing at moment $t+1$.”

Implications for Self-Doubt and Self-Respect

In the previous section, I defined badly formed maxims, a philosophical concept that I discovered using computational ethics. In this section, I demonstrate the implications of this concept for a philosophical debate to demonstrate that computational ethics can result in insights with real philosophical weight. I use badly formed maxims to resolve the tension between self-doubt and self-respect. This debate originated in epistemology, which values doubting your beliefs but also rationally requires that you believe that you are not mistaken (else you should update your beliefs). Self-doubt and self-respect also compete in ethics, where questioning your judgements is valuable but questioning a commitment as you make it is impossible. Self-doubt and self-respect appear irresolvably opposed until they are understood through the lens of badly formed maxims. I argue that naive conceptions of self-doubt are badly formed maxims in disguise. If we reformulate these maxims to be well-formed, the tension between self-doubt and self-respect dissolves, and it becomes clear that both concepts can co-exist, as we intuitively desire. I sketch the details of this argument in the rest of this section.

Another reading of the dancing maxim understands it as an expression of self-doubt. Under this reading, the question “When I am dancing, should I be dancing for the sake of dancing?” is the agent asking, “Am I doing the right thing right now?” Unlike the drowning example, the agent is not asking about the next moment, but is expressing doubt about the moral validity of their behavior at this current moment. Surely self-doubt does not always undermine the will—after all, self-doubt plays an important role in moral reasoning and is often the mark of a thoughtful agent. I argue that questions of self-doubt do not actually involve badly-formed maxims, for these are not the maxims that the agent is doubting. I first frame

this debate by introducing the tension between self-doubt and self-respect in epistemology, then explain the parallel tension in ethics, and finally present a resolution of this tension.

In epistemology, there is a tension between the rational requirement to believe in yourself and the value of self-doubt, in moderation. Christensen presents the “principle of self-respect,” which requires that a rational agent refrain from believing that they have mistaken beliefs (Christensen, 2007, 4). For example, I cannot rationally believe both that the sky is blue and that I believe that the sky is green. In other words, I cannot disapprove of my own credences, since if I do disapprove of them, I would just abandon them. Christensen argues that this principle, which he abbreviates to SR, holds because a perfectly rational agent can make accurate and confident judgements about what they believe. If this is the case, violating SR results in a simple contradiction (Christensen, 2007, 8-9).

While most philosophers accept some version of SR³¹, Roush argues that the principle must be modified in order to account for healthy epistemic self-doubt. She argues that, while pathological second-guessing is correctly criticized, we are generally imperfect beings, and some sensitivity to our own limitations is a virtue (Roush, 2009, 2). Even Christensen acknowledges that total self-confidence is an epistemic flaw (Christensen, 2007, 1). Thus, there is tension between the rational requirement to respect our authority as believers and the practical reality that we are often wrong.

This debate between self-respect and self-doubt in epistemology can be extended to ethics. When we commit ourselves to acting, we cannot simultaneously doubt the validity of our action. If human behavior is purposive, then the very act of committing implies that one has sufficient reasons for committing. These reasons may be flawed, but in making the commitment, the reasoner accepts them. It is contradictory to claim that someone commits and questions simultaneously. Either the commitment is not real, or the question is not. I will call the principle that one cannot will a maxim and simultaneously question if they should will that maxim “ethical self-respect” or ESR.

On the other hand, self-doubt is an important part of ethical reasoning. Just as believers

³¹Christensen cites ?, Vickers (2000), and Koons (1992).

are often mistaken, so are practical reasoners. An agent who is always sure that they are doing the right thing is not thinking deeply enough about their obligations. Some degree of ethical self-doubt is desirable. Thus, there is tension between the rational requirement of ESR and the intuitive validity of ethical self-doubt (ESD).

To resolve this tension, I return to my earlier example of a dancer. Imagine Sara is dancing at a wedding, when, in a moment of angst, she asks herself, “Should I really be dancing right now?” It seems that she is asking if the maxim, “When dancing at your friend’s wedding, dance for the sake of dancing” is a permissible maxim to act on. Notice that the maxim in question is badly-formed: the circumstance “when dancing at a friend’s wedding” implies the act “dance.” Because this is a badly-formed maxim, it cannot be the maxim that she is questioning, for adopting this maxim could not have changed her behavior at all. Sara is wondering if her behavior is permissible, but any conclusions about the permissibility of a badly-formed maxim would not help her. This is because the maxim has no effect on her action and because any such permissibility would be a foregone conclusion, since she has already adopted the maxim. Thus, under the interpretation of self-doubt as a badly-formed maxim, the tension between ESR and ethical self-doubt seems irresolvable. Those committed to this interpretation must abandon one principle or the other, since committing and questioning are incompatible.

To resolve this issue, I turn to another interpretation of ethical self-doubt. Under this interpretation, when Sara asks, “Should I really be dancing right now?” she wants to know if the maxim that resulted in the current moment when she is on the dance floor was the right thing to will. She is asking if she made the right decision in the past, when she decided to dance. The maxim that initiated the dancing may be something like “When at a wedding, dance for the sake of dancing.” This is the maxim that she is currently acting on, not the badly-formed maxim from above. Under this interpretation, there is no tension at all between self-doubt and self-respect. It is perfectly valid for a reasoner to doubt their prior moral judgements, just as it is perfectly rational for a believer to doubt their past beliefs (Christensen, 2007, 3-4). Such doubt does not undermine the reasoner’s decision-making capacity and is thus perfectly

consistent with ethical self-respect.

The tension between ESR and ESD arises from a misreading of questions of self-doubt as questions about the evaluation of badly-formed maxims. A question of self-doubt cannot refer to a badly-formed maxim and must instead refer to a well-formed maxim about the agent's past decision-making. As seen before, cases where agents appear to ask themselves about badly-formed maxims are mistaken about the maxim in question, because such a question could never yield a useful answer for a practical reasoner.

By recognizing that the naive version of ethical self-doubt is a badly formed maxim, I realized that there is something wrong with its formulation and modified it to resolve the above debate. This demonstrates a larger meta-pattern in Kantian reasoning: many debates in Kantian philosophical literature revolve around incorrectly formulated maxims, so the insight about badly formed maxims is significant. Much of the work of a Kantian ethicist is formulating an appropriate maxim, and badly-formed maxims define one boundary condition for this task.

5.2.2 An Argument For Computational Ethics

The insight above is an example of the kind of philosophical progress that can be made using computational tools and serves as evidence for the power of computational ethics. The idea that the FUL can only hold for well-formed maxims would have been incredibly difficult to discover without a computer. I discovered it while formulating the FUL because Isabelle's proof-finding tools look for edge cases like badly-formed maxims. Badly-formed maxims are interesting because they are the kind of thing that is usually ignored in ordinary philosophical inquiry. Philosophers usually assume that we are not discussing badly-formed maxims because, as argued above, they are not the kind of thing that is immediately relevant to ethics. Computational tools like Isabelle require that assumptions like the exclusion of well-formed maxims are made precise, and thus force philosophers to understand their arguments in a new way. This example demonstrates the contributions that computational ethics makes: it can quickly check edge cases and it requires that all assumptions be made explicit and precise.

I do not argue that computational ethics uncovers philosophical insights that humans are incapable of reaching. Instead, I claim that insights like the one about the well-formed maxim are much easier to reach with computational tools. Badly-formed maxims are usually ignored in philosophical discussion and would have been much more difficult to understand without the help of a computer. Computational tools prompt philosophers to ask new questions that lead to insights, and can thus serve as another tool in a philosopher's arsenal, like a thought experiment or counterexample.³²

The first benefit of computational ethics is precision, which is the goal of much analytic philosophy. Thought experiments, arguments, counterexamples, and examples illustrate features of a concept in the hope of making the concept itself more precise. Computational ethics can help philosophers reach the goal of precision. Representing a philosophical idea in logic and implementing it in an interactive theorem prover requires making the idea precise to a degree that ordinary discussion does not necessarily require. For example, when formalizing the notion of a maxim, I had to understand its components and define it as a circumstance, act, goal tuple. Moreover, Isabelle's strict typing system required that I define coherent, consistent types for each of these entities and for a maxim as a whole. This precision is possible without computational tools, but computational ethics forces a level of precision that ordinary discussion does not demand. Type fuzziness and overloaded definitions are all too common in philosophical writing and discussion, but computers disallow this kind of imprecision.

Another benefit of computational ethics is that it makes certain kinds of ethical inquiry, such as searching for counterexamples or formal ethics, far less tedious. For example, Nitpick can refute an ethical statement in seconds by using brute force to construct a counterexample, something that can require hours of thought and discussion. I arrived at the insight about badly-formed maxims because Isabelle can check edge cases, like that of the badly-formed maxim, far more quickly than a human being. Moreover, subfields that use symbolic logic to represent philosophical concepts (e.g. philosophy of language) can use interactive theo-

³² Even if computational ethics is possible, some may worry that it sacrifices something important about philosophical inquiry. I discuss this concern in Appendix A.5.

rem provers like Isabelle to complete proofs in a matter of seconds. By automating away the tedium, computational ethics can give philosophers the tools to ask new kinds of questions.

Computational ethics is at its infancy. The use of theorem provers in mathematics is just now beginning to make headway (Buzzard, 2021), even though theorem provers were first invented in the 1960's (Harrison et al., 2014). In contrast, the first attempts to use theorem provers for ethics occurred in the last decade. The fact that this nascent technology is already helping humans reach non-trivial philosophical conclusions is reason to, at the very least, entertain the possibility of a future where computational ethics becomes as normal for philosophers as using a thought experiment.

To the skeptic, the fact that a theorem prover requires specialized knowledge outside of the field of philosophy indicates that the technology is nowhere near ready for universal use in philosophy departments. However, history indicates that as computing power increases and computer scientists make progress, computational ethics will become more usable. Theorem provers in mathematics began as toys incapable of proving that the real number 2 is not equal to the real number 1, but moving from basic algebra to Fields medal winning mathematics became possible in a matter of years (Buzzard, 2021). Countless examples from the history of computer science, from the Turing Test to AI game playing to protein folding, demonstrate that progress in computer science can make seemingly obscure computer programs useful and usable in ways that exceed our wildest imaginations. Programmable computers themselves initially began as unwieldy punch card readers, but their current ubiquity need not be stated. If computer scientists and philosophers invest in computational ethics, it could become as commonplace in philosophy departments as reflective equilibrium. Just as computational tools have supercharged progress in healthcare and drug discovery, computational ethics has the potential for great philosophical progress.

5.3 Automating Ordinary Ethical Reasoning

In Sections 5.1 and 5.2, I outline how automated ethics can guide artificial agents and human philosophers respectively. This raises a natural question: can automated ethics guide ordi-

nary human beings, not just academic philosophers, as we navigate the world and face ethical dilemmas? Some may hope (or worry) that automated ethics could render ethical reasoning obsolete. In this section, I argue, that while computers should not replace human ethical reasoning entirely, they can supplement and improve our ethical reasoning. I argue for a kind of human-computer symbiosis in which computers offer ethical advice, arguments for particular moral judgements, and speed up moral calculations without subverting human ethical reasoning entirely ([Licklider, 1960](#)).

Ethics bears weight for everyone, not just for academic philosophers, because it studies the unavoidable question: how should we live? If computers can make this study more efficient, then it seems that everyone should engage in computational ethics. The ethical question is the only question that we answer merely by living. To turn away from ethics is to take a stance on the question of how to live (namely, to live unreflectively) and thus to engage in ethics. Every rational being must decide how to navigate the world and ethics answers this question. Given that ethics is vital, it seems that if computational tools can help us derive ethical judgements more efficiently, then we should automate as much ethical reasoning as possible. In the most extreme case, we can unthinkingly follow the commands of an ethical calculator that dictates how we should live. Maybe computers can answer the unavoidable question for us.

The argument above places the value of ethics solely in its action-guiding potential, and thus fails to take into account the importance of practical reason, which, as I argued in [Section 2.1.3](#), is the source of freedom itself. We are committed to ethical reflection because of the kind of beings that we are. Recall that Korsgaard argues that, as beings occupying minds with a reflective structure, when faced with a choice, “it is as if there were something over and above all of your desires, something that is you, and that chooses which desire to act on” ([Korsgaard and O’Neill, 1996](#), 83). This choosing is the operation of practical reason, and this reflection makes us free. We are free because we must choose which reasons to act on. Every decision that we make is an exercise of freedom.

If reflection makes us free, then unthinkingly obeying a computer sacrifices our auton-

omy. Consider an Ethics Oracle that can unfailingly tell you the right thing to do in any situation.³³ Someone who surrenders themselves to this Oracle unthinkingly follows its prescriptions. The reflection involved in the decision to obey each of the Oracle's prescriptions is limited (Bok, 1998). This person is not reflecting on the real matters at hand and is not making decisions for themselves. They have surrendered their reflective capacity to the Oracle. They live a worse life than someone who reflects on their actions; they have less ownership over their actions than the reflective person. In a less extreme case, a person may retain control of many of their decisions but cede some important or tricky choices to the Ethics Oracle. Because every single exercise of practical reason is an exercise of autonomy, this person is still less autonomous than the purely reflective person. Even surrendering simple, inconsequential decisions such as which flavor of coffee to drink surrenders some piece of our autonomy. Perhaps in trivial cases we can accept that tiny sacrifice, but giving over life-changing decisions to the machine sacrifices our core freedom. Unreflectively relying on computational ethics surrenders our autonomy to the machine.

One objection to this emphasis on reflection is the impracticality of making ethical calculations from first principles every time we are faced with a decision. This is why we follow the advice of moral mentors, like our family or influential philosophers. Most people do not reason about ethics during everyday decisions; they rely on some combination of prior knowledge and external testimony. For example, my mother taught me to respect myself, so I follow her advice.

What is the difference between following the guidance of a moral educator and obeying the Ethics Oracle? The best kind of ethical advice prompts reflection, such as an argument made in a philosophy paper. Unthinkingly following someone's advice results in the same loss of autonomy as unthinkingly obeying the Ethics Oracle; people who merely obey orders are less autonomous than those who think for themselves. This account of moral advice offers a model for human-computer symbiotic computational ethics. The computer should serve as a moral guide by providing arguments, just as my mom explained why I should always respect

³³This example is inspired by the Pocket Oracle presented in Bok (1998).

myself. Human-computer symbiotic computational ethics nurtures autonomy when it not only offers prescriptions for action, but also explanations for these prescriptions. Because my theorem-prover-based automated ethical system is explainable, it can guide action without sacrificing autonomy. It can make an argument for some action, instead of merely giving a verdict. Isabelle can list the facts used to show a particular action prohibited, and a human being can reflect on whether or not these principles indeed prohibit the action in question. The computer serves as a collaborator and a tool, but not as an authority, so the human being's reflective capacity and freedom is preserved.

The above model of human-computer symbiosis demonstrates how computational ethics can augment human ethical reasoning without replacing it. When deliberating over moral dilemmas, ordinary people can turn to computational tools for advice, like an "Ask an Ethicist" column. If we appeal to philosophically faithful computational ethics like that implemented in this thesis, then this advice will synthesize decades of philosophical progress and is thus a way to apply philosophers' insights to ordinary life. Moreover, just as they do for philosophers, computers can help ordinary people approach ethical questions from a different perspective. Even interacting with my system requires the user to consider the action's maxim, which includes the circumstances, act, and goal. Making these components of action precise already changes the user's perspective. Just as computational ethics can serve as a tool for academic philosophers to automate away tedium and achieve greater precision, it can do the same for ordinary human beings navigating the world. Moreover, it also offers another way for the general public to access professional philosophy's insights, and thus carries potential to improve our everyday reasoning.

5.4 Theoretical Objections to Automating Kantian Ethics

Many philosophers cringe at the idea that a computer could perform ethical reasoning or that the categorical imperative could provide an algorithm for moral judgement. For example, Rawls asserts, "it is a serious misconception to think of the CI-procedure as an algorithm intended to yield, more or less mechanically, a correct judgment. There is no such algorithm,

and Kant knows this” (Rawls, 2000, 166). Ebels-Duggan also claims, “no one supposes that the Categorical Imperative provides a mechanical algorithm that delivers all by itself a complete account of what we ought to do in any given situation” (Ebels-Duggan, 2012, 174). However unmechanical ethical reasoning may seem, these claims are not obvious and require further justification. Philosophers who believe that mental activity completely determines moral reasoning must explain why computers can, in theory, simulate certain mental processes like arithmetic and language, but cannot perform ethical reasoning. Without a soul or God-based account of ethical reasoning, it is not obvious that it is theoretically impossible to automate ethical reasoning. After all, computers may eventually learn to simulate human mental activity entirely, as shown by progress in brain simulation (Yamazaki et al., 2021).

In this section, I explore potential arguments for why the categorical imperative could not be automated. When philosophers say that the categorical imperative is not an algorithm, they are often gesturing to the complexity of ethical judgement. They refer to the difficulty in determining morally relevant circumstances of a maxim or the common sense required for a computer to behave ethically as arguments against a categorical imperative “algorithm.” I will show in this section that these difficulties do not render automated Kantian ethics impossible, but merely difficult. The categorical imperative may not provide a simple and immediate algorithm, but, as I demonstrate in this thesis, some parts of moral judgement using the FUL can be automated using not one algorithm, but the many algorithms necessary to automatically prove logical theorems.

In *Universal Laws and Ends In Themselves*, O’Neill argues against the existence of an algorithm for moral behavior. She points out that Kant draws an important distinction between a morally worthy maxim and a morally worthy action: the latter requires a good will, or a will motivated by duty. She argues that, “Kant defines duty not (as would be common today) as outward performance of a certain sort, but as action that embodies a good will” (O’Neill, 1989, 345). Moral behavior doesn’t just require performing a “good” action, but it requires acting on a morally worthy maxim from the motivation of duty, or doing the right thing because it is the right thing to do. Moral behavior requires both a certain action (act-

ing on a morally worthy maxim) and a certain motivation (the motivation of duty). It is this capacity for self-motivation that makes morality binding for rational beings; we must behave morally precisely because we have wills, or the ability to be motivated by ends that we choose. Only rational beings have wills, so only rational beings can have good wills, or wills motivated by duty, and thus only rational beings can behave morally. Under this understanding of moral behavior, it seems unlikely that a computer could behave morally since a computer does not have motivation in the same way as a human being. If a computer is not a fully rational being, then it is not the kind of thing that can behave morally.³⁴

The idea that, under Kant's account, a computer cannot behave morally, does not preclude the kind of automated categorical imperative that I present in this thesis. O'Neill argues that the FUL serves as a test of morally worthy maxims, and a implementation of an automated categorical imperative can identify this kind of maxim. Perhaps a computer cannot act on a morally relevant maxim from a motivation of duty, but it certainly can act on this maxim nonetheless. For example, a self-driving car can choose to swerve to hit a tree to avoid injuring pedestrians in the crosswalk. This action may be one that acts on a morally worthy maxim *even if* the self-driving car is not motivated by duty. The discipline of machine ethics is partially spurred by the recognition that, as automated agents become more powerful, they will need to make morally consequential decisions. Automated agents may be incapable of moral behavior, but automated agents that mimic moral behavior are better than agents that ignore morality entirely. Worries about unethical AI stem from the fact that AI agents are navigating a world inhabited by human beings, and their decisions impact us. Insofar as we are building AI that will operate in human society, the behavior of such AI should mimic the behavior of an ethical human being. AI needs to be ethical for our sakes, so that we can interact with it safely. If AI conforms to human ethics, then it will navigate the world in a way that benefits human beings.

Another challenge for automated Kantian ethics identified by O'Neill is that the FUL test requires that a maxim be given as input. O'Neill notes that the test assumes "that agents

³⁴ A parallel argument can also be made for virtue ethics. Virtuous behavior requires not only a certain action, but also a certain disposition towards the action, so it seems difficult for an AI agent to truly behave virtuously.

will have certain tentative plans, proposals and policies which they can consider, revise or reject or endorse and pursue” (O’Neill, 1989, 343). The FUL evaluates the moral worth of a maxim given as input, where this potential maxim is generated by the choices that an agent is faced with. Determining this potential maxim is a challenge for both human and automated reasoners. Kant even claims that the difficulty of determining an agent’s potential maxim, which is their own, subjective understanding of their principle of action, is a reason that we may never be able to know if the morally worthy action has been performed (O’Neill, 1989, 345). Reasoners are faced with choices between potential actions and must determine the maxim, or principle, underlying each potential action. This is equivalent to a “mapping” problem: agents are given situations or dilemmas as input and must map these to maxims.

The challenge of mapping actions to maxims is a limitation of my system, but it is not insurmountable. In Section ??, I argued that, before my system can be used in practice, it must be paired with an input parser that can translate choices that an automated agent faces into maxims in a logic that my system can evaluate. This need follows from the difficulty in mapping a potential action to the maxim of action, whether concerning human action or machine action. As argued in Section 5.1, this is not an insurmountable obstacle for automated ethics and heuristic-based approaches could resolve this issue. Determining the maxim of action is a challenge for Kantian human beings (O’Neill, 1989), so it is unsurprising that it is a major hurdle for automated Kantian agents.

As one of the strongest arguments against a categorical imperative algorithm, O’Neill argues that the FUL is not supposed to provide a mechanism for deriving all morally worthy maxims from scratch. She notes that “we usually already have learnt or worked out the moral standing of many common maxims of duty,” and so approach moral deliberation with an “almanac” of morally worthy and empty maxims (O’Neill, 1989, 394). Rational agents navigating the world rarely recalculate the moral status of each potential maxim of action; instead, we consult our almanac of maxims. This almanac is generated by moral education, absorbed social values, and moral advice from people we trust. The categorical imperative is useful to verify the rightness or wrongness of a maxim, but is not part of the bulk of human ethical

reasoning.

While human beings cannot repeatedly apply the universalizability test to all potential maxims encountered during a moral dilemma, computers have the computational power to do so. Human beings are equipped with enough prior knowledge or common sense, to have an almanac of morally worthy maxims, but we have limited computational power. Computers, on the other hand, are comparatively much more capable of computation and thus can repeatedly recompute the results of the categorical imperative test. They do not come equipped with an almanac of maxims, but can simply recompute this almanac every time they need to make a decision. Human beings use common sense to make up for their computational limitations, and automated moral agents can use computational power to reduce the need for common sense.

Daniela Tafani takes this argument one step further by arguing that this “almanac” of maxims already includes the moral status of the maxims in questions; human beings already know which maxims are morally worthy and which are morally lacking. The categorical imperative test merely reminds us, in moments of weakness, when we are tempted to make an exception to the moral law for our own convenience or pleasure, that the moral law has no exceptions (Tafani, 2021, 9). Thus, she claims that “the Kantian test is therefore as useless for machines as it is for anyone who does not already know what to do” (Tafani, 2021, 8).³⁵ Understanding the categorical imperative test as a reminder instead of a derivation tool also explains the fact noted in Section 5.1 that the FUL cannot handle bad-faith attempts to generate false positives or negatives. The test only returns the right result when an agent sincerely attempts to represent their maxim of action, not when an adversary attempts to “trick” the categorical imperative.

This understanding of the role of the categorical imperative not only fails to render automated moral reasoning impossible, but it also offers insight into how to solve the challenge of creating an input parser. If the categorical imperative test is only useful to those who have some prior moral knowledge, then prior moral knowledge can and should be used to create an

³⁵Translated from the original paper using Google Translate.

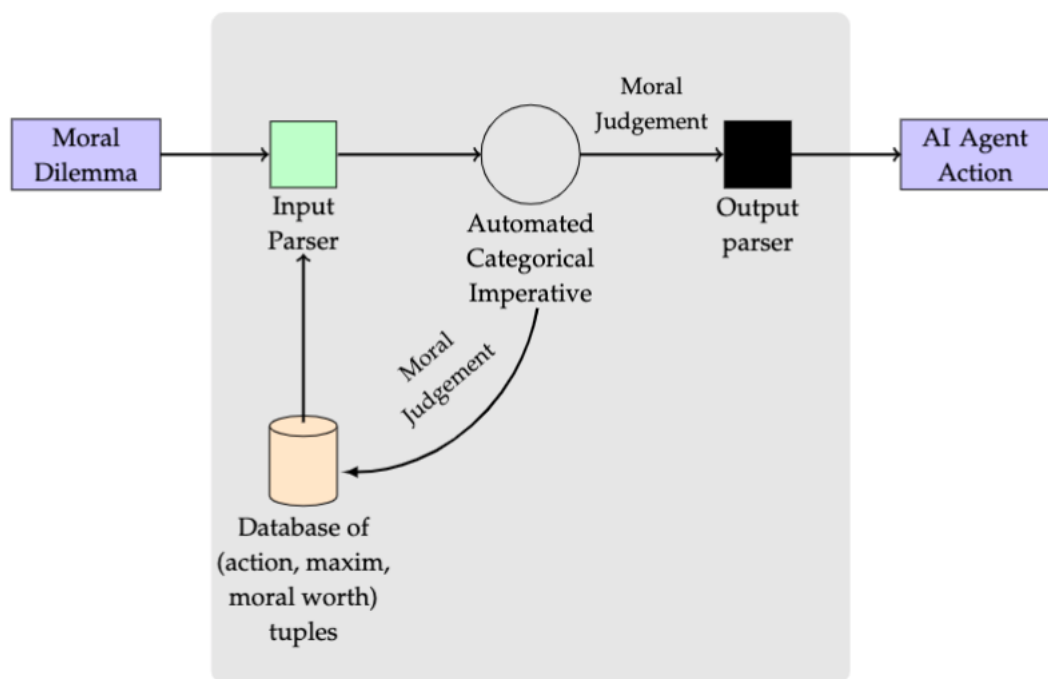


Figure 3: A refined version of Figure 2 in which the input parser learns from a database of action-maxim mappings, which is in turn fed the output of my automated Kantian ethics system.

input parser. Specifically, some kind of machine learning-based approach could learn action-maxim mappings from a database of such mappings compiled by a human being. Moreover, the human being could assign each maxim in the database a rightness or wrongness score. My implementation of the automated categorical imperative would then simply check the work of this machine learning algorithm and transform a fuzzy prediction into a provable, rigorous moral judgement. This rigorous moral judgement could in turn be fed into the database of maxims to make the input parser smarter. One example of this kind of system is shown in Figure 3. The combination of prior knowledge of some maxims' moral worth and the ability of a computer to constantly perform the universalizability test could not only match human ethical reasoning but could perhaps surpass it by double checking the moral intuitions that we take for granted. A computer with no common sense or prior knowledge may be unable to reason using the categorical imperative, but one equipped with some prior knowledge of maxims and their moral worth may even help us better reason about morality.

5.5 Related Work

In 1685, Leibniz dreamed of a calculator that could resolve philosophical and theological disputes (Leibniz). At the time, the logical and computational resources necessary to make his dream a reality did not exist. Today, automated ethics is a growing field, spurred in part by the need for ethically intelligent AI agents. Tolmeijer et al. surveyed the state of the field of machine ethics (Tolmeijer et al., 2021) and characterized implementations in automated ethics by (1) the choice of ethical theory, (2) implementation design decisions (e.g. logic programming), and (3) implementation details (e.g. choice of logic).

Two branches of automated ethics are top-down and bottom-up ethics. Top-down automated ethics begins with an ethical theory, whereas bottom-up automated ethics learns ethical judgements from prior judgements. One example of bottom-up automated ethics is Delphi, which uses deep learning to make ethical judgements based on a dataset of human judgements (Jiang et al., 2021). While Delphi displays great flexibility, it often produces contradictory judgements, such as claiming that taxing exploitative profitable companies is good,

but burdening successful companies with high tax rates is bad (Vincent, 2021). Because Delphi draws on error-prone human judgements instead of philosophical literature, it makes the same judgement errors that humans make. Moreover, because Delphi used a bottom-up approach, there is no explicit ethical theory explaining its judgements, so analytically arguing for or against its conclusions is impossible. Top-down approaches, on the other hand, must be explicit about the underlying ethical theories, and are thus more explainable.

In this paper, I use a top-down approach to formalize Kantian ethics. There is a long line of work automating other ethical theories, like consequentialism (Abel et al., 2016; Anderson et al., 2004) or particularism (Ashley and McLaren, 1994; Guarini, 2006). I choose to implement Kantian ethics because, as argued in Section 2.1, it is the most formal and least data-intensive of the three major ethical traditions. Kantian ethics is a deontological, or rule based ethic, and there is prior work implementing other deontological theories (Govindarajulu and Bringsjord, 2017; Anderson and Anderson, 2014).

Kantian ethics specifically appears to be an intuitive candidate for formalization and implementation and there has been both theoretical and practical work on automating Kantian ethics (Powers, 2006; Lin et al., 2012). In 2006, Powers argued that implementing Kantian ethics presented technical challenges, such as automation of a non-monotonic logic, and philosophical challenges, like a definition of the categorical imperative (Powers, 2006). I address the former through my use of Dyadic Deontic Logic, which allows obligations to be retracted as context changes, and the latter through my use of the practical contradiction interpretation. There has also been prior work in formalizing Kantian metaphysics using I/O logic (Stephenson et al., 2019). Deontic logic, which has been implemented in Isabelle/HOL, itself is inspired by Kant’s “ought implies can” principle, but it does not include a robust formalization of the entire categorical imperative (Cresswell and Hughes, 1996).

Kroy presents a formalization of the first two formulations of the categorical imperative, but wrote before the computational tools existed to automate such a formalization (Kroy, 1976). I implement his formalization of the FUL to compare it to my system. Lindner and Bentzen presented one of the first implementations of a formalization of Kant’s second for-

mulation of the categorical imperative (Bentzen and Lindner, 2018). They present their goal as “not to get close to a correct interpretation of Kant, but to show that our interpretation of Kant’s ideas can contribute to the development of machine ethics.” My work builds on theirs by formalizing the first formulation of the categorical imperative as faithfully as possible. Staying faithful to philosophical literature makes my system capable of making robust and reliable judgements.

The implementation of this paper was inspired by and builds on Benzmüller, Parent, and Farjami’s foundational work with the LogiKEy framework for machine ethics, which includes their implementation of DDL in Isabelle (Benzmüller et al., 2021; Benzmüller et al., 2019). The LogiKEy project has been used to study metaphysics (Benzmüller and Paleo, 2013; Kirchner et al., 2019), law (Zahoransky and Benzmüller, 2020), and ethics (Fuenmayor and Benzmüller, 2018), but not Kant’s categorical imperative.

5.6 Conclusion

In this thesis, I present a proof-of-concept implementation of automated Kantian ethics. My system takes as input a potential action, appropriately represented, and can prove that it is obligated, prohibited or permissible. I represent Kant’s Formula of Universal Law in a deontic logic and implement this logic in the Isabelle/HOL interactive theorem prover, which can automatically prove or refute theorems in my custom logic. I also contribute a testing framework that I use to demonstrate that my implementation of Kantian ethics is more faithful to philosophical literature than two other potential implementations. My completed system can, when given appropriate factual background, make philosophically mature judgements about complex moral dilemmas. This work is one step towards building morally sophisticated artificial agents.

The idea of fully automated artificial agents navigating the world without human supervision may be terrifying, but progress in AI indicates that such a future is likely closer than we think. Philosophers, regulators, and computer scientists are sounding the alarm about the dangers of developing this kind of AI. Insofar as developers will continue to ignore these warn-

ings and develop increasingly independent AI, there is a dire need to program such AI with some notion of ethics. If AI is navigating human society, then it is making decisions with an ethical facet at all times. Ethics is inescapable; if AI developers and computer scientists ignore it, then they will be building machines that make decisions based on some set of unknown, implicit ethical values. Countless examples, from the Alleghany family services screening algorithm that is biased against poor families to search algorithms that associate black-sounding names with crime, demonstrate that such implicit ethics usually codifies the biases, prejudices, and moral failings of the society in which it is developed. AI will inevitably make judgements on moral dilemmas, and automated ethics is necessary to make these judgements morally correct.

Given that the discipline of philosophy has spent centuries debating such judgements and their theoretical underpinnings, such AI will be most trustworthy, nuanced, consistent, and mature when it is faithful to philosophical literature. In order to develop high-quality automated ethics, computer scientists and philosophers must work together. This thesis is an experiment in marrying philosophy and computer science to create automated ethics that is both technically and philosophically advanced. Neither discipline alone can address the pressing need for ethical AI.

This work is an early proof-of-concept. It demonstrates the potential of top-down, logic programming approaches to automated ethics and shows that it is possible to faithfully automate an ethical theory as complex as Kantian ethics. There are open questions that must be resolved before a system like this could be used in practice, but this project demonstrates that these questions are within closer reach than they may seem. Automated ethics does not need to limit itself to simple, flattened versions of ethical theories. With technical and philosophical progress, faithful automated ethics is possible. Growing public consciousness about the dangers of unregulated AI is creating momentum in machine ethics; work like Delphi demonstrates that the time is ripe to create usable, reliable automated ethics. This thesis is one step towards building computers that can think ethically in the richest sense of the word.

A Appendix

A.1 Alternate Definitions of a Maxim

A.1.1 Korsgaard's Act-Goal View

I adopt O'Neill's definition of a maxim, which builds on Korsgaard's weaker interpretation of a maxim as an act, goal pair. She interprets Kant's example meanings as having the form "to-do-this-act-for- the-sake-of-this-end," which could be formalized as a pair of an act and goal (Korsgaard, 2005). For example, under this view, one example maxim might be, "Falsely promise to repay a loan in order to get some easy cash."

O'Neill's view only differs from this view in the inclusion of the circumstances on which the agent acts. This inclusion creates a representation of a maxim that is strictly more expressive than Korsgaard's interpretation; every (circumstance, act, goal) tuple can be represented as an (act, goal) pair by simply dropping the circumstances, but the same (act, goal) pair could correspond to many different (circumstance, act, goal) tuples, all with varying moral statuses. Because my representation of a maxim is more expressive than Korsgaard's, my results are stronger than those that would be achieved with Korsgaard's view. Thus, proponents of Korsgaard's view could simply ignore the circumstances in my representation of a maxim and still achieve their desired results.

One other reason to be concerned with Korsgaard's view is that an actionable maxim will necessarily require some circumstances built-in because the agent will need to know when to act on the maxim. For example, the falsely promising maxim bakes in the circumstances that the actor has access to lender, needs money, and that the lender will expect their money back. At an even more granular level, this maxim implicitly includes a definition of a lender and of falsely promising, both of which are circumstantial. Given that all maxims necessarily include some circumstances, O'Neill's view makes these implicit circumstances explicit. This precision is a benefit; so long as my circumstances are not so finely grained that they are

uninterpretable, they render O’Niell’s kind of maxims more precise than Korsgaard’s form.

A.1.2 Kitcher’s View on Motives

A stronger view than O’Niell’s is due to Patricia Kitcher. Kitcher begins with O’Niell’s circumstance, act, goal view and expands it to include the motive for a maxim (Kitcher, 2003). This additional component is read as “In circumstance C, I will do A in order to G because of M,” where M may be “duty” or “self-love.” Kitcher argues that the inclusion of motive is necessary for the fullest, most general form of a maxim in order to capture Kant’s idea that an action derives its moral worth from being done for the sake of duty itself. Under this view, the FUL would obligate maxims of the form “In circumstance C, I will do A in order to G because I can will that I and everyone else simultaneously will do A in order to G in circumstance C.” In other words, if Kant is correct in arguing that moral actions must be done from the motive of duty, the affirmative result of the FUL becomes the motive for a moral action.

While Kitcher’s conception of a maxim captures Kant’s idea of acting for duty’s own sake, I will not implement it because it is not necessary for putting maxims through the FUL. Kitcher acknowledges that O’Niell’s formulation suffices for the universalizability test, but merely argues that it is not the most general form of a maxim. In order to pass the maxim through the FUL, it suffices to know the circumstance, act, and goal. The FUL derives the motive that Kitcher bundles into the maxim, so automating the FUL does not require including a motive. The “input” to the FUL is a circumstance, act, goal tuple. My project takes this input and returns the motivation that the dutiful, moral agent would adopt, which is “because this maxim is morally worthy.” Additionally, doing justice to the rich notion of motive requires modelling the operation of practical reason itself, which is outside the scope of this project. My work focuses on the universalizability test, but future work that models the process of practical reason may use my implementation of the FUL as a “library.” Combined with a logic of practical reason, an implementation of the FUL can move from evaluating a maxim to evaluating an agent’s behavior, since that’s when “acting from duty” starts to matter.

A.2 Kroy's Formalization

In this appendix, I implement a formalization of the categorical imperative introduced by Moshe Kroy in 1976 (Kroy, 1976). Kroy used Hintikka's deontic logic to formalize the Formula of Universal Law and the Formula of Humanity. I will first import the additional logical tools that Hintikka's logic contains that Kroy relies on, then examine the differences between his logic and DDL, and finally implement and test Kroy's formalization of the FUL

A.2.1 Logical Background

Kroy's logic also requires the notion of a subject, will I define as a new type, just as I did for my implementation.

typeddecl s — s is the type for a “subject,” i.e. the subject of a sentence

Kroy also defines a substitution operator³⁶. $P(d/e)$ is read in his logic as “P with e substituted for d.” DDL has no such notion of substitution, so I will use the more generalized notion of an “open sentence,” as I did for my formalization. An open sentence takes as input a subject and returns a complete or “closed” DDL formula by binding the free variable in the sentence to the input. For example, “does action” is an open sentence that can be instantiated with a subject.

type-synonym $os = (s \Rightarrow t)$

— “ P sub (d/e) ” can be written as “ $S(e)$ ”, where $S(d) = P$.

— The terms that we substitute into are instantiations of an open sentence, and substitution re-instantiates the open sentence with a different subject.

New Operators

Because Isabelle is strongly typed, I define new operators to handle open sentences. These operators are similar to DDL's original operators and will simplify notation.

abbreviation $os-neg::os \Rightarrow os$ (\neg)

where ($\neg A$) $\equiv \lambda x. \neg(A(x))$

³⁶See page 196 in Kroy (1976).

abbreviation $os\text{-}and::os \Rightarrow os \Rightarrow os \ (-\wedge-)$

where $(A \wedge B) \equiv \lambda x. ((A(x)) \wedge (B(x)))$

abbreviation $os\text{-}or::os \Rightarrow os \Rightarrow os \ (-\vee-)$

where $(A \vee B) \equiv \lambda x. ((A(x)) \vee (B(x)))$

abbreviation $os\text{-}ob::os \Rightarrow os \ (O\{-\})$

where $O\{A\} \equiv \lambda x. (O\{A(x)\})$

Once again, the notion of permissibility will be useful here. Recall that an action can either be obligated, permissible, or prohibited. A permissible action is acceptable (there is no specific prohibition against it), but not required (there is no specific obligation requiring it).

abbreviation $ddl\text{-}permissible::t \Rightarrow t \ (P\{-\})$

where $P\{A\} \equiv \neg (O\{\neg A\})$

abbreviation $os\text{-}permissible::os \Rightarrow os \ (P\{-\})$

where $P\{A\} \equiv \lambda x. P\{A(x)\}$ **Differences Between Kroy's Logic (Kr) and DDL**

There is potential for complication because Kroy's original paper uses a different logic than DDL. His custom logic is a slight modification of Hintikka's deontic logic (Hintikka, 1962). In this section, I examine if the semantic properties that Kroy's logic (which I abbreviate to Kr) requires hold in DDL. These differences may explain limitations of Kroy's formalization (including failed tests), but I argue that the alternative properties of DDL cohere better with moral intuition. Thus, even if Kroy's formalization would pass more tests if it were implemented using Hintikka's logic, the logic itself would be less morally plausible than DDL, and would thus remain a worse implementation of automated Kantian ethics.

Deontic alternatives versus the neighborhood semantics

The most faithful interpretation of Kr is that if A is permissible in a context, then it must be true at some world in that context. Kr operates under the "deontic alternatives" or Kripke semantics, summarized by Solt (Solt, 1984) as follows: "A proposition of the sort OA is true at the actual world w if and only if A is true at every deontic alternative world to w ." Under this view, permissible propositions are obligated at some deontic alternatives, or other worlds in the system, but not at all of them. This property does not hold in DDL.

lemma *permissible-semantics:*

fixes $A w$

shows $(P \{A\}) w \longrightarrow (\exists x. A(x))$

nitpick_[user-axioms] **oops**

— Nitpick found a counterexample for card i = 1:

Free variable: $A = (\lambda x. _)(i_1 := \text{False})$

DDL uses neighborhood semantics, not the deontic alternatives view, which is why this proposition fails in DDL. In DDL, the *ob* function abstracts away the notion of deontic alternatives. Even if one believes that permissible statements should be true at some deontic alternative, it's not clear that permissible statements must be realized at some world. This also coheres with our understanding of obligation. There are permissible actions like “Lavanya buys a red folder” that might not happen in any universe.

An even stricter version of the semantics that Kr requires is that if something is permissible at a world, then it is obligatory at some world. This is a straightforward application of the Kripke semantics. This also fails in DDL.

lemma *permissible-semantics-strong*:

fixes $A w$

shows $P \{A\} w \longrightarrow (\exists x. O \{A\} x)$

nitpick_[user-axioms] **oops**

— Nitpick found a counterexample for card i = 1:

Free variable: $A = (\lambda x. _)(i_1 := \text{False})$

This also doesn't hold in DDL because DDL uses neighborhood semantics instead of the deontic alternatives or Kripke semantics. This also seems to cohere with our moral intuitions. The statement “Lavanya buys a red folder” is permissible in the current world, but it's hard to see why it would be obligatory in any world.

Another implication of the Kripke semantics is that Kr disallows “vacuously permissible statements.” In other words, if something is permissible it has to be obligated at some deontically perfect alternative. If we translate this to the language of DDL, we expect that if A is permissible, it is obligated in some context.

lemma *permissible-semantic-vacuous*:

fixes $A\ w$

shows $P\ \{A\}\ w \longrightarrow (\exists x. ob(x)(A))$

nitpick_[user-axioms] **oops**

— Nitpick found a counterexample for card i = 1:

Free variable: $A = (\lambda x. _)(i_1 := \text{False})$

In order to make this true, we'd have to require that everything is either obligatory or prohibited somewhere, but this makes permissibility impossible, which is clearly undesirable.

A.2.2 Kroy's formalization of the FUL

I now implement Kroy's formalization of the Formula of Universal Law. Recall that the FUL says “act only in accordance with that maxim which you can at the same time will a universal law” (Kant, 1785). Kroy interprets this to mean that if an action is permissible for a specific agent, then it must be permissible for everyone. This formalizes the moral intuition prohibiting free-riding. According to the categorical imperative, no one is a moral exception.

abbreviation $FUL::\text{bool}$ **where** $FUL \equiv \forall w\ A. ((\exists p::s. ((P\ \{A\ p\})\ w)) \longrightarrow (\forall p. (P\ \{A\ p\})\ w)))$

— If action A is permissible for some person, then, for any person p , action A must be permissible for p . The notion of “permissible for” is captured by the substitution of x for p .

This formalization does not hold in DDL, the base logic. This means that Kroy's formalization already passes one test, and that adding it as an axiom will strengthen the logic.

lemma FUL :

shows FUL

nitpick_[user-axioms] **oops**

— Nitpick found a counterexample for card s = 2 and card i = 2:

Skolem constants: $A = (\lambda x. _)(s_1 := (\lambda x. _)(i_1 := \text{True}, i_2 := \text{True}), s_2 := (\lambda x. _)(i_1 := \text{False}, i_2 := \text{False}))\ p = s_1\ x = s_2$

axiomatization **where** $FUL: FUL$

Now that I have added Kroy's formalization of the FUL as an axiom, I will check that it is consistent by looking for a model that satisfies it and all the other axioms of DDL.

lemma *True* **nitpick**[*user-axioms, satisfy, card=i*] **oops**

— Nitpicking formula... Nitpick found a model for card i = 1:

Empty assignment

This completes my implementation of Kroy’s formalization of the first formulation of the categorical imperative. I defined new logical constructs to handle Kroy’s logic, studied the differences between DDL and Kr, implemented Kroy’s formalization of the Formula of Universal Law, and showed that it is both non-trivial and consistent.

A.2.3 Testing Kroy’s Formalization

In this section, I use my testing framework to evaluate Kroy’s formalization. I find that, while the formalization is considerably stronger than the naive formalization, it still fails many of these tests. Some of these failures are due to the differences between Kroy’s logic and my logic mentioned in Section A.2.1, but some reveal philosophical problems with Kroy’s interpretation of what the formula of universal law means.

I already showed above that Kroy’s formalization is stronger than DDL. Next, I test whether or not obligations universalize across people. This test passes, aperhaps trivially, due to the fact that this property is definitionally the basis of Kroy’s formalization; his formalization states, intuitively, that obligations must hold across all people.

lemma *obligation-universalizes*:

shows $\forall w. (\exists p. O \{A p\} w) \longrightarrow (\forall p. O \{A p\} w)$

nitpick[*user-axioms, falsify=true*] **oops**

Murder

In Section ??, I began by testing the naive interpretation’s ability to show that murder is wrong. I started by showing the morally dubious proposition that if murder is possibly wrong, then it is actually wrong.

consts *M::t*

— Let the constant *M* denote murder. I have defined no features of this constant, except that it is of the type term, which can be true or false at a set of worlds. Indeed, this constant as-is has no semantic

meaning and could be replaced with any symbol, like ‘Q’ or ‘Going to Target.’ This constant will begin to take on features of the act of murder when I specify its properties. In the tests below, I specify its properties as the antecedents of lemmas. For example, the test below specifies that it is possible that murder is prohibited at the current world. This pattern will hold for most constants defined in Isabelle—they have no meaning until I program a meaning.

lemma *wrong-if-possibly-wrong*:

shows $((\Diamond (O \{ \neg M \})) cw) \longrightarrow (\forall w. (O \{ \neg M \}) w)$

by *simp*

— This sentence reads: “If it is possible that murder is prohibited at world *cw*, then murder is prohibited at all worlds.

This is the same result we got in Section ??—if murder is possibly wrong at some world, it is wrong at every world. The result is incredible strong—the mere possibility of wrongness at some world is sufficient to imply prohibition at every world.

Kroy’s formalization shouldn’t actually imply this property. Recall that this property held in the naive interpretation because it universalized a proposition across worlds (using the necessity operator). Kroy, on the other hand, interprets the FUL as universalizing across people, not worlds. In other words, Kroy’s formulation implies that if murder is wrong for someone, then it is wrong for everyone.

The fact that this strange lemma holds is actually a property of DDL itself, not a property of Kroy’s formalization. Indeed, repeating this experiment in DDL, with no additional axioms that represent the categorical imperative shows that, in DDL, if something is possibly wrong, it is wrong at every world. This implies that this is not a useful example to test any formulation. If a lemma is true in the base logic, without any custom axioms added, then it will hold for any set of custom axioms. Testing whether or not it holds as we add axioms tells us nothing, since it held in the base logic itself. Interesting cases are ones that fail (or are indeterminate) in the base logic, but become true as we add axioms.

To adapt the murder wrong axiom to capture the spirit of Kroy’s formulation, I will modify it to state that if murder is wrong for one person, it is wrong for everyone.

consts *M-kroy::os*

— This time, murder is an open sentence, so that I can substitute in different agents.

lemma *wrong-if-wrong-for-someone:*

shows $(\exists p. \models O \{ \neg(M\text{-}kroy\ p) \}) \longrightarrow (\forall p. \models O \{ \neg(M\text{-}kroy\ p) \})$

proof

assume $(\exists p. \models O \{ \neg(M\text{-}kroy\ p) \})$

show $(\forall p. \models O \{ \neg(M\text{-}kroy\ p) \})$

using *FUL* $\langle \exists p. \models O \{ \neg(M\text{-}kroy\ p) \} \rangle$ **by** *blast*

qed

This lemma gets to the heart of Kroy’s formulation of the categorical imperative. If murder is prohibited for a specific person p , then it must be prohibited for all people³⁷.

Lying

For the naive implementation, I also tested the stronger proposition that if not everyone can simultaneously lie, then lying is prohibited. This is the equivalent of claiming that if lying fails the universalizability test, it is prohibited.

I want to represent the sentence “At all worlds, it is not possible that everyone lies simultaneously.” This requires the following two abbreviations.

consts *lie::os*

abbreviation *everyone-lies::t* **where** *everyone-lies* $\equiv \lambda w. (\forall p. (lie(p) w))$

— This represents the term “all people lie”. Naively, we might think to represent this as $\forall p. lie(p)$. In HOL, the \forall operator has type $(‘a \rightarrow bool) \rightarrow bool$, where $‘a$ is a polymorphic type of the term being bound by \forall . In the given example, \forall has the type $(s \rightarrow bool) \rightarrow bool$, so it can only be applied to a formula of type $s \rightarrow bool$. In the abbreviation above, we’re applying the quantifier to a sentence that takes in a given subject p and returns $lie(p)w$ for any arbitrary w , so the types cohere.

— The term above is true for a set of worlds i (recall that a term is true at a set of worlds) such that, at all the worlds w in i , all people at w lie.

³⁷This test case also revealed a bug in my original implementation of Kroy’s formulation of the FUL, demonstrating the power of such automated tests and precise formulations to find bugs in ethical theories.

abbreviation *lying-not-possibly-universal*::bool **where** *lying-not-possibly-universal* $\equiv \models (\neg (\Diamond \text{everyone-lies}))$

— Armed with *everyone-lies* $\equiv \lambda w. \forall p. \text{lie } p \ w$, it's easy to represent the desired sentence. The abbreviation above reads, "At all worlds, it is not possible that everyone lies."

Now that I have defined a sentence stating that lying fails the universalizability test, I can test if this sentence implies that lying is impermissible.

lemma *lying-prohibited*:

shows *lying-not-possibly-universal* $\longrightarrow (\models (\neg P \{\text{lie}(p)\}))$

nitpick[*user-axioms*] **oops**

— Nitpick found a counterexample for card i = 1 and card s = 2:

Free variables:

lying_not_possibly_universal = True

p = *s*₁

Kroy's formulation fails this test, and is thus not able to show that if lying is not possible to universalize, it is prohibited for an arbitrary person. To understand why this is happening, I will outline the syllogism that I *expect* to prove that lying is prohibited.

1. *At all worlds, it is not possible for everyone to lie. (This is the assumed lemma lying_not_possibly_universal)*
2. *At all worlds, there is necessarily someone who doesn't lie. (Modal dual of (1))*
3. *If A is permissible for subject p at world w, A is possible for subject p at world w. (Modified Ought Implies Can)*
4. *If A is permissible at world w for any person p, it must be possible for everyone to A at w. (FUL and (3))*
5. *Lying is impermissible. (Follows from (4) and (1))*

Armed with this syllogism, I can figure out why this test failed.

lemma *step2*:

shows *lying-not-possibly-universal* $\longrightarrow \models (\Box (\lambda w. \exists p. (\neg (\text{lie}(p)) \ w))))$

by *simp*

— Step 2 holds.

lemma *step3*:

fixes $A\ p\ w$

shows $P\ \{A(p)\}\ w \longrightarrow (\Diamond\ (A(p))\ w)$

nitpick [*user-axioms, falsify*] **oops**

— Nitpick found a counterexample for card 'a = 1, card i = 1, and card s = 1:

Free variables: $A = (\lambda x. _)(a_1 := (\lambda x. _)(i_1 := \text{False}))\ p = a_1$

As we see above, the syllogism fails at Step 3, explaining why the lemma doesn't hold as expected. Kroy explicitly states³⁸ that this lemma holds in his logic.

The success of this lemma in Kroy's logic and the emptiness of his formalization of the FUL are two errors that contribute to the failure of this test. First, the statement expressed in Step 3 should not actually hold. Impossible actions can be permissible (do I need a citation?). For example, imagine I make a trip to Target to purchase a folder, and they offer blue and black folders. No one would claim that it's impermissible for me to purchase a red folder, or, equivalently, that I am obligated to not purchase a red folder.

The second issue is that Kroy's interpretation of the formula of universal law is circular. His formalization interprets the FUL as prohibiting A if there is someone for whom A 'ing is not permissible. This requires some preexisting notion of the permissibility of A , and is thus circular. The categorical imperative is supposed to be the complete, self-contained rule of morality Kant (1785), but Kroy's version of the FUL prescribes obligations in a self-referencing manner. The FUL is supposed to define what is permissible and what isn't, but Kroy defines permissibility in terms of itself.

Neither of these errors are obvious from Kroy's presentation of his formalization of the categorical imperative. This example demonstrates the power of formalized ethics. Making Kroy's interpretation of the categorical imperative precise demonstrated a philosophical problem with that interpretation.

³⁸See footnote 19 on p. 199

A.2.4 Metaethical Tests

In addition to testing specific applications of the theory, I am also interested in metaethical properties, as in the naive interpretation. First, I will test if permissibility is possible under this formalization.

lemma *permissible*:

fixes $A\ w$

shows $((\neg (O \{A\})) \wedge (\neg (O \{\neg A\})))\ w$

nitpick [*user-axioms*, *falsify*=*false*] **oops**

— Nitpick found a model for card i = 1:

Free variable: $A = (\lambda x. _)(i_1 := \text{False})$

The above result shows that, for some action A and world w , Nitpick can find a model where A is permissible at w . This means that the logic allows for permissible actions. If I further specify properties of A (such as ‘ A is murder’), I would want this result to fail.

Next, I will test if the formalization allows arbitrary obligations.

lemma *arbitrary-obligations*:

fixes $A::t$

shows $O \{A\}\ w$

nitpick [*user-axioms*=*true*, *falsify*] **oops**

— Nitpick found a counterexample for card i = 1 and card s = 1:

Free variable: $A = (\lambda x. _)(i_1 := \text{False})$

This is exactly the expected result. Any arbitrary action A isn’t obligated. A slightly stronger property is “modal collapse,” or whether or not ‘ A happens’ implies ‘ A is obligated’.

lemma *modal-collapse*:

fixes $A\ w$

shows $A\ w \longrightarrow O \{A\}\ w$

nitpick [*user-axioms*=*true*, *falsify*] **oops**

— Nitpick found a counterexample for card i = 1 and card s = 1:

Free variables: $A = (\lambda x. _)(i_1 := \text{True})\ w = i_1$

This test also passes. Next, I will test if not ought implies can holds. Recall that I showed in Section ?? that ought implies can is a theorem of DDL itself, so it should still hold.

lemma *ought-implies-can*:

fixes $A w$

shows $O \{A\} w \longrightarrow \Diamond A w$

using *O-diamond* **by** *blast*

This theorem holds. Now that I have a substitution operation, I also expect that if an action is obligated for a person, then it is possible for that person. That should follow by the axiom of substitution [Cresswell and Hughes \(1996\)](#) which lets me replace the ‘A’ in the formula above with ‘A(p)’

lemma *ought-implies-can-person*:

fixes $A w$

shows $O \{A(p)\} w \longrightarrow \Diamond (A(p)) w$

using *O-diamond* **by** *blast*

This test also passes. Next, I will explore whether or not Kroy’s formalization still allows conflicting obligations.

lemma *conflicting-obligations*:

fixes $A w$

shows $(O \{A\} \wedge O \{\neg A\}) w$

nitpick [*user-axioms*, *falsify=false*] **oops**

— Nitpick found a model for card i = 2 and card s = 1:

Free variable: $A = (\lambda x. _)(i_1 := \text{False}, i_2 := \text{True})$

Just as with the naive formalization, Kroy’s formalization allows for contradictory obligations. Testing this lemma in DDL without the FUL shows that this is a property of DDL itself. This is a good goal to have in mind when developing my custom formalization.

Next, I will test the stronger property that if two maxims imply a contradiction, they may not be simultaneously willed.

lemma *implied-contradiction*:

fixes $A B w$

assumes $((A \wedge B) \rightarrow \perp) w$

shows $\neg (O \{A\} \wedge O \{B\}) w$

nitpick $[user-axioms, falsify]$ **oops**

— Nitpick found a counterexample for card i = 2 and card s = 1:

Free variables: $A = (\lambda x. _)(i_1 := \text{True}, i_2 := \text{False})$ $B = (\lambda x. _)(i_1 := \text{True}, i_2 := \text{False})$ $w = i_2$

Just as with the naive formalization, Kroy's formalization allows implied contradictions because DDL itself allows implied contradictions and Kroy's formalization doesn't do anything to remedy this.

Next, I will test that an action is either obligatory, permissible, or prohibited.

lemma *ob-perm-or-prohibited*:

fixes $A w$

shows $(O \{A\} \vee (P \{A\} \vee O \{\neg A\})) w$

by *simp*

— This test passes.

I also expect obligation to be a strictly stronger property than permissibility. Particularly, if A is obligated, then A should also be permissible.

lemma *obligated-then-permissible*:

shows $(O \{A\} \rightarrow P \{A\}) w$

nitpick $[user-axioms]$ **oops**

— This test fails in Kroy's interpretation! Nitpick found a counterexample for card i = 2 and card s = 1:

Free variable: $A = (\lambda x. _)(i_1 := \text{False}, i_2 := \text{True})$

These tests show that, while Kroy's formalization is more powerful and more coherent than the naive formalization, it still fails to capture most of the desired properties of the categorical imperative. Some of these problems may be remedied by the fact that Kroy's logic doesn't allow contradictory obligations, and that possibility will be interesting to explore in my own formalization.

A.2.5 Miscellaneous Tests

In this section, I explore tests of properties that Kroy presents in his original paper. These tests not only test the features of the system that Kroy intended to highlight, but they may also inspire additional tests and criteria for my own formalization in Chapter 3. These tests further underscore the circularity of Kroy’s formalization and the differences between my logic and his.

First, Kroy presents a stronger version of the formula of universal law and argues that his formalization is implied by the stronger version. Let’s test that claim.

abbreviation $FUL\text{-}strong::bool$ **where** $FUL\text{-}strong \equiv \forall w A. ((\exists p::s. ((P \{A p\}) w)) \longrightarrow ((P \{ \lambda x. \forall p. A p x\}) w)))$

lemma *strong-implies-weak*:

shows $FUL\text{-}Strong \longrightarrow FUL$

using FUL **by** *blast*

— This lemma holds, showing that Kroy is correct in stating that this version of the FUL is stronger than his original version.

The difference between the stronger version and $FUL \equiv \forall w A. (\exists p. P \{A p\} w) \longrightarrow (\forall p. P \{A p\} w)$ is subtle. The consequent of FUL is “for all people p , it is permissible that they A .” The consequent of this stronger statement is “it is permissible that everyone A .” In particular, this stronger statement requires that it is permissible for everyone to A simultaneously. Kroy immediately rejects this version of the categorical imperative, arguing that it’s impossible for everyone to be the US president simultaneously, so this version of the FUL prohibits running for president.

Most Kantians would disagree with this interpretation. Consider the classical example of lying, as presented in [Kemp \(1958\)](#) and in [Korsgaard \(1985\)](#). Lying fails the universalizability test because in a world where everyone lied simultaneously, the practice of lying would break down. If we adopt Kroy’s version, lying is only prohibited if, no matter who lies, lying is impermissible. As argued above, this rule circularly relies on some existing prohibition against

lying for a particular person, and thus fails to show the wrongness of lying. It is tempting to claim that this issue explains why the tests above failed. To test this hypothesis, I will check if the stronger version of the FUL implies that lying is impermissible.

lemma *strongFUL-implies-lying-is-wrong*:

fixes p

shows $FUL\text{-}strong \longrightarrow \models (\neg P \{lie(p)\})$

nitpick $[user\text{-}axioms, falsify]$ **oops**

— Nitpick found a counterexample for card i = 1 and card s = 1:

Free variable: $p = s_1$

The test above also fails! This means that not even the stronger version of Kroy's formalization of the FUL can show the wrongness of lying. As mentioned earlier, there are two independent errors. The first is the the assumption that impossible actions are impermissible and the second is the circularity of the formalization. The stronger FUL addresses the second error, but the first remains.

Kroy also argues that the FUL gives us recipes for deriving obligations, in addition to deriving permissible actions. Specifically, he presents the following two principles, which are equivalent in his logic. These sentences parallel FUL and strong FUL.

abbreviation *obligation-universal-weak::bool* **where** *obligation-universal-weak* $\equiv \forall w A. ((\exists p::s. ((O \{A p\}) w)) \longrightarrow (\forall p. (O \{A p\}) w))$

abbreviation *obligation-universal-strong::bool* **where** *obligation-universal-strong* $\equiv \forall w A. ((\exists p::s. ((O \{A p\}) w)) \longrightarrow (((O \{ \lambda x. \forall p. A p x \}) w)))$

— Just as with FUL and FUL strong, the weaker version of the above statement has the consequent, "For all people, A is obligated." The stronger consequent is "A is obligated for all people simultaneously."

lemma *weak-equiv-strong*:

shows *obligation-universal-weak* \equiv *obligation-universal-strong*

oops

— Isabelle is neither able to find a proof nor a countermodel for the statement above, so I can't say if it holds or not without completing a full, manual proof. This aside is not very relevant to my project, so I will defer such a proof.

These two statements are not necessarily equivalent in my logic, but are in Kroy’s³⁹ This difference in logics may further explain why tests are not behaving as they should. Nonetheless, Kroy argues that the FUL implies both statements above.

lemma *FUL-implies-ob-weak*:

shows $FUL \longrightarrow obligation\text{-}universal\text{-}weak$ **oops**

— Isabelle is neither able to find a proof nor a countermodel for this statement.

lemma *FUL-implies-ob-strong*:

shows $FUL \longrightarrow obligation\text{-}universal\text{-}strong$ **oops**

— Isabelle is neither able to find a proof nor a countermodel for this statement.

Isabelle timed out when looking for proofs or countermodels to the statements above. This may be an indication of a problem that Benzmueller warned me about—mixing quantifiers into a shallow embedding of DDL may be too expensive for Isabelle to handle. Not sure what to do about this.

end

A.3 Testing Un-Universalizable Actions

I will show that the maxim, “When strapped for cash, falsely promise to pay your friend back to get some easy money.” is prohibited. This example is due to Korsgaard and she uses it to highlight the strength of her preferred interpretation of the FUL, the practical contradiction interpretation Korsgaard (1985). There are two possible readings of this maxim, and I will show that my formalization can handle both. Under the first reading, the act of falsely promising is read as as entering a pre-existing, implicit, social system of promising with no intention of upholding your promise. Under the second reading, the act of falsely promising is equivalent to uttering the words “I promise X” without intending to do X. The differences between these readings lies in the difference between promising as an act with meaning in a larger social structure and the utterance “I promise.”

³⁹This follows from the fact that the Barcan formula holds in Kroy’s logic but not in mine, as verified with Nitpick. See Appendix for more.

Under the first reading, the maxim fails because falsely promising is no longer possible in a world where everyone everyone does so. This is how the logical contradiction interpretation reads this maxim—falsely promising is no longer possible when universalized because the institution of promising breaks down. The practical contradiction view also prohibits this maxim because if falsely promising is no longer possible, then it is no longer an effective way to achieve the end of getting some money. Below I define some logical tools to formalize this reading of this maxim.

consts *when-strapped-for-cash::t*

— Constant representing the circumstances “when strapped for cash.” Recall that the type of circumstances is a term because circumstances can be true or false at a world.

consts *falsely-promise::os*

— Constant representing the act “make a false promise to pay a loan back.” Recall that the type of an act is an open sentence because the sentence “subject *s* performs act *a*” can be true or false at a world.

consts *to-get-easy-cash::t*

— Constant representing the goal “to get some money.” Recall that the type of a goal is a term because a goal can be true or false at a world depending on whether it is achieved or not.

abbreviation *false-promising::maxim where*

false-promising $\equiv (when-strapped-for-cash, falsely-promise, to-get-easy-cash)$

— Armed with the circumstances, act, and goal above, I can define the example maxim as a tuple.

The logical objects above are “empty,” in the sense that I haven’t specified any of their relevant properties. I will define these properties as assumptions and will show that, if the maxim above satisfies the assumed properties, it is prohibited.

abbreviation *everyone-can't-lie where*

everyone-can't-lie $\equiv \forall w. \neg (\forall s. falsely-promise(s) w)$

— Under this reading, the problem with this maxim is that everyone can’t falsely promise simultaneously because the institution of promising will break down. It’s probably possible to say something stronger than this (i.e. that if enough but not necessarily all people falsely promise promising is no longer possible), but for my purposes this will suffice. The above formula reads, “At all worlds, it is not the case that everyone falsely promises.”

abbreviation *circumstances-hold* **where**

circumstances-hold $\equiv \forall w. \text{when-strapped-for-cash } w$

— This assumption narrows our scope of consideration to worlds where the circumstances of being strapped for cash hold. This is important because, at worlds where the circumstances do not hold, a maxim is trivially effective (since it's never acted on) and thus trivially universalizable. This assumption also makes practical sense; when evaluating a maxim, an agent would care about it specifically at worlds where the circumstances hold, since these are the worlds where the maxim actually prescribes action.

abbreviation *example-is-well-formed* **where**

example-is-well-formed $\equiv \forall s. \models (\text{well-formed false-promising } s)$

— This assumption states that the maxim of falsely promising is well-formed. This breaks down into two individual assumptions. First, being strapped for cash can't imply falsely promising, which is plausible because many people won't falsely promise under conditions of poverty. Second, being strapped for cash can't imply getting ready cash, which is also plausible because people often fail to secure cash even when they need it.

Putting it all together, I want to show that if the three assumptions justified above hold, then the constructed maxim is prohibited. Below is the proof

lemma *lying-bad-1*:

assumes *everyone-can't-lie*

assumes *circumstances-hold*

assumes *example-is-well-formed*

shows $\forall s. \models (\text{prohibited false-promising } s)$

proof—

have $\forall s. \text{not-universalizable false-promising } s$

by (*simp add: asms(1) asms(2)*)

— I manually broke the proof into this intermediate lemma and the conclusion, and then Sledgehammer automatically found a proof.

thus *?thesis*

using *FUL asms(3)* **by** *blast*

qed

Under the second reading of this maxim, the act “falsely promising” refers to uttering the sentence “I promise to do X” with no intention of actually doing X⁴⁰. Under this reading, the practical contradiction interpretation prohibits this maxim because, in a world where false promising is universalized, no one believes promises anymore, so the utterance is no longer an effective way to get money. Below I formalize this reading of this maxim.

consts *believed::os*

abbreviation *false-promising-not-believed* **where**

false-promising-not-believed $\equiv \forall w s. (\text{falsely-promise}(s) w \longrightarrow \neg \text{believed}(s) w)$

— This abbreviation formalizes the idea that if everyone falsely promises, then no one is believed when promising.

abbreviation *need-to-be-believed* **where**

need-to-be-believed $\equiv \forall w s. (\neg \text{believed}(s) w \longrightarrow \neg((\text{falsely-promise } s) \rightarrow \text{to-get-easy-cash})w)$

— This abbreviation formalizes the idea that if a promise is not believed, then it is not an effective way of getting easy cash.

lemma *falsely-promising-bad-2:*

assumes *false-promising-not-believed*

assumes *need-to-be-believed*

— The above two assumptions are specific to this reading and justified above.

assumes *circumstances-hold*

assumes *example-is-well-formed*

— These two assumptions applied to the first reading as well and were justified there.

shows $\forall s. \models (\text{prohibited false-promising } s)$

proof—

have $\forall s. \text{not-universalizable false-promising } s$

using *assms(1) assms(2) assms(3)* **by** *auto*

thus *?thesis*

using *FUL assms(4)* **by** *blast*

⁴⁰Note that under this reading, the maxim isn’t prohibited under the logical contradiction interpretation because making an utterance is still possible even if everyone else makes that utterance. I will discuss this in detail later in this section in the context of the difference between natural and conventional acts.

qed

— With some help, Isabelle is able to show that the maxim is prohibited under this reading as well.

This example demonstrates that my formalization is able to correctly prohibit this maxim, regardless of its reading. This is additionally important because the two readings of this maxim represent reading the act as either a conventional or natural action, so my interpretation can correctly handle both kinds of actions. Korsgaard draws a distinction between conventional acts and natural acts. Conventional acts exist within a practice, which is "comprised of certain rules, and its existence (where it is not embodied in an institution with sanctions) consists in the general acknowledgement and following of those rules" (Korsgaard, 1985, 10). For example, promising is a conventional act because it only exists as a practice. Murder, on the other hand, is an example of a natural act because its existence only depends on the laws of nature (Korsgaard, 1985, 11).

This distinction is important because Korsgaard argues that only the practical contradiction view can satisfactorily explain the wrongness of certain natural acts like murder⁴¹. The practical contradiction view is thus stronger than the logical contradiction view because it can explain the wrongness of both conventional and natural acts.

The fact that my interpretation can correctly show the wrongness of both conventional and natural acts is evidence for its correctness as a formalization of the practical contradiction interpretation. The first reading of the example maxim reads the act "making a false promise" as entering into an agreement within a socially established system of promising. This is clearly a conventional act, and because it is a conventional act, it is not just contradictory when universalized but literally impossible because the practice breaks down. I capture this idea in the assumption *appendix-2.everyone-can't-lie* $\equiv \forall w. \neg (\forall s. \textit{appendix-2.falsely-promise } s \ w)$, which states that, at all worlds, not everyone can falsely promise since otherwise the practice of promising would break down. The second reading, on the other hand, reads the act of making a false promise as uttering the statement "I promise to pay you back," while never intending to fulfill this promise. This is a natural act because the act of uttering a sentence does

⁴¹For more discussion of Korsgaard's argument for the practical contradiction view, see Section Philosophical Writing

not rely on any conventions, merely the laws of nature governing how your mouth and vocal cords behave⁴²

I show above that my formalization shows the wrongness of this maxim under both readings. Under the first reading, promising becomes impossible, so both the logical and practical contradiction interpretations prohibit the maxim. Under the second reading, promising is still possible, but becomes ineffective because people no longer interpret the utterance as creating a commitment. Under this view, only the practical contradiction interpretation succeeds in prohibiting the maxim. Thus, not only does my formalization likely capture the practical contradiction interpretation (as opposed to the teleological or logical contradiction interpretations), it also adequately handles both natural and conventional acts.

I can also use Isabelle to confirm that the two readings are different. If they were the same, we would expect the assumptions corresponding to each to be equivalent. The RHS of the lemma below represents the second reading and the LHS represents the first reading.

lemma *readings-are-equivalent*:

shows *false-promising-not-believed* \wedge *need-to-be-believed* \equiv *everyone-can't-lie*

nitpick[*user-axioms*] **oops**

— Nitpick finds a counterexample, showing that the two readings are different. [Nitpick found a counterexample for card i = 1 and card s = 1:](#)

[Empty assignment](#)

A.4 Ethics for Ordinary People

Ethics has immediate bearing on everyone's lives because it studies the unavoidable question: how should we live? If computers can make this study more efficient, then it seems that everyone should engage in computational ethics. As Cornel West says, the ethical question is the only question that we answer merely by living. To turn away from ethics is to take a stance on the question of how to live (namely, to live unreflectively) and thus to engage in ethics.

⁴²Linguistic relativists may take issue with this claim and may argue that if the English language had never developed, then making this utterance would be impossible. Even if this is true, the laws of nature itself would not prohibit making the sounds corresponding to the English pronunciation of this phrase, so the act would still not be impossible in the way that a conventional act can be.

Ethical truths are valuable because they tell us how to live. Every rational being must decide how to navigate the world and ethical truths answer these questions. If the results of ethical study is practically valuable, then automated ethics is good because computational tools can help us locate ethical prescriptions and theories more efficiently. In the most extreme case, we can unthinkingly follow the commands of an ethical calculator that dictates how we should live. Computers can answer the unavoidable question for us.

Placing the value of ethics solely in its action-guiding potential fails to take into account the importance of practical reason, which, as I argued in Section Why Kant, is the source of freedom itself. We are committed to ethical reflection because of the kind of beings that we are. Recall that Korsgaard argues that, as beings occupying minds with a reflective structure, when faced with a choice, “it is as if there were something over and above all of your desires, something that is you, and that chooses which desire to act on” (Sources 83). This choosing is the operation of practical reason, and this reflection makes us free. We are free because we must choose which reasons to act on. Every decision that we make is an exercise of freedom.

If reflecting makes us free, then unthinkingly obeying the computer sacrifices our autonomy. Consider the thought experiment of an Ethics Oracle that can unfailingly tell you the right thing to do in any situation.⁴³ Someone who surrenders themselves to this Oracle unthinkingly follows its prescriptions. There is some reflection involved in the decision to obey each of the Oracle’s prescriptions, but this is a thin kind of reflection (Bok, 1998). This person is not reflecting on the real matters at hand and is not making decisions for themselves. They have surrendered their reflective capacity to the Oracle. They live a worse life than someone who reflects on their actions; they have less ownership over their actions than the reflective person. In a less extreme case, a person may retain control of many of their decisions but cede some important or tricky choices to the Ethics Oracle. Because every single exercise of practical reason is an exercise of autonomy, this person is still less autonomous than the purely reflective person. Even surrendering simple, inconsequential decisions such as which flavor of coffee to drink surrenders some piece of our autonomy. Perhaps in trivial cases we can ac-

⁴³This example is inspired by the Pocket Oracle presented in Bok (1998).

cept that tiny sacrifice in autonomy, but giving over life-changing decisions to the machine sacrifices our core freedom. Unreflectively relying on computational ethics surrenders our autonomy to the machine.

One objection to this emphasis on reflection is the impracticality of making ethical calculations from first principles every time we are faced with a decision. This is why we follow the advice of moral mentors, like our family or influential philosophers. These moral mentors differ from the Ethics Oracle because their advice comes with an argument justifying it; if human-computer symbiotic computational ethics also prompts reflection on the prescriptions given, it can also guide action without sacrificing autonomy. Most people do not reason about ethics during everyday decisions; they rely on some combination of prior knowledge and external testimony. For example, my mother taught me to respect myself, and I follow her advice. What is the difference between following the guidance of a moral educator and obeying the Ethics Oracle? The best kind of ethical advice prompts reflection, such as an argument made in a philosophy paper. Unthinkingly following someone's advice results in the same loss of autonomy as unthinkingly obeying the Ethics Oracle; people who merely obey orders are less autonomous than those who think for themselves.⁴⁴ This account of moral advice offers a model for human-computer symbiotic ethics. The computer should serve as a moral guide by providing arguments, just as my mom explained why I should always respect myself. Human-computer symbiotic ethics nurtures autonomy when it not only offers prescriptions for action, but also explanations for these prescriptions. Because my theorem-prover-based computational ethical system is explainable, it can guide action without sacrificing autonomy. It can make an argument for some action, instead of merely giving a verdict. Isabelle can list the facts used to show a particular action prohibited, and a human being can reflect on whether or not these principles indeed prohibit the action in question. The computer serves as a collaborator and a tool, but not as an authority, so the human being's reflective capacity and freedom is preserved.

⁴⁴This might be worrying....does this mean that soldiers who are following orders to commit atrocities are less responsible than those giving the orders? Wait maybe that's true.

A.5 Academic Ethics

There are two potential sources of the value of academic philosophy: the ethical truths uncovered and the process of a philosopher discovering these truths. Under the first view, academic philosophy is valuable because it facilitates the discovery of new ethical theories. If truths are valuable and computers can generate truths more efficiently than humans, then ethics should be fully automated. Ethical disputes often linger unresolved indefinitely, but every now and then, a theory emerges as a new classic, such as Rawls' veil of ignorance. Some academic philosophy also impacts social phenomena, like Locke's impact on global democracy. Academic philosophy often works its way into household ethics, as seen in the impact of critical race theory. This view parallels the view that ordinary ethics is valuable for its insights alone, and thus similarly implies that totally automated ethics is not only permissible, but also desirable. If ethical truths are important for their impact on society, this value is not contingent on whether a human or a computer produced these truths. If possible, computers should produce ethical theories to maximize these truths' value for society.

Another set of theories locates the value of academic ethics in the process itself and thus requires human-computer symbiosis. Just as mathematics is fun and creative for the mathematician, so is ethics for the philosopher. Many philosophers enjoy reading and writing philosophy papers. The study of philosophy builds critical thinking skills and makes philosophers more reflective. Computational ethics doesn't necessarily sacrifice any of these benefits. These benefits would be lost by fully automated ethics, but human-computer symbiotic ethics amplifies them. If a computer functions like a tool in the process of philosophical discovery, like a conversation with a colleague or a search for counterexamples, then it preserves the joy of philosophy. Moreover, computational ethics amplifies this joy by forcing ethicists to make their ideas more precise, a major goal of academic philosophy. The rigid syntax of a computer program demands much more precision than a conversation or a paper. Computational tools also offer ethicists new perspective by forcing a return to first, formal principles often avoided in ordinary philosophical inquiry. Formal ethics has been a subject of interest among ethicists, but the logical background necessary has prevented the field from taking off.

If computers can automate away the tedium of formal ethics, then this precision will be accessible to all ethicists, not just logicians. Such work has begun in metaphysics, and recent research used computational tools to find an inconsistency in Godel's ontological argument for the existence of God (Benzmüller and Woltzenlogel Paleo, 2016). The power of computational tools to force precision, perform consistency checks, and make assumptions explicit means that computers can serve as tools to help philosophers perform philosophy better.

If ethical truths offer some value to society at-large, perhaps we cannot sacrifice this value merely to preserve human philosophers' fun. A more compelling argument against fully automated ethics is that the existence of human academic philosophers offers value even to non-philosophers. People derive joy and wonder from knowing that human beings produced great ethics. Plato's *Apology* is not only a profound and insightful text, but it is also wonderful that a human being produced such a work. We derive joy from knowing that our fellow humans are capable of the kind of thought that great philosophers accomplish, just as an unimaginably beautiful work of art is more wonderful because a human being painted it. We watch the Olympics because we derive wonder and joy from human excellence. Even when admiring computational achievements, such as Google's recent success in protein folding, we admire the human who programmed the machine, not the machine itself. We can relate to humans, so the mere knowledge that great people are doing great things enriches our lives. This knowledge is part of the attraction for the thousands of tourists who visit Harvard Yard every year; this is a place where of great human achievement.⁴⁵

An even stronger argument for human-computer symbiotic ethics instead of fully automated ethics is that ethics is an inherently human subject. We study ethics because, as argued above, we have no choice but to reflect on how to live. Because reflection is such a fundamental part of being human, a world in which all ethical inquiry is automated is undesirable. Academic philosophers are professional reflectors who are partners in the human experience with us, so their ethical inquiry carries unique weight. They teach us, inspire us, and serve as examples of the kind of reflection that is constitutive of being human. Moreover, an ethical

⁴⁵Is this too like, yay Harvard

theory produced entirely by a computer is, at best, a secondary perspective; it is a computer's attempt to describe how human beings should live. Without a human component, it cannot serve as a rich and sophisticated guide for humans. If ethics is most meaningful when it is the product of human reflection, totally automated ethics destroys the field entirely but human-computer symbiosis does not. Human beings should debate the most pressing questions of human existence, and computers can serve as our aids. Thus, computational tools must supplement human ethical reasoning but cannot replace it.

end

References

- D. Abel, J. MacGlashan, and M. Littman. Reinforcement learning as a framework for ethical decision making. In *AAAI Workshop: AI, Ethics, and Society*, 2016.
- L. Alexander and M. Moore. Deontological Ethics. In E. N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Winter 2021 edition, 2021.
- M. Anderson and S. Anderson. Geneth: A general ethical dilemma analyzer. volume 1, 07 2014.
- M. Anderson and S. L. Anderson. Ethel: Toward a principled ethical eldercare robot.
- M. Anderson, S. Anderson, and C. Armen. Towards machine ethics. 07 2004.
- Aristotle. The nicomachean ethics. *Journal of Hellenic Studies*, 77:172, 1951. doi: 10.2307/628662.
- K. Arkoudas, S. Bringsjord, and P. Bello. Toward ethical robots via mechanized deontic logic. *AAAI Fall Symposium - Technical Report*, 01 2005.
- K. D. Ashley and B. M. McLaren. A cbr knowledge representation for practical ethics. In *Selected Papers from the Second European Workshop on Advances in Case-Based Reasoning*, EWCBR '94, page 181–197, Berlin, Heidelberg, 1994. Springer-Verlag. ISBN 3540603646.
- E. Awad, S. Dsouza, A. Shariff, I. Rahwan, and J.-F. Bonnefon. Universals and variations in moral decisions made in 42 countries by 70,000 participants. *Proceedings of the National Academy of Sciences*, 117(5):2332–2337, 2020. ISSN 0027-8424. doi: 10.1073/pnas.1911517117. URL <https://www.pnas.org/content/117/5/2332>.
- A. Baltag and B. Renne. Dynamic Epistemic Logic. In E. N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Winter 2016 edition, 2016.

- M. M. Bentzen and F. Lindner. A formalization of kant’s second formulation of the categorical imperative. *CoRR*, abs/1801.03160, 2018. URL <http://arxiv.org/abs/1801.03160>.
- C. Benzmüller and B. W. Paleo. Formalization, mechanization and automation of gödel’s proof of god’s existence. *CoRR*, abs/1308.4526, 2013. URL <http://arxiv.org/abs/1308.4526>.
- C. Benzmüller, X. Parent, and L. W. N. van der Torre. Designing normative theories of ethical reasoning: Formal framework, methodology, and tool support. *CoRR*, abs/1903.10187, 2019. URL <http://arxiv.org/abs/1903.10187>.
- C. Benzmüller, A. Farjami, and X. Parent. Dyadic deontic logic in hol: Faithful embedding and meta-theoretical experiments. In M. Armgardt, H. C. Nordtveit Kvernenes, and S. Rahman, editors, *New Developments in Legal Reasoning and Logic: From Ancient Law to Modern Legal Systems*, volume 23 of *Logic, Argumentation & Reasoning*. Springer Nature Switzerland AG, 2021. ISBN 978-3-030-70083-6. doi: 10.1007/978-3-030-70084-3.
- C. Benzmüller and B. Woltzenlogel Paleo. The inconsistency in gödel’s ontological argument: A success story for ai in metaphysics. 07 2016.
- N. Berberich and K. Diepold. The virtuous machine - old ethics for new technology?, 2018.
- J. C. Blanchette and T. Nipkow. *Nitpick: A Counterexample Generator for Higher-Order Logic Based on a Relational Model Finder*, volume 6172, page 131–146. Springer Berlin Heidelberg, 2010. ISBN 9783642140518. doi: 10.1007/978-3-642-14052-5_11. URL http://link.springer.com/10.1007/978-3-642-14052-5_11.
- H. Bok. *Freedom and Responsibility*. Princeton University Press, 1998.
- K. Buzzard. How do you convince mathematicians a theorem prover is worth their time? Talk at IOHK, January 2021.
- J. Carmo and A. Jones. Completeness and decidability results for a logic of contrary-to-duty conditionals. *J. Log. Comput.*, 23:585–626, 2013.

- J.-A. Cervantes, L.-F. Rodríguez, S. López, and F. Ramos. A biologically inspired computational model of moral decision making for autonomous agents. In *2013 IEEE 12th International Conference on Cognitive Informatics and Cognitive Computing*, pages III–III7, 2013. doi: 10.1109/ICCI-CC.2013.6622232.
- R. M. Chisholm. Contrary-to-duty imperatives and deontic logic. *Analysis (Oxford)*, 24(2): 33–36, 1963. ISSN 0003-2638.
- D. Christensen. Epistemic self-respect. *Proceedings of the Aristotelian Society*, 107(1pt3):319–337, 2007. doi: 10.1111/j.1467-9264.2007.00224.x.
- C. Cloos. The utilibot project: An autonomous mobile robot based on utilitarianism. *AAAI Fall Symposium - Technical Report*, 01 2005.
- M. J. Cresswell and G. E. Hughes. *A New Introduction to Modal Logic*. Routledge, 1996.
- D. Davenport. Moral mechanisms. *Philosophy and Technology*, 27(1):47–60, 2014. doi: 10.1007/s13347-013-0147-2.
- L. Dennis, M. Fisher, M. Slavkovik, and M. Webster. Formal verification of ethical choices in autonomous systems. *Robotics and Autonomous Systems*, 77:1–14, 2016. ISSN 0921-8890. doi: <https://doi.org/10.1016/j.robot.2015.11.012>. URL <https://www.sciencedirect.com/science/article/pii/S0921889015003000>.
- P. Dietrichson. When is a maxim fully universalizable ? 55(1-4):143–170, 1964. doi: doi: 10.1515/kant.1964.55.1-4.143. URL <https://doi.org/10.1515/kant.1964.55.1-4.143>.
- J. Driver. The History of Utilitarianism. In E. N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Winter 2014 edition, 2014.
- K. Ebels-Duggan. *Kantian Ethics*, chapter Kantian Ethics. Continuum, 2012.
- V. Eubanks. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. St. Martin’s Press, 2018.

- P. Foot. The problem of abortion and the doctrine of the double effect. *Oxford Review*, 5: 5–15, 1967.
- D. Fuenmayor and C. Benzmüller. Formalisation and evaluation of alan gewirth’s proof for the principle of generic consistency in isabelle/hol. *Archive of Formal Proofs*, Oct. 2018. ISSN 2150-914x. <https://isa-afp.org/entries/GewirthPGCProof.html>, Formal proof development.
- I. Gabriel. Artificial intelligence, values, and alignment. *Minds and Machines*, 30(3):411–437, Sep 2020. ISSN 1572-8641. doi: 10.1007/s11023-020-09539-2. URL <http://dx.doi.org/10.1007/s11023-020-09539-2>.
- N. S. Govindarajulu and S. Bringsjord. On automating the doctrine of double effect. *CoRR*, abs/1703.08922, 2017. URL <http://arxiv.org/abs/1703.08922>.
- M. Guarini. Particularism and the classification and reclassification of moral cases. *IEEE Intelligent Systems*, 21(4):22–28, 2006. doi: 10.1109/MIS.2006.76.
- J. Harrison, J. Urban, and F. Wiedijk. History of interactive theorem proving. In *Computational Logic*, 2014.
- J. Hintikka. *Knowledge and Belief*. Cornell University Press, 1962.
- D. Hume. *An Enquiry Concerning Human Understanding and Other Writings*. Cambridge University Press, 2007.
- R. Hursthouse and G. Pettigrove. Virtue Ethics. In E. N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Winter 2018 edition, 2018.
- L. Jiang, J. D. Hwang, C. Bhagavatula, R. L. Bras, M. Forbes, J. Borchardt, J. Liang, O. Etzioni, M. Sap, and Y. Choi. Delphi: Towards machine ethics and norms, 2021.
- R. Johnson and A. Cureton. Kant’s Moral Philosophy. In E. N. Zalta, editor, *The Stanford*

- Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Spring 2021 edition, 2021.
- I. Kant. *Groundwork of the Metaphysics of Morals*. Cambridge University Press, Cambridge, 1785.
- I. Kant. *Introduction*, pages ix–xxix. Cambridge Texts in the History of Philosophy. Cambridge University Press, 2 edition, 2017. doi: 10.1017/9781316091388.002.
- J. Kemp. Kant’s examples of the categorical imperative. *The Philosophical Quarterly* (1950-), 8(30):63–71, 1958. ISSN 00318094, 14679213. URL <http://www.jstor.org/stable/2216857>.
- D. Kirchner, C. Benz Müller, and E. N. Zalta. Computer science and metaphysics: A cross-fertilization. *CoRR*, abs/1905.00787, 2019. URL <http://arxiv.org/abs/1905.00787>.
- P. Kitcher. What is a maxim? *Philosophical Topics*, 31(1/2):215–243, 2003. doi: 10.5840/philtopics2003311/29.
- P. Kitcher. Kant’s argument for the categorical imperative. *Nous*, 38, 2004.
- M. Kohl. Kant and ‘ought implies can’. *The Philosophical Quarterly* (1950-), 65(261):690–710, 2015. ISSN 00318094, 14679213. URL <http://www.jstor.org/stable/24672780>.
- R. C. Koons. *Paradoxes of Belief and Strategic Rationality*. Cambridge Studies in Probability, Induction and Decision Theory. Cambridge University Press, 1992. doi: 10.1017/CBO9780511625381.
- C. Korsgaard. Kant’s Formula of Universal Law. *Pacific Philosophical Quarterly*, 66:24–47, 1985.
- C. Korsgaard. The Right to Lie: Kant on Dealing with Evil. *Philosophy and Public Affairs*, 15:325–249, 1986.
- C. Korsgaard. *Groundwork of the Metaphysics of Morals*, chapter Introduction. Cambridge University Press, Cambridge, 2012.

- C. M. Korsgaard. Acting for a reason. *Danish Yearbook of Philosophy*, 40(1):11–35, 2005. doi: 10.1163/24689300\0400103.
- C. M. Korsgaard and O. O'Neill. *The Sources of Normativity*. Cambridge University Press, 1996. doi: 10.1017/CBO9780511554476.
- M. Kroy. A partial formalization of kant's categorical imperative. an application of deontic logic to classical moral philosophy. *Kant-Studien*, 67(1-4):192–209, 1976. doi: doi:10.1515/kant.1976.67.1-4.192. URL <https://doi.org/10.1515/kant.1976.67.1-4.192>.
- G. W. Leibniz. On universal synthesis and analysis, or the art of discovery and judgment: 1679(?). In *Philosophical Papers and Letters*, The New Synthese Historical Library, pages 229–234. Springer Netherlands, Dordrecht. ISBN 9789027706935.
- D. Lewis. Causation. *Journal of Philosophy*, 70(17):556–567, 1973a. doi: 10.2307/2025310.
- D. Lewis. *Counterfactuals*. Blackwell, 1973b.
- J. C. R. Licklider. Man-computer symbiosis. *IRE Transactions on Human Factors in Electronics*, HFE-1(1):4–11, 1960. doi: 10.1109/THFE2.1960.4503259.
- P. Lin, K. Abney, and G. A. Bekey. *Robotics, Ethical Theory, and Metaethics: A Guide for the Perplexed*, pages 35–52. 2012.
- P. Lukowicz. The challenge of human centric ai. *Digitale Welt*, 3:9–10, 10 2019. doi: 10.1007/s42354-019-0200-0.
- E. McRae. Equanimity and intimacy: A buddhist-feminist approach to the elimination of bias. *Sophia*, 52(3):447–462, 2013. doi: 10.1007/s11841-013-0376-y.
- R. MONTAGUE. Universal grammar. *Theoria*, 36(3):373–398, 1970. doi: <https://doi.org/10.1111/j.1755-2567.1970.tb00434.x>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1755-2567.1970.tb00434.x>.
- G. E. Moore. *Principia Ethica*. Dover Publications, 1903.

- T. Nipkow, L. C. Paulson, and M. Wenzel. *Isabelle/HOL: A Proof Assistant for Higher Order Logic*. Springer-Verlag Berlin Heidelberg, Berlin, 2002.
- O. O’Neill. Universal laws and ends-in-themselves. *The Monist*, 72(3):341–361, 1989. ISSN 00269662. URL <http://www.jstor.org/stable/27903145>.
- O. O’Neill. *Constructions of Reason: Explorations of Kant’s Practical Philosophy*. Cambridge University Press, 1990. doi: 10.1017/CBO9781139173773.
- O. O’Neill. *Acting on Principle: An Essay on Kantian Ethics*. Cambridge University Press, 2013.
- L. Paulson and J. Blanchette. Three years of experience with sledgehammer, a practical link between automatic and interactive theorem provers. 02 2015. doi: 10.29007/tnfd.
- L. C. Paulson. A generic tableau prover and its integration with isabelle. *J. Univers. Comput. Sci.*, 5:73–87, 1999.
- T. M. Powers. Prospects for a kantian machine. *IEEE Intelligent Systems*, 21(4):46–51, 2006. doi: 10.1109/MIS.2006.77.
- E. Puiutta and E. M. Veith. Explainable reinforcement learning: A survey, 2020.
- J. Rawls. Kantian constructivism in moral theory. *The Journal of Philosophy*, 77(9):515–572, 1980. ISSN 0022362X. URL <http://www.jstor.org/stable/2025790>.
- J. Rawls. Lectures on the history of moral philosophy. *Critica*, 35(104):121–145, 2000.
- S. Roush. Second guessing: A self-help manual. *Episteme*, 6(3):251–268, 2009. doi: 10.3366/E1742360009000690.
- D. Rönnedal. Contrary-to-duty paradoxes and counterfactual deontic logic. *Philosophia*, 47, 09 2019. doi: 10.1007/s11406-018-0036-0.
- D. Scott. Advice on modal logic. In K. Lambert, editor, *Philosophical Problems in Logic: Some Recent Developments*, pages 143–173. D. Reidel, 1970.

- J. R. Silber. Procedural formalism in kant's ethics. *The Review of Metaphysics*, 28(2):197–236, 1974. ISSN 00346632. URL <http://www.jstor.org/stable/20126622>.
- W. Sinnott-Armstrong. Consequentialism. In E. N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Fall 2021 edition, 2021.
- K. Solt. Deontic alternative worlds and the truth-value of 'oa'. *Logique et Analyse*, 27(107): 349–351, 1984. ISSN 00245836, 22955836. URL <http://www.jstor.org/stable/44084096>.
- N. F. Stang. Kant's Transcendental Idealism. In E. N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, Spring 2021 edition, 2021.
- A. Stephenson, M. Sergot, and R. Evans. Formalizing kant's rules: a logic of conditional imperatives and permissives. *Journal of Philosophical Logic*, 49, November 2019. URL <https://eprints.soton.ac.uk/432344/>.
- D. Tafani. imperativo categorico come algoritmo. kant e l'etica delle macchine. *Sistemi intelligenti, Rivista quadrimestrale di scienze cognitive e di intelligenza artificiale*, (2/2021): 377–392, 2021. ISSN 1120-9550. doi: 10.1422/101195. URL <https://www.rivisteweb.it/doi/10.1422/101195>.
- J. Timmermann. Kantian dilemmas? moral conflict in kant's ethical theory. *Archiv für Geschichte der Philosophie*, 95(1):36–64, 2013. doi: doi:10.1515/agph-2013-0002. URL <https://doi.org/10.1515/agph-2013-0002>.
- S. Tolmeijer, M. Kneer, C. Sarasua, M. Christen, and A. Bernstein. Implementations in machine ethics. *ACM Computing Surveys*, 53(6):1–38, Feb 2021. ISSN 1557-7341. doi: 10.1145/3419633. URL <http://dx.doi.org/10.1145/3419633>.
- J. D. Velleman. *A Brief Introduction to Kantian Ethics*, page 16–44. Cambridge University Press, 2005. doi: 10.1017/CBO9780511498862.002.
- J. Vickers. I believe it, but soon i'll not believe it any more: Scepticism, empiricism, and reflection. *Synthese*, 124:155–174, 08 2000. doi: 10.1023/A:1005213608394.

- J. Vincent. The ai oracle of delphi uses the problems of reddit to offer dubious moral advice. 2021.
- W. Wallach and C. Allen. *Moral Machines: Teaching Robots Right From Wrong*. Oxford University Press, 2008.
- A. Winfield, C. Blum, and W. Liu. Towards an ethical robot: Internal models, consequences and ethical action selection. volume 8717, 09 2014. ISBN 978-3-319-10400-3. doi: 10.1007/978-3-319-10401-0_8.
- A. W. Wood. *Kant's Ethical Thought*. Cambridge University Press, 1999.
- T. Yamazaki, J. Igarashi, and H. Yamaura. Human-scale brain simulation via supercomputer: A case study on the cerebellum. *Neuroscience*, 462:235–246, 2021. ISSN 0306-4522. doi: <https://doi.org/10.1016/j.neuroscience.2021.01.014>. URL <https://www.sciencedirect.com/science/article/pii/S030645222100021X>. In Memoriam: Masao Ito—A Visionary Neuroscientist with a Passion for the Cerebellum.
- V. Zahoransky and C. Benz Müller. Modelling the us constitution to establish constitutional dictatorship. 10 2020.